

# Relational Databases: Internals

## Responsible:

Dr. Tegawendé F. BISSYANDE

[tegawende.bissyande@uni.lu](mailto:tegawende.bissyande@uni.lu)

## Course Author:

Eduardo Cunha de Almeida

[eduardo@inf.ufpr.br](mailto:eduardo@inf.ufpr.br)

## Teacher Assistant:

Médéric Hurier

[mederic.hurier@uni.lu](mailto:mederic.hurier@uni.lu)

# Course Features

- Sep. 20th - **Introduction to Big Data**

## **Part 1. Databases and Query Models for Big Data**

- Sep. 27th - **Relational Databases: Reminders**
- **Oct. 4th - Relational Databases: Internals**
- Oct. 11th - **NoSQL Databases**
- Oct. 18th - **MapReduce Model**
- Oct. 25th - **Hadoop and Spark**
- Nov. 8th - **Datalog Model**

## **Part 2. Data Analysis and Machine Learning**

- Nov. 15th - **Statistics and Probabilities**
- Nov. 22th - **Communication and Visualization**
- Nov. 29th - **Features Engineering and Supervised Learning**
- Dec. 5st - **Unsupervised and Reinforcement Learning**
- Dec. 12th - **Homework Time**


# Section Features

- History
- Database Layout
- Query Processing
- Query Optimization
- Transactions

# History

# History

1970's




- INGRES Project at Berkeley

INGRES corp, Sybase, MS SQL Server, PostgreSQL

- System-R Project at IBM

DB2, Non-StopSQL, HP Allbase

# History



- DB2
- Oracle
- Informix
- Teradata
- Sybase

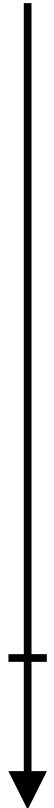
**SQL as the standard query language !**

# History

↓  
1990's

- Postgres
- MySQL
- Illustra (from Postgres)
- BerkeleyDB (embedded K/V)

# History

- 
- 2000's
- MonetDB (as open-source)
  - Vertica
  - Greenplum
  - EnterpriseDB
  - Infobright



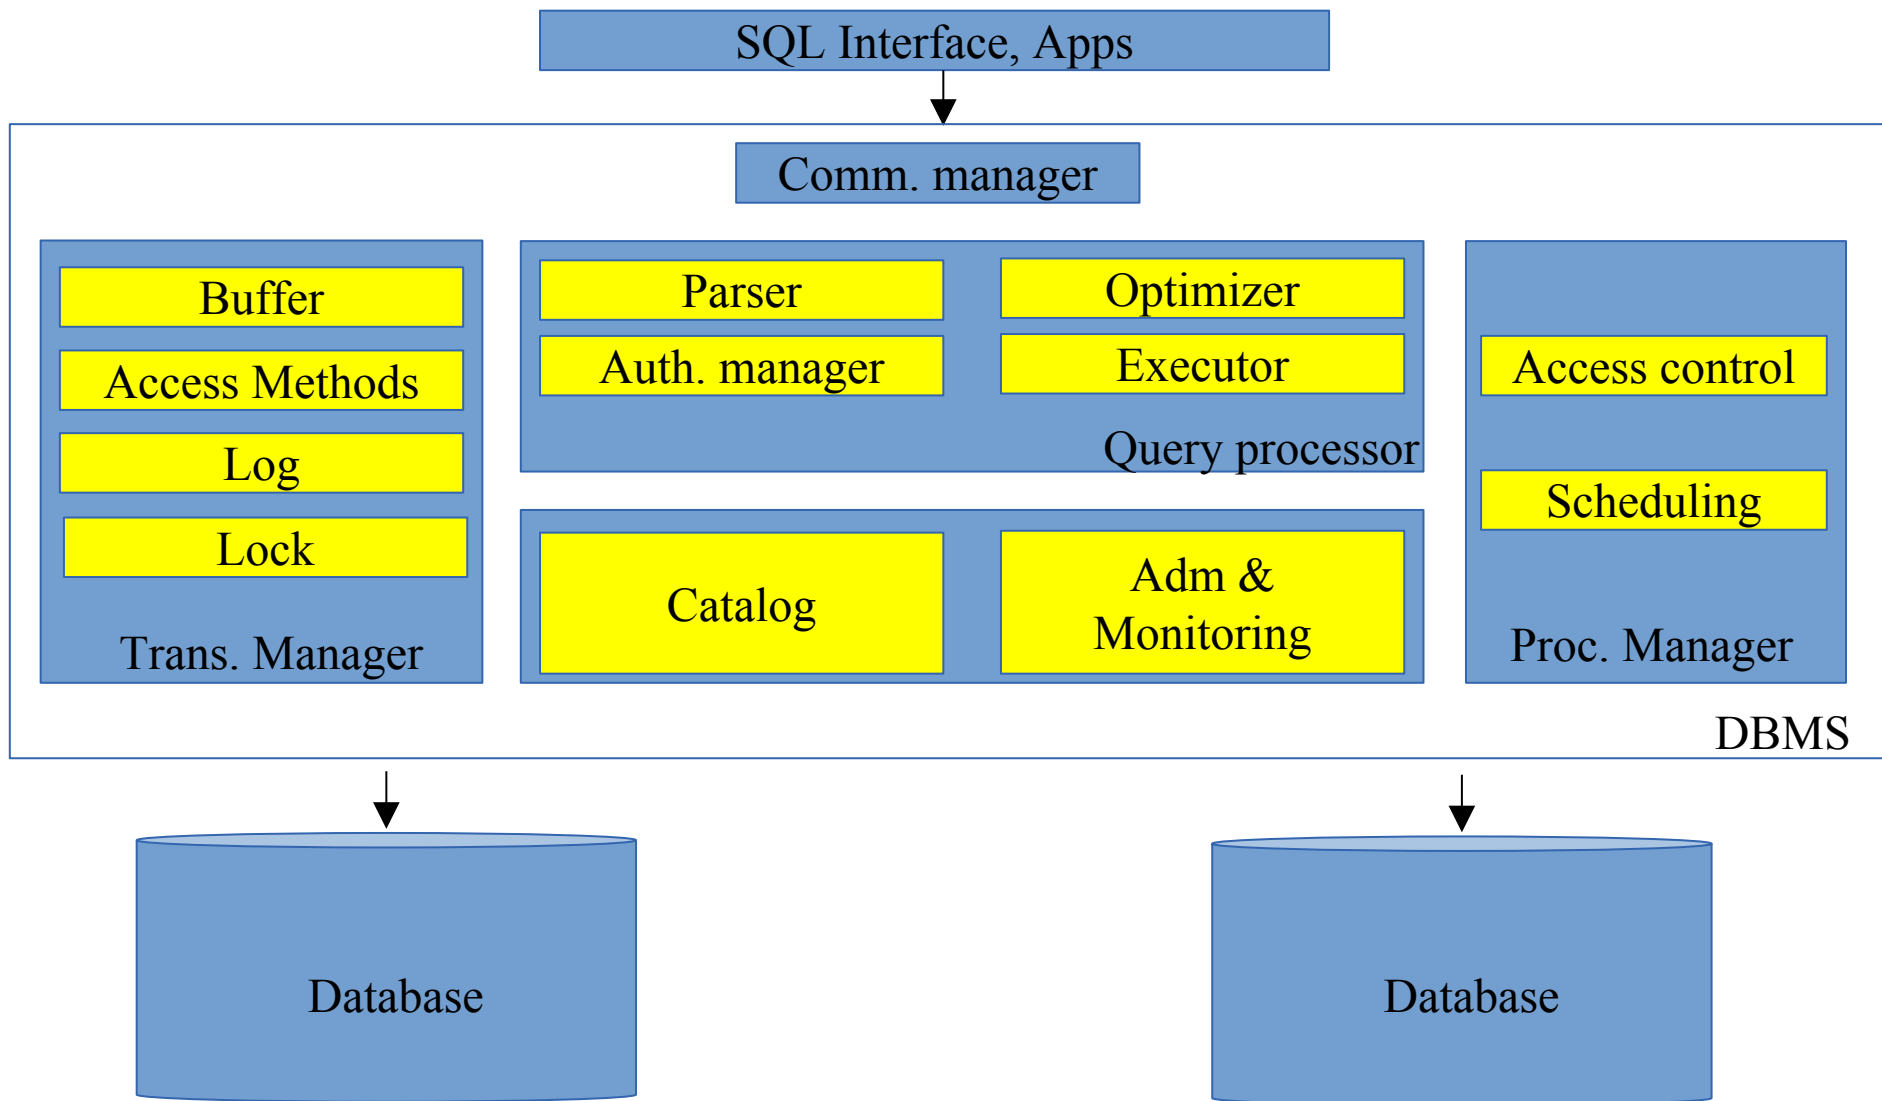
# History

↓ 2010's

- VoltDB (from H-Store)
- Facebook Presto
- Google F1
- Google Megastore

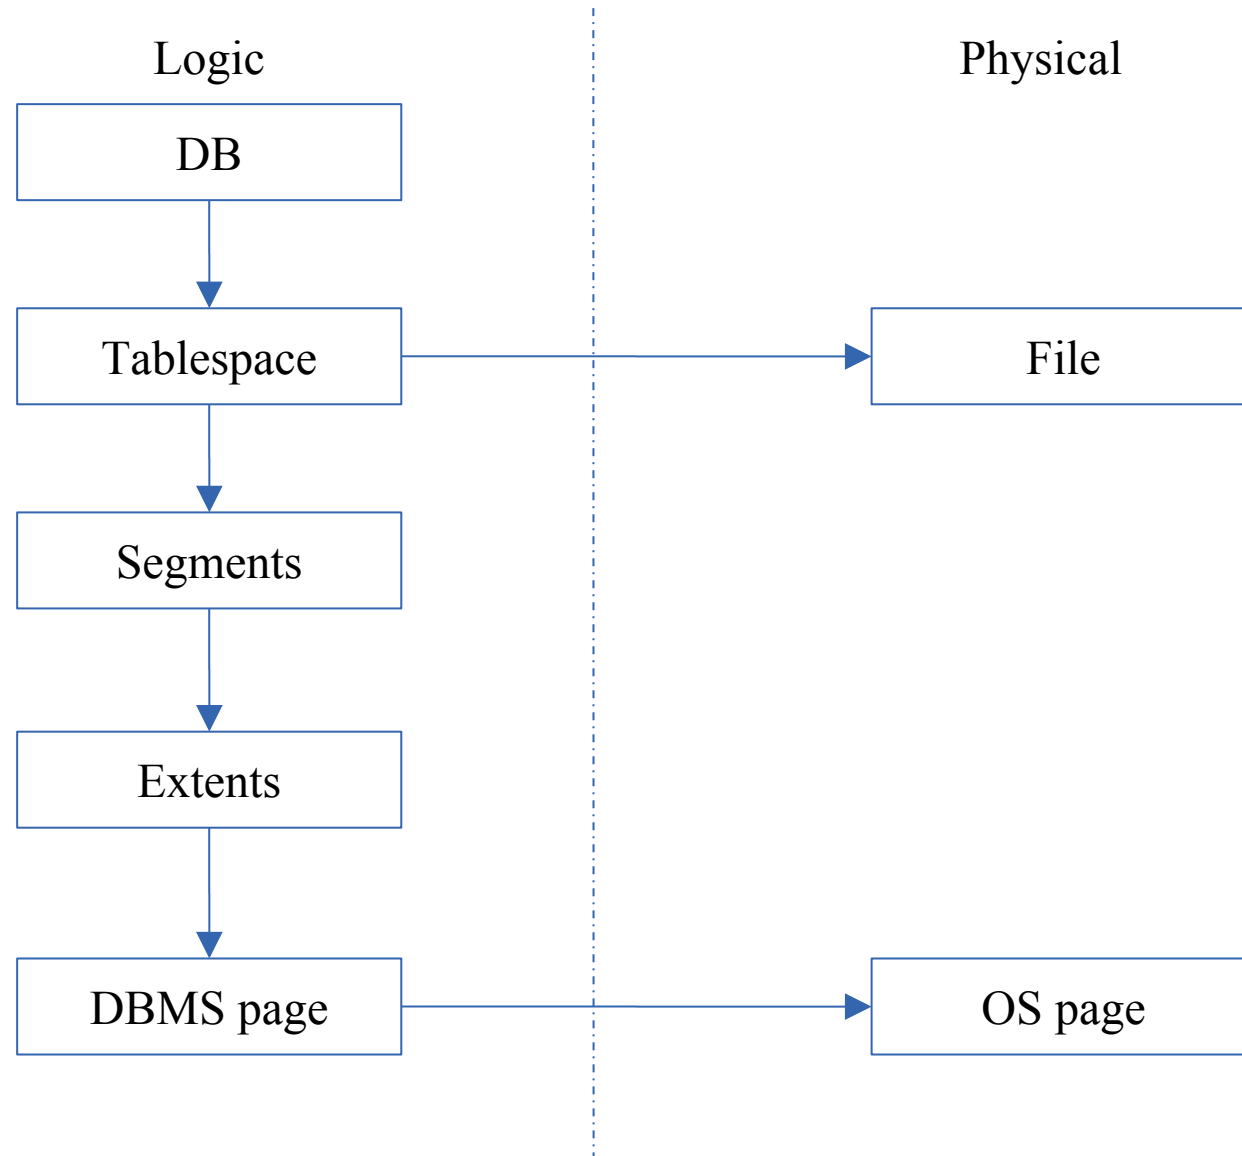
**NewSQL and NoSQL architectures !**

# General DBMS Architecture




# Database Layout

# Database Layout



Tab customers


Index





Tab products

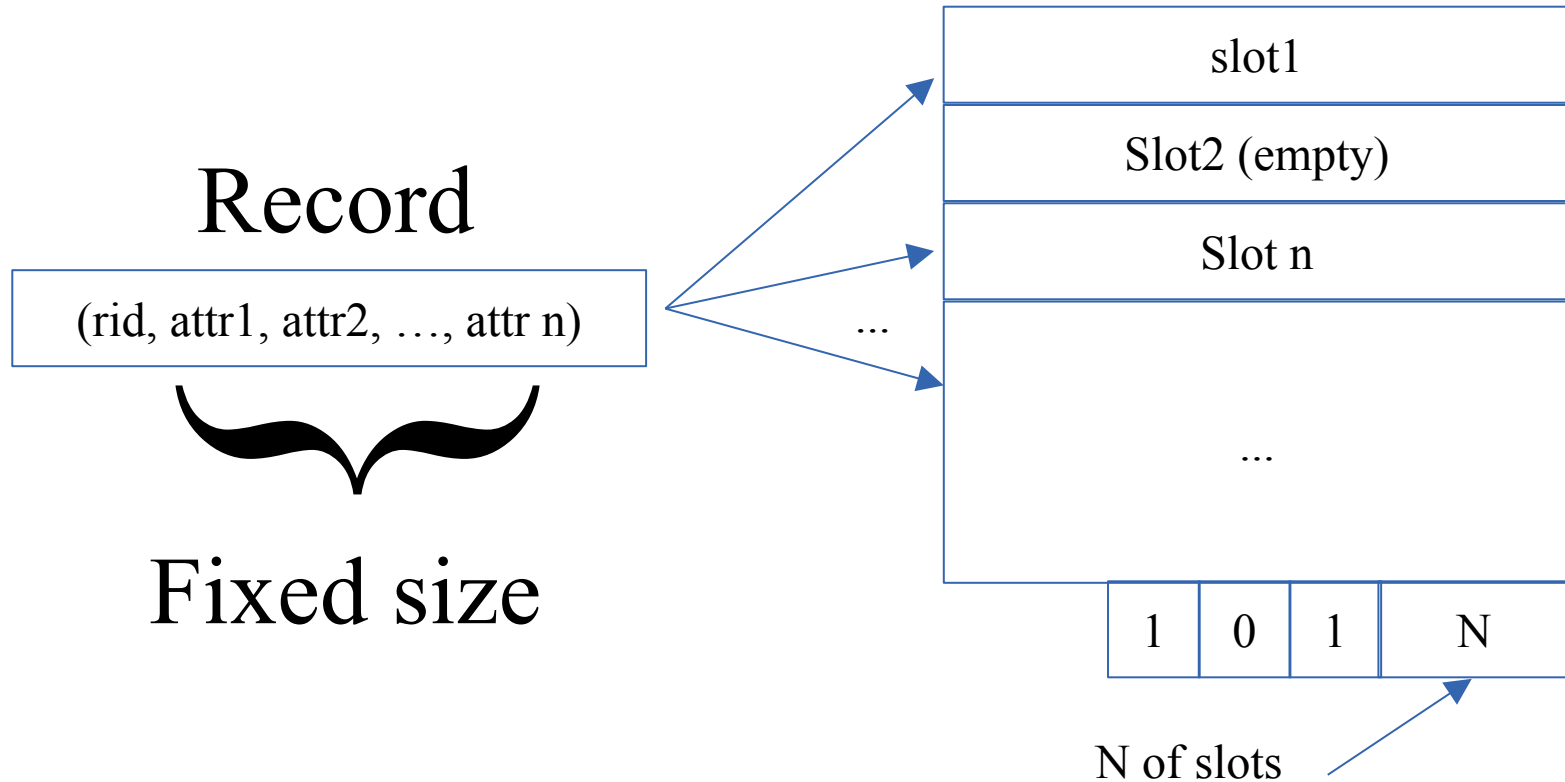

Tab customers


Tab products


Index on prodID

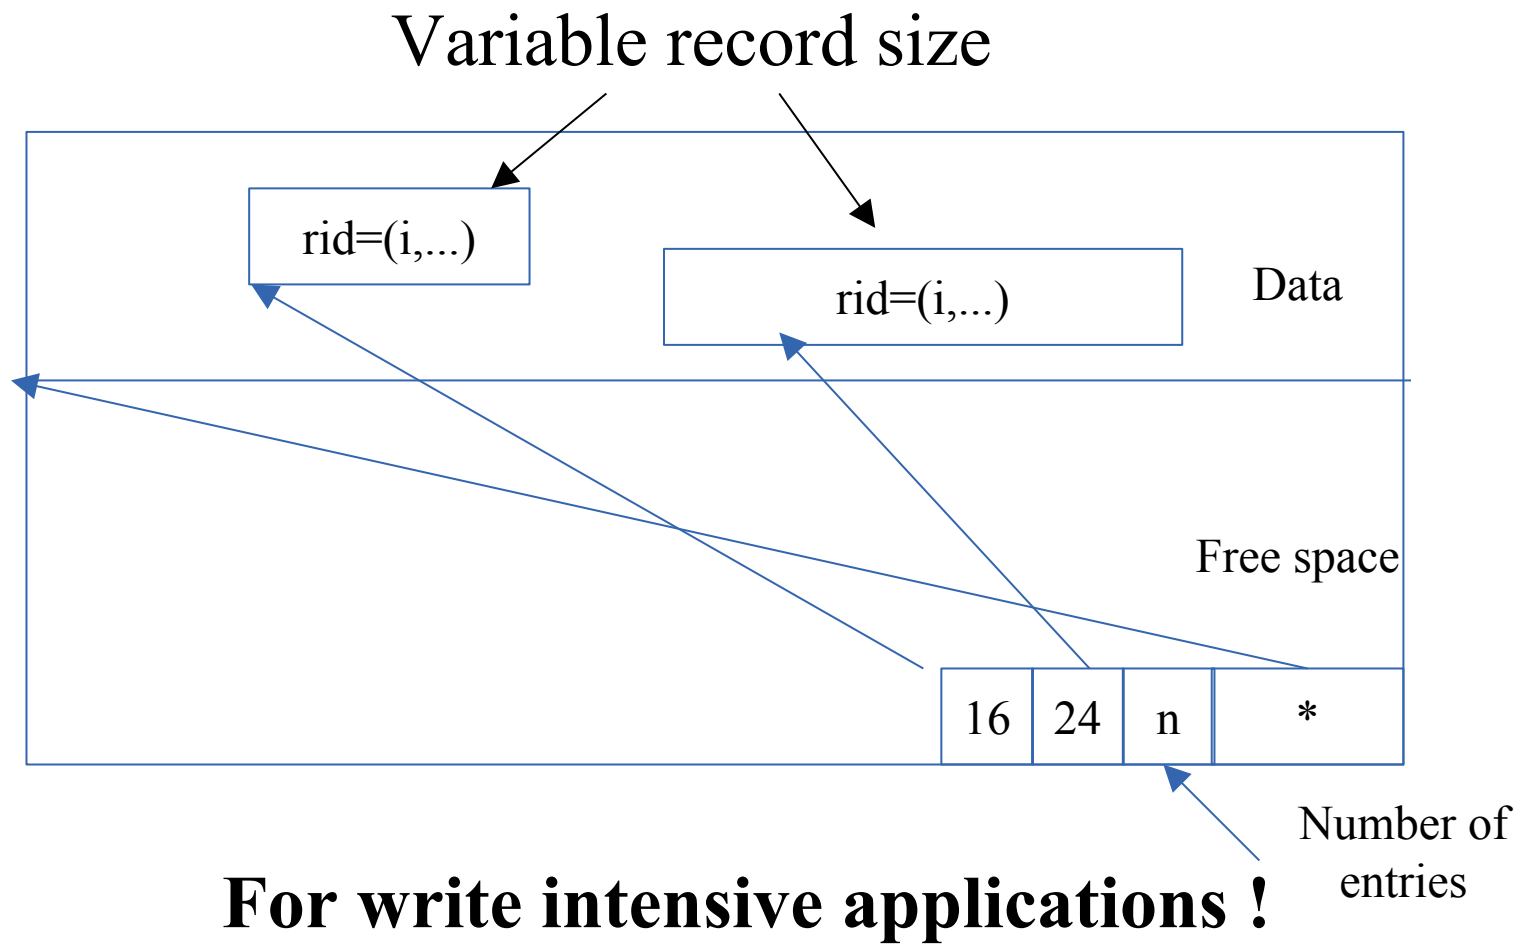


# Fixed size pages



# Variable size pages

## N-ary Storage Model (NSM)



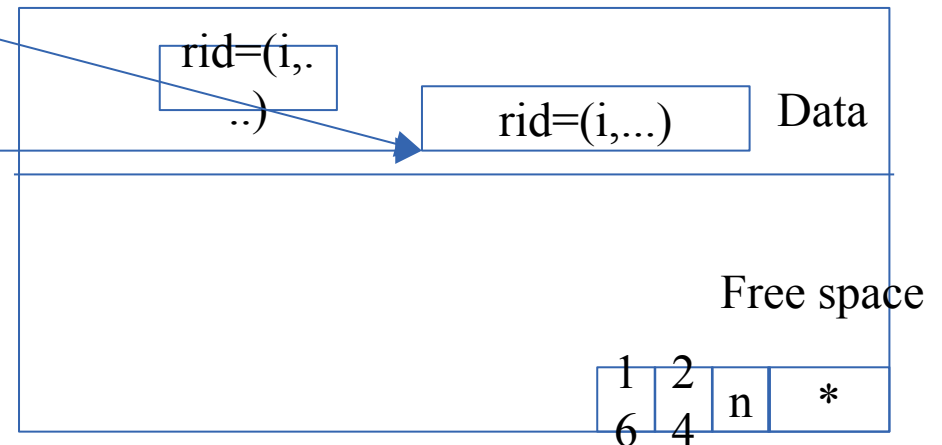


# NSM Example

INSERT INTO T  
VALUES(2, 'MARIA', 18, CURITIBA, F);

SELECT \*  
FROM CUSTOMER  
WHERE ID=2;

SELECT AVG(AGE)  
FROM CUSTOMER;

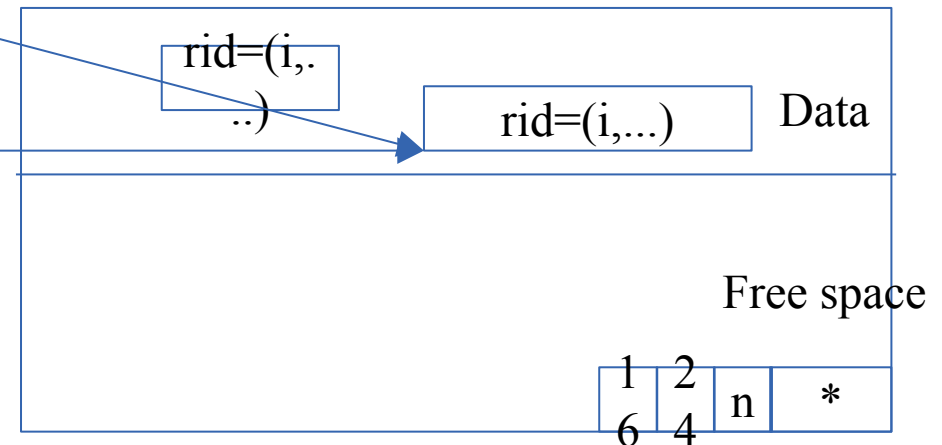


# NSM Example

INSERT INTO T  
VALUES(2, 'MARIA', 18, CURITIBA, F);

SELECT \*  
FROM CUSTOMER  
WHERE ID=2;

SELECT AVG(AGE)  
FROM CUSTOMER;



- Few write operations
- Read the entire tuple at once
- Aggregations may read the entire DB!!

# Columnar page

Decomposition Storage Model (DSM) [Ailamaki, VLDB02]

Page: Attr1	Page: Attr2	Page: Attr3
(rid1,attr1)	(rid1,attr2)	(rid1,attr3)
(rid2,attr1)	(rid2,attr2)	(rid2,attr3)
...	...	...
(rid_n,attr1)	(rid_n,attr2)	(rid_n,attr3)

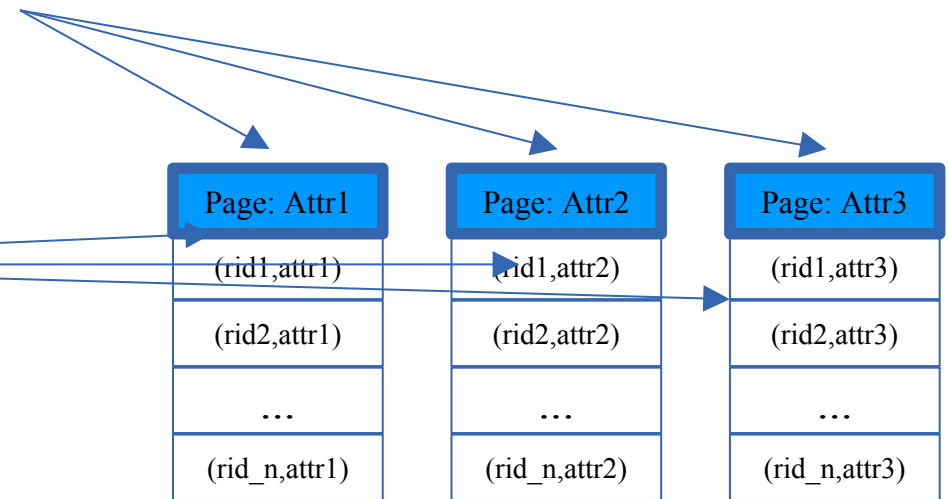
**Improves aggregation and compression  
For read intensive applications !**

# DSM Example

```
INSERT INTO T  
VALUES(2, 'MARIA', 18, CURITIBA, F);
```

```
SELECT *  
FROM CUSTOMER  
WHERE ID=2;
```

```
SELECT AVG(AGE)  
FROM CUSTOMER;
```

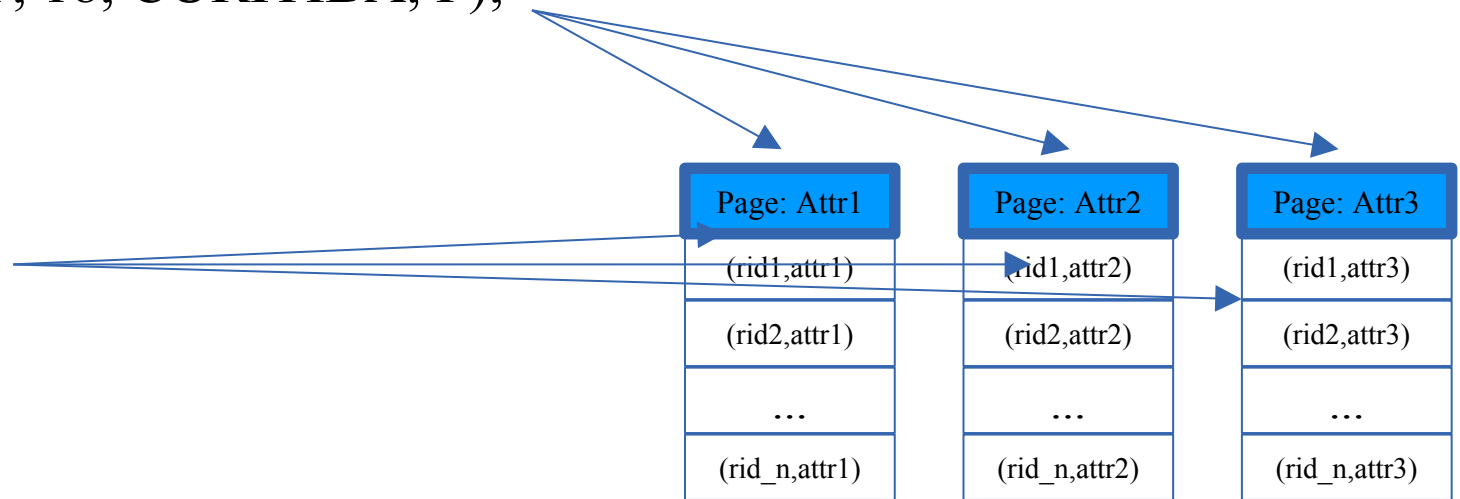


# DSM Example

INSERT INTO T  
VALUES(2, 'MARIA', 18, CURITIBA, F);

SELECT \*  
FROM CUSTOMER  
WHERE ID=2;

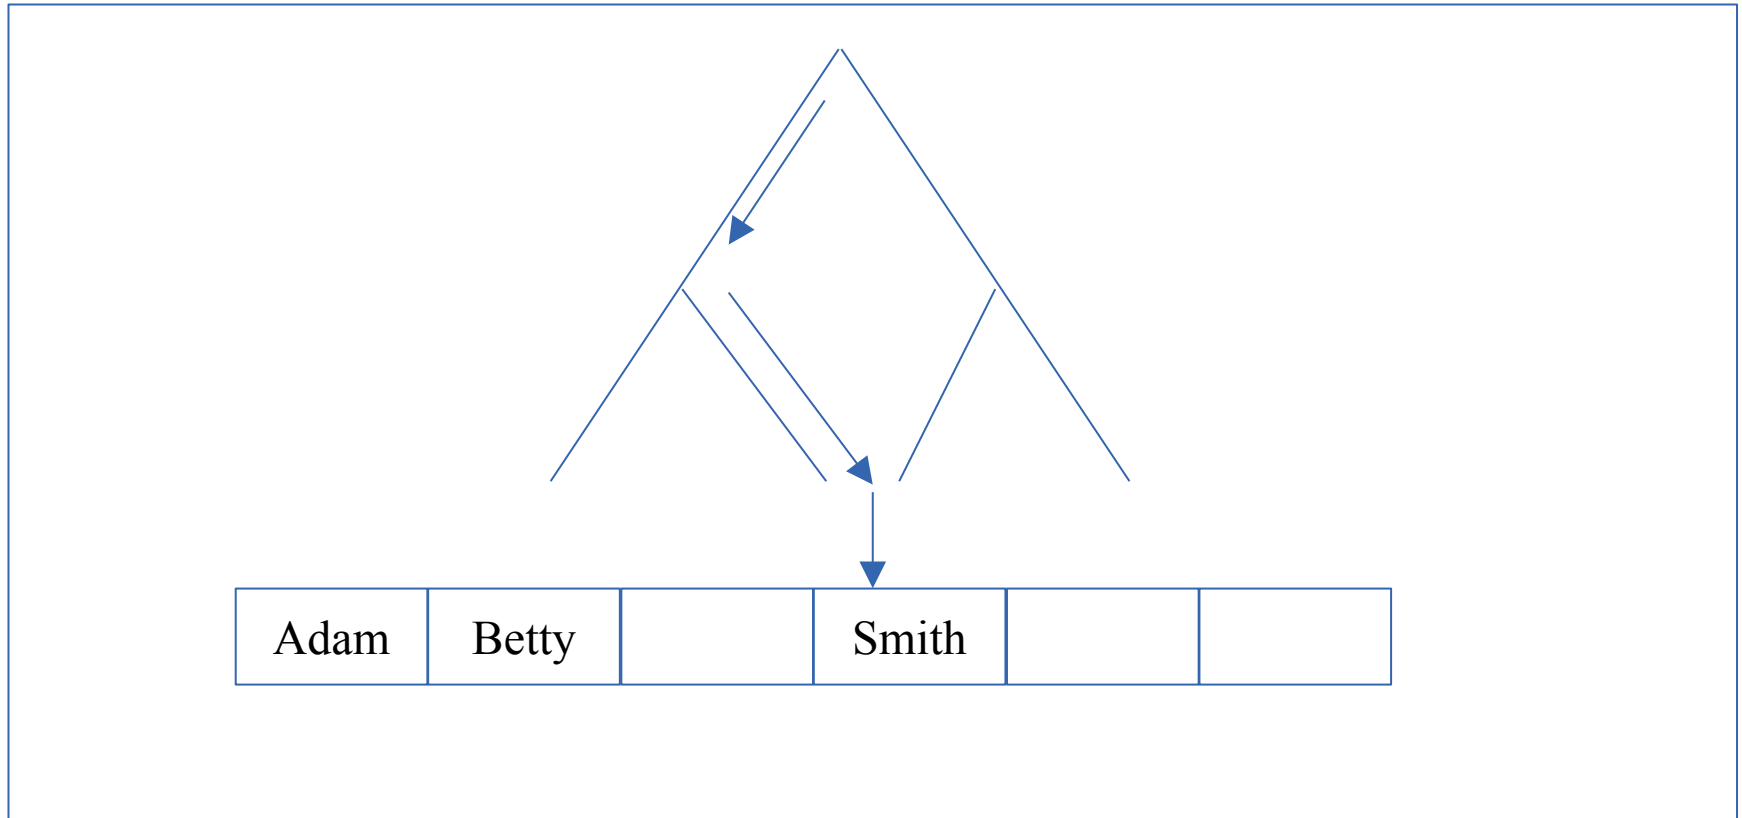
SELECT AVG(AGE)  
FROM CUSTOMER;

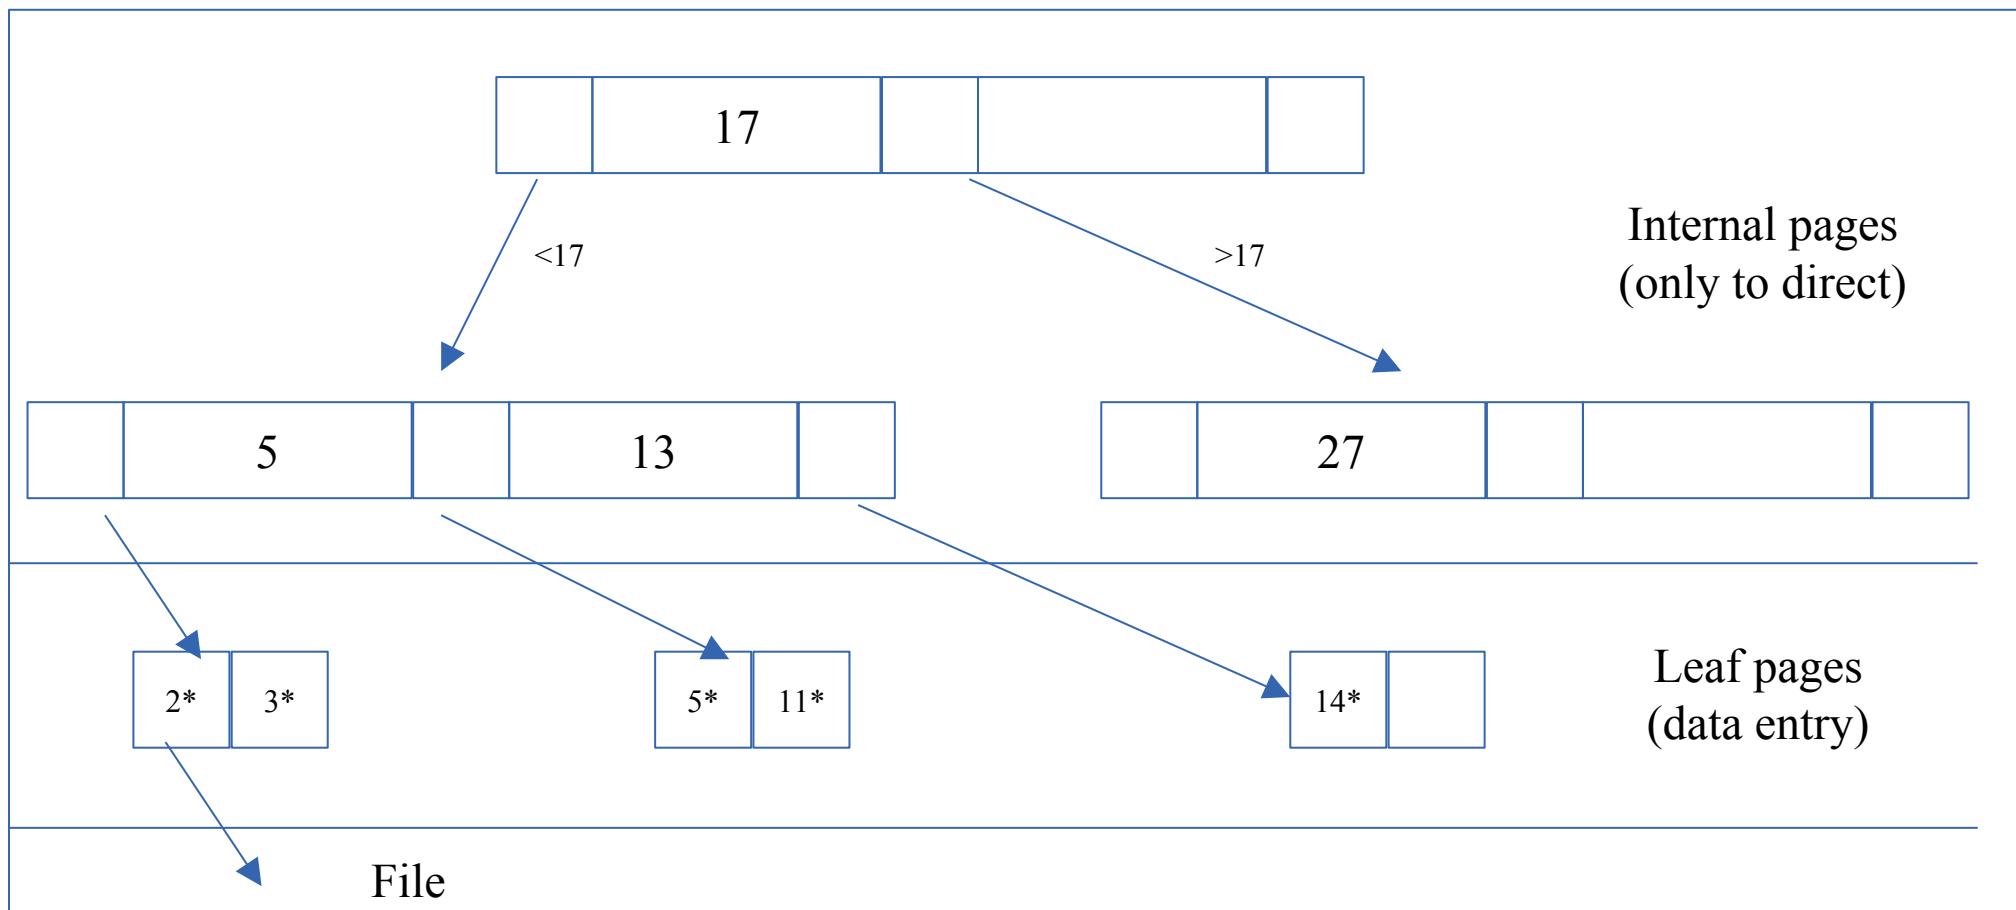



- More write operations than NSM
- Cannot read an entire tuple at once
- Aggregations are fast, do not read the entire DB!!

# Database Index


# Access method: B-tree Index





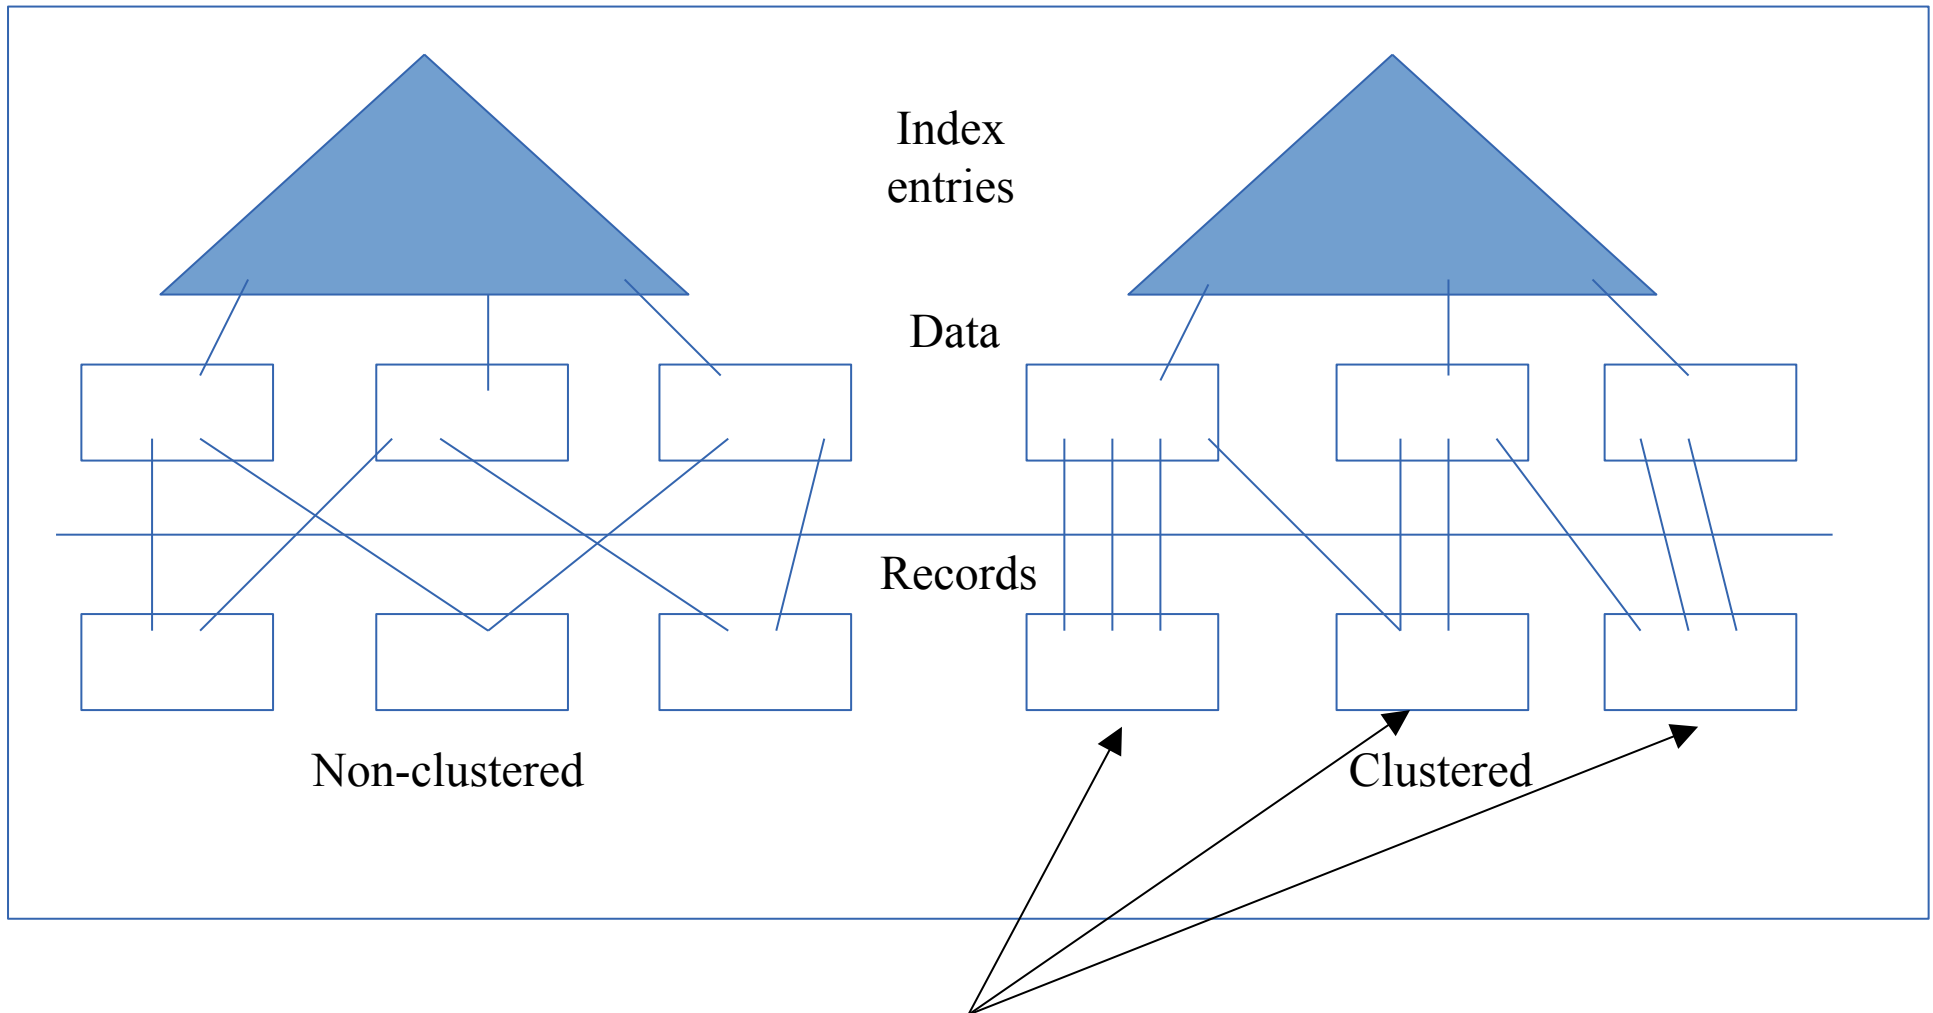


# Access method: Hash Index



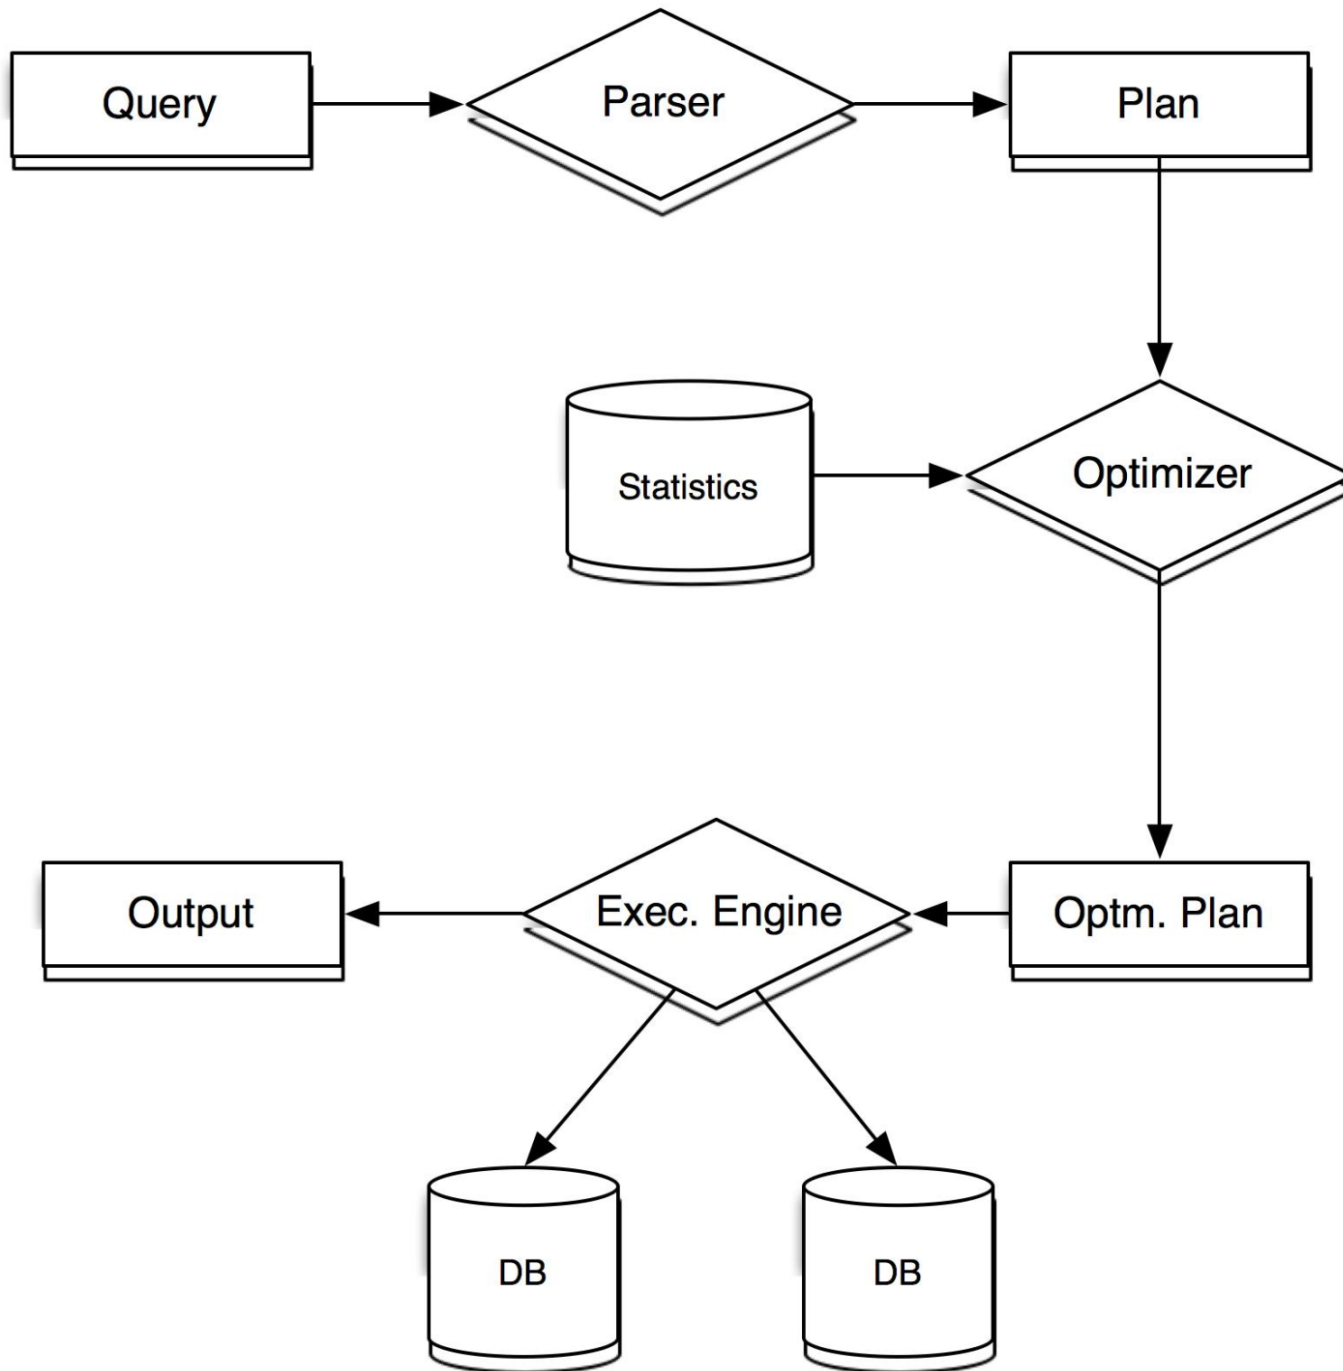
If insert 5\*, then 5=101, bucket '01' (read backwards)

# Clustered VS Non-clustered index



**Physically sorted:** only one clustered index per table

# Query Processing



(T.F. Bissyandé & M. Hurier)

# Query processing

Query

Query expression

SELECT BALANCE  
FROM ACCOUNT  
WHERE BALANCE > 2500;

$\Pi$  BALANCE ( $\sigma$   
BALANCE > 2500  
(ACCOUNT))

$\sigma$  BALANCE > 2500

$\Pi$  BALANCE (ACCOUNT)

# Query processing

Query

Query expression

SELECT BALANCE  
FROM ACCOUNT  
WHERE BALANCE > 2500;

$\Pi$  BALANCE ( $\sigma$   
BALANCE > 2500  
(ACCOUNT))

$\sigma$  BALANCE > 2500


$\Pi$  BALANCE (ACCOUNT)

May have n plans!!!

**How does the DBMS rewrite a query?**

# Query Optimization


# Query rewrite



**Rule 1:** Selection operations are commutative




# Query rewrite




**Rule 2:** Join operations are commutative  
if the order of the attributes is not taken into account

# Query rewrite



**Rule 3:** Natural-joins are associative


# Query rewrite



**Rule 4:** Natural-joins are associative.


For theta-joins, this only holds  
if the following join involves attributes only from T2 and T3.

# Query rewrite



**Rule 5:** Join distributes when all the attributes in the selection condition involve only the attributes of one of the expressions

# Query rewrite



**Rule 6:** Conjunctive selection operations can be split into a sequence of individual selections

Further rules can be found  
in the readings

# Database Catalog

**The catalog is a meta-database that stores informations about:**

## **DB and DBMS:**

Data structures, Buffer and Page sizes, Indices, Views

## **Statistics:**

Cardinality of tables and indices, Domain values

Page sizes for tables and indices

# Database Catalog

## ALL\_TABLES view from the Oracle catalog

### A.1.10 ALL\_TABLES

This view provides the following information about tables accessible to the user. The parameters for this view are listed in [Table A-10](#):

**Table A-10 ALL\_TABLES Parameter**

Column	Datatype	NULL ALLOWED	Description
OWNER	VARCHAR2(128)	No	User name of the owner of the table.
TABLE_NAME	VARCHAR2(128)	No	Name of the table.
TABLESPACE_NAME	VARCHAR2(128)	Yes	Name of the catalog or database file containing the table.
CLUSTER_NAME*	VARCHAR2(128)	Yes	Name of the cluster, if any, to which the table belongs.
PCT_FREE*	NUMBER(10)	Yes	Minimum percentage of free space in a block.
PCT_USED*	NUMBER(10)	Yes	Minimum percentage of used space in a block.
INI_TRANS*	NUMBER(10)	Yes	Initial number of transactions.
MAX_TRANS*	NUMBER(10)	Yes	Maximum number of transactions.
INITIAL_EXTENT*	NUMBER(10)	Yes	Size of the initial extent in bytes.
NEXT_EXTENT*	NUMBER(10)	Yes	Size of secondary extents in bytes.
MIN_EXTENTS*	NUMBER(10)	Yes	Minimum number of extents allowed in the segment.
MAX_EXTENTS*	NUMBER(10)	Yes	Maximum number of extents allowed in the segment.
PCT_INCREASE*	NUMBER(10)	Yes	Percentage increase in extent size.
BACKED_UP*	VARCHAR2(1)	Yes	If the table was backed up since last change.
NUM_ROWS*	NUMBER(10)	Yes	Number of rows in the table.
BLOCKS*	NUMBER(10)	Yes	Number of data blocks allocated to the table.
EMPTY_BLOCKS*	NUMBER(10)	Yes	Number of data blocks allocated to the table that contain no data.
AVG_SPACE*	NUMBER(10)	Yes	Average amount of free space (in bytes) in a data block allocated to the table.
CHAIN_CNT*	NUMBER(10)	Yes	Number of rows in the table that are chained from one data block to another, or that have migrated to a new block, requiring a link to preserve the old ROWID.
AVG_ROW_LEN*	NUMBER(10)	Yes	Average length of a row in the table in bytes.



# Operators (I/O cost for Selection)

- **Scan on unsorted data:**
  - $O(M)$ , where  $M$  = number of pages
- **Scan on sorted data:**
  - $O(\log M)$ , where  $M$  = number of pages
- **Index scan (clustered index):**
  - $O(h+1)$ , where  $h$  = height of the index tree
- **Index scan (non-clustered index):**
  - $O(h+n)$ , where  $n$  = number of fetched tuples
  - $n = M$  if the tuples are spread across all pages
    - (worst case)

# Transactions

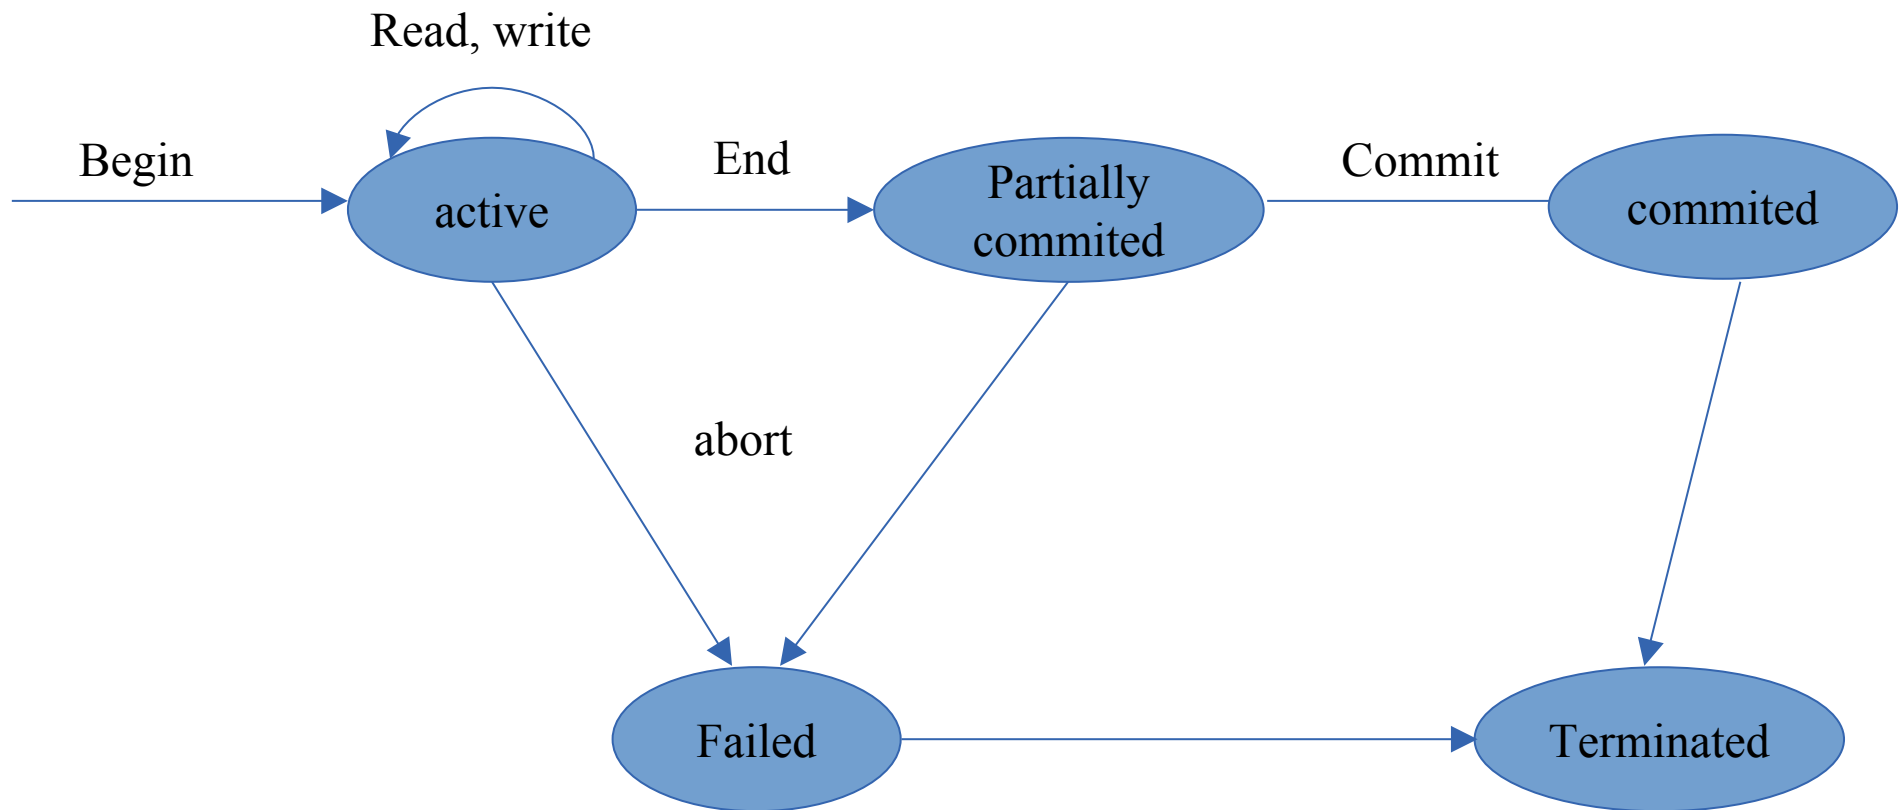
# Transactions

**A transaction is an atomic unit of work that is either completed in its entirety or not done at all**  
[Navathe and Elmasri, 2005]

T1	Database
r(balance)	balance = 200
balance += 100	r(x)
w(balance)	balance = 300

# Transaction State Transition

[Elmasri and Navathe, 2005]



# ACID Properties

**ACID = Atomicity, Consistency, Isolation, Durability**

**Atomicity:** if one part of the transaction fails, the entire transaction fails, and the database state is left unchanged

**Consistency:** ensures that any transaction will bring the database from one valid state to another (constraint rules)

**Isolation:** the concurrent execution of transactions results in a state that would be obtained if they were executed serially

**Durability:** once a transaction has been committed, it will remain so, even in the event of power loss, crashes, or errors

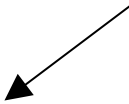
# Schedules

**A schedule orders the execution of concurrent transactions in an interleaving fashion**

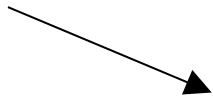
T1	T2
r(x)	
	r(x)
	w(x)
w(x)	

# Concurrency Conflicts

**Concurrency problems:**  
dirty read and lost update



T1	T2
r(balance)	
balance += 100	r(x)
w(balance)	w(x)
w(x)	r(balance)
abort	balance-=100
	w(balance)



T1	T2
r(balance)	
	r(balance)
balance += 100	w(x)
w(x)	balance -= 100
w(balance)	w(balance)

# Testing Conflict Serializability

## Algorithm:

Create a node for each transaction T in schedule S;

Create an edge  $T_i \rightarrow T_j$  for each  $r(x)$  in  $T_j$  after  $w(x)$  in  $T_i$


Create an edge  $T_i \rightarrow T_j$  for each  $w(x)$  in  $T_j$  after  $r(x)$  in  $T_i$

Create an edge  $T_i \rightarrow T_j$  for each  $w(x)$  in  $T_j$  after  $w(x)$  in  $T_i$

S is serializable if and only if the precedence graph has no cycles

## For example:

T1	T2
r(x)	
	r(x)
	w(x)
w(x)	



Not serializable



# Readings

- Joseph Hellerstein, Michael Stonebraker and James Hamilton.
- Architecture of a Database System, material at:  
<http://db.cs.berkeley.edu/papers/fntdb07-architecture.pdf>
- Hector Garcia-Molina, Jeffrey Ullman, Jennifer Widom.  
Database Systems: The Complete Book, material at:  
<http://infolab.stanford.edu/~ullman/dscb.html>

**Thank You !**  
**15 minutes break**

# Quiz

1. What types of page layout are more likely for the following queries? Why?

- 1.Q1    1.SELECT \*  
         2.FROM CUSTOMERS;
- 1.Q2    1.SELECT A.ID, A.BALANCE, C.NAME  
         2.FROM CUSTOMERS C, ACCOUNT A  
         3.WHERE A.ID=C.ID
- 1.Q3    1.SELECT SUM(BALANCE)  
         2.FROM ACCOUNT;  
         3.  
            SELECT C.CITY, SUM(A.BALANCE)  
1.Q4    FROM CUSTOMERS C, ACCOUNT A  
         WHERE A.ID=C.ID  
         GROUP BY C.CITY;

BigData

(T.F. Bissyandé & M. Hurier)

# Quiz

1.Considering the following queries, what types of index need to be implemented? Why?

1.Q1      1.SELECT NAME, AGE  
            2.FROM CUSTOMER  
            3.WHERE AGE BETWEEN 18 AND 25;

1.Q2      1.SELECT NAME, AGE  
            2.FROM CUSTOMER  
            3.WHERE CUST\_ID= 1234;

1.Q3      1.SELECT NAME, CITY, AGE, SSN  
            2.FROM CUSTOMER  
            3.WHERE AGE BETWEEN 18 AND 25  
            4.AND CITY ;

# Quiz

1. Please, rewrite the following expressions using the presented rules and the following relations:

1.Q1

ACCOUNT (BALANCE > 2500  
B\_ID),  
CUSTOMER (C\_ID > 1000  
BRANCH (B\_ID > 1000  
B\_NAME))  
ACCOUNT X

1.Q2

CUSTOMER (BALANCE > 2500  
( $\sigma$  B\_ID = 1000  
(ACCOUNT X  
BALANCE, C\_ID  
BRANCH))  
( $\sigma$  BALANCE >

1.Q3

2500 (ACCOUNT X  
BALANCE)  
CUSTOMER (C\_ID > 1000  
BRANCH (B\_ID > 1000  
B\_NAME))  
ACC\_ID

1.Q4

(CUSTOMER X  
BALANCE,  
ACCOUNT)  
ACC\_ID

1.Q5

(CUSTOMER X  
(ACCOUNT X  
BRANCH))

# Quiz

1. Please, rewrite the following expressions using the presented rules and the following relations:

- 1.Q1  $\pi_{B\_ID}(\sigma_{C\_ID > 1000}(\sigma_{BALANCE > 2500}(\pi_{C\_ID, BALANCE}(\text{CUSTOMER} \times \text{ACCOUNT}))))$
- 1.Q2  $\pi_{C\_ID}(\sigma_{BALANCE > 2500}(\pi_{C\_ID, BALANCE}(\text{CUSTOMER} \times \text{ACCOUNT})))$
- 1.Q3  $\pi_{C\_ID}(\pi_{BALANCE}(\pi_{C\_ID, BALANCE}(\text{CUSTOMER} \times \text{ACCOUNT})))$
- 1.Q4  $\pi_{C\_ID}(\pi_{BALANCE}(\pi_{C\_ID, BALANCE}(\text{CUSTOMER} \times \text{ACCOUNT})))$
- 1.Q5  $\pi_{C\_ID}(\pi_{BALANCE}(\pi_{C\_ID, BALANCE}(\text{CUSTOMER} \times \text{ACCOUNT})))$

# Quiz

1. Compute of the cost for a query selecting 10% of a table 'R' based on the following information from the catalog:

Relation 'R'	10,000 tuples
Tuples/page ratio	100 tuples

Scan	
Scan on sorted	
Clustered index scan	
Non-clustered index scan	

# Quiz

1. Compute of the I/O cost for a query selecting 10% of a table 'R' based on the following information from the catalog:

Relation 'R'	10,000 tuples
Tuples/page ratio	100 tuples

Scan	100 I/O
Scan on sorted	6.6 I/O
Clustered index scan	3 I/O
Non-clustered index scan	103 I/O



# Hands-on

Which of the following schedules is conflict serializable?

S1 - T1 : r(x), T3 : r(x), T1 : w(x), T2 : r(x), T3 : w(x)

S2 - T1 : r(x), T2 : r(x), T1 : w(x), T2 : w(x)

S3 - T1 : r(x), T2 : r(y), T3 : w(x), T2 : r(x), T1 : r(x)