

NoSQL & NewSQL Databases

Responsible:

Dr. Tegawendé F. BISSYANDE

tegawende.bissyande@uni.lu

Course Author:

Dongsun Kim

dongsun.kim@uni.lu

Teacher Assistant:

Médéric Hurier

mederic.hurier@uni.lu

Course Features

- Sep. 20th - **Introduction to Big Data**

Part 1. Databases and Query Models for Big Data

- Sep. 27th - **Relational Databases: Reminders**
- Oct. 4th - **Relational Databases: Internals**
- Oct. 11th - **NoSQL & NewSQL Databases**
- Oct. 18th - **MapReduce Model**
- Oct. 25th - **Hadoop and Spark**
- Nov. 8th - **Datalog Model**

Part 2. Data Analysis and Machine Learning

- Nov. 15th - **Statistics and Probabilities**
- Nov. 22th - **Communication and Visualization**
- Nov. 29th - **Features Engineering and Supervised Learning**
- Dec. 5st - **Unsupervised and Reinforcement Learning**
- Dec. 12th - **Homework Time**

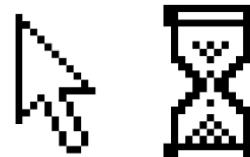
Section Features

- Motivation
- CAP Theorem
- Characteristics
- Database Landscape
- Example: MongoDB

Motivation

In a time far away ...

Everybody was happy about RDBMS/SQL !



- Centralized
- Easy Maintenance
- Available anytime
- Stable performance
- Quite reliable (ACID)



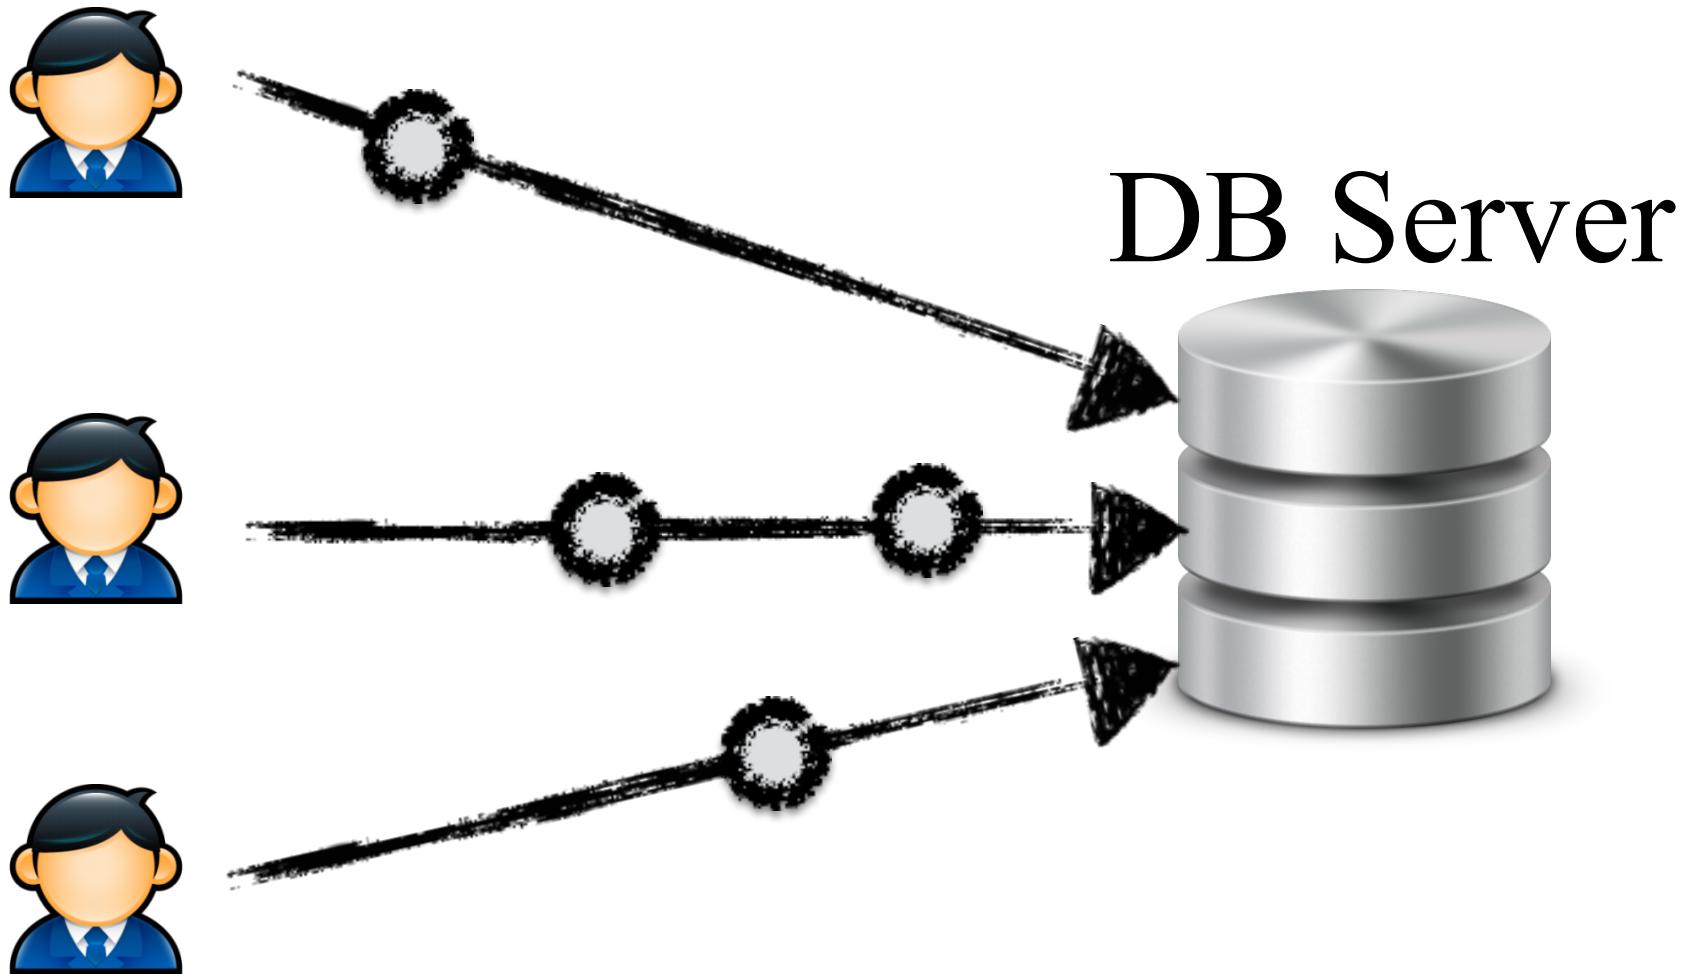
ORACLE



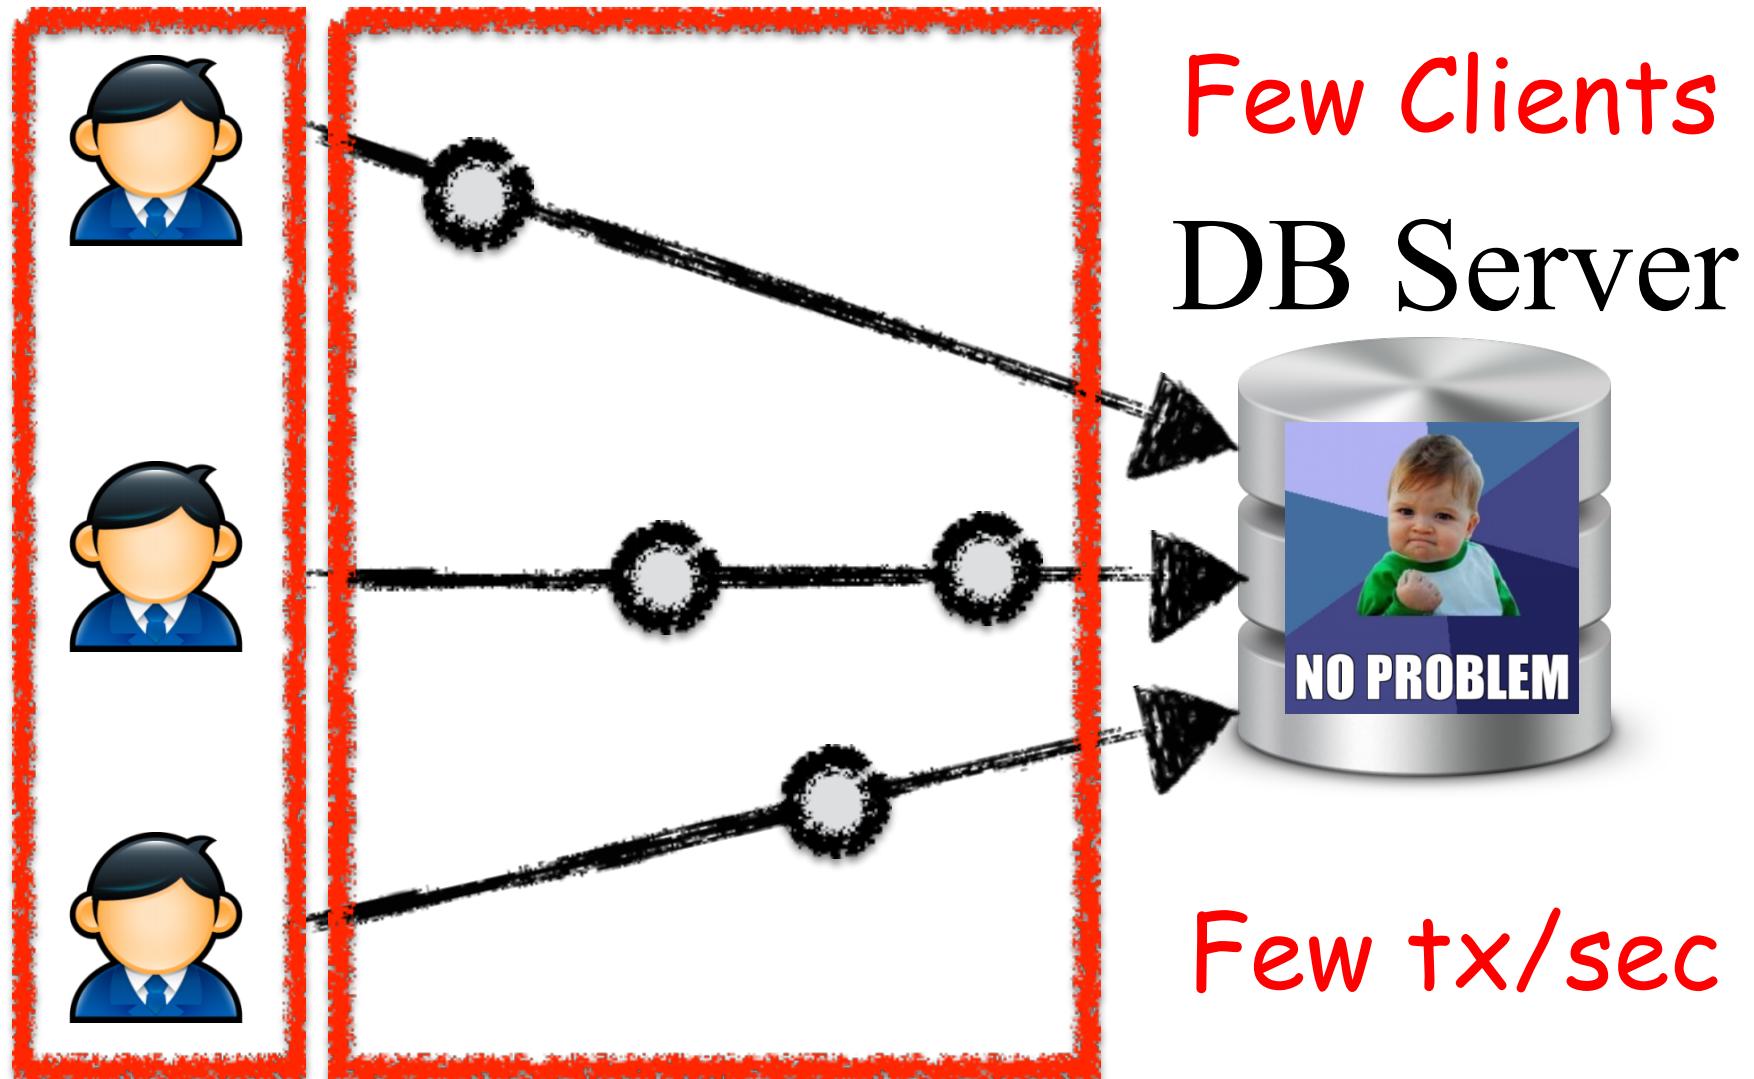
That ... was enough ?



Traditional RDBMS: Past



Traditional RDBMS: Past



Traditional RDBMS: Present



RDBMS is now overloaded !



- Large number of sources: users, devices, sensors
- Swarm of transactions: e.g. 1,000tx/sec/source
- Huge data size: e.g. 10MB/sec stream

RDBMS is not agile enough !

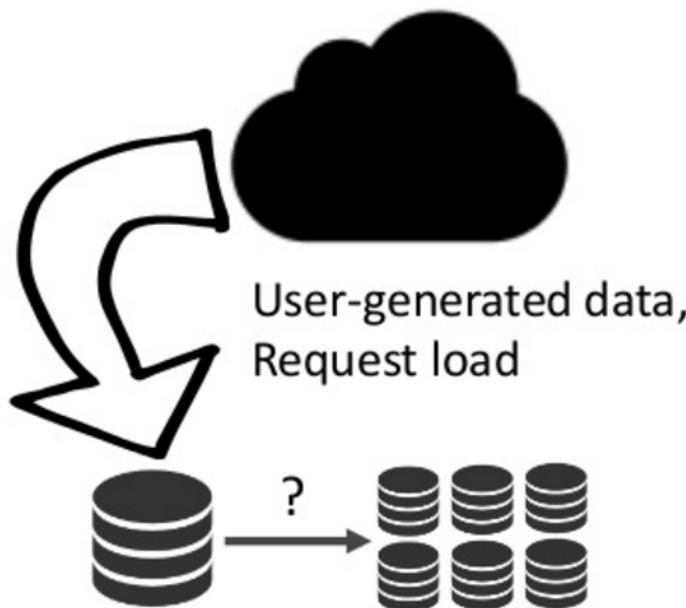
Facebook added a new feature: replying with a photo
(June 2013)



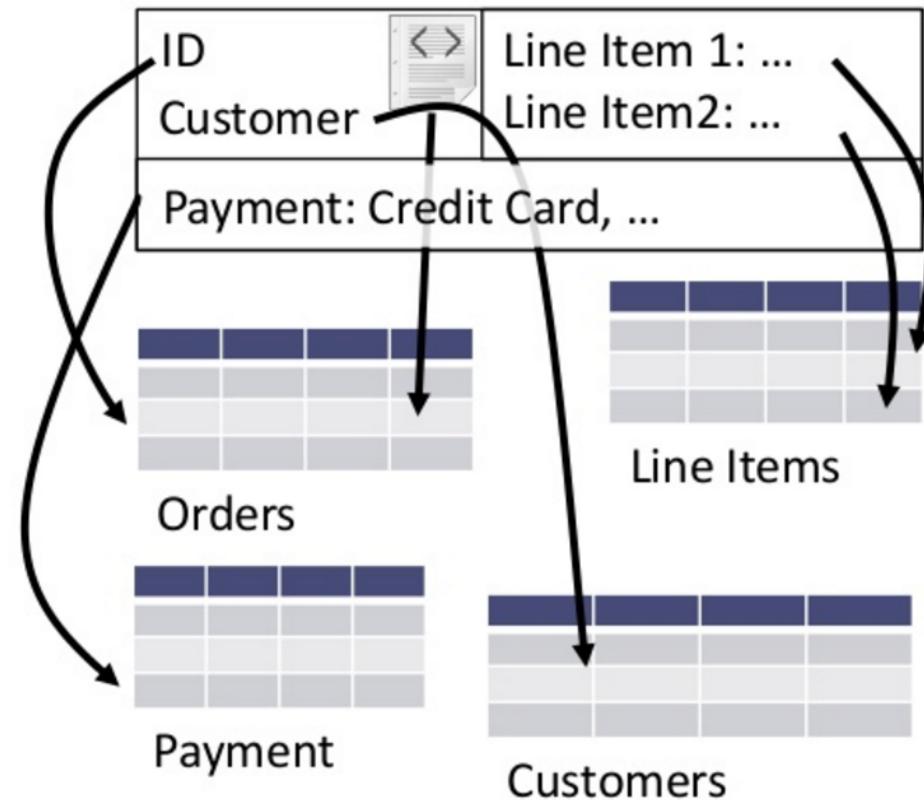
source: <http://www.technixupdate.com/how-to-use-new-facebook-photo-reply-comment-feature/>

Main motivations for NoSQL

Scalability



Impedance Mismatch



Characteristics

Key features of NewSQL

- Conform ACID
- High-availability
- Support SQL naturally
- In-Memory Transactions
- Scalable (e.g. horizontal scaling)
- Easy replication through network



Key features of NoSQL

Term coined in 2009 (Not Only SQL)



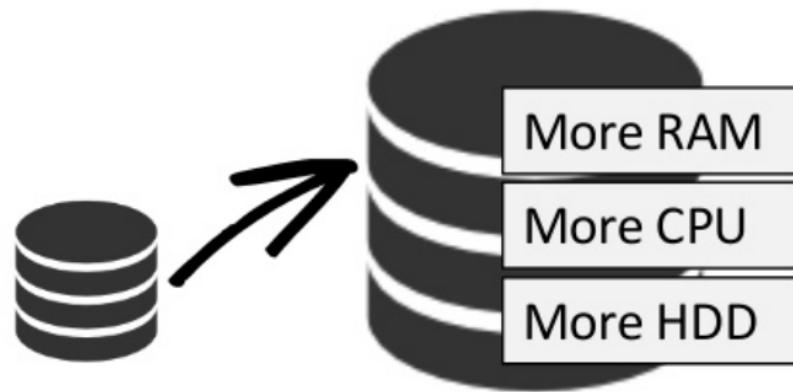
- Scalable (e.g. horizontal scaling)
- Easy replication through the network
- “Simple” API language (NO query language)
- Tunable consistency model (NO ACID properties)
- Ability to add new attributes to record (NO schema)

OldSQL vs. NewSQL vs. NoSQL

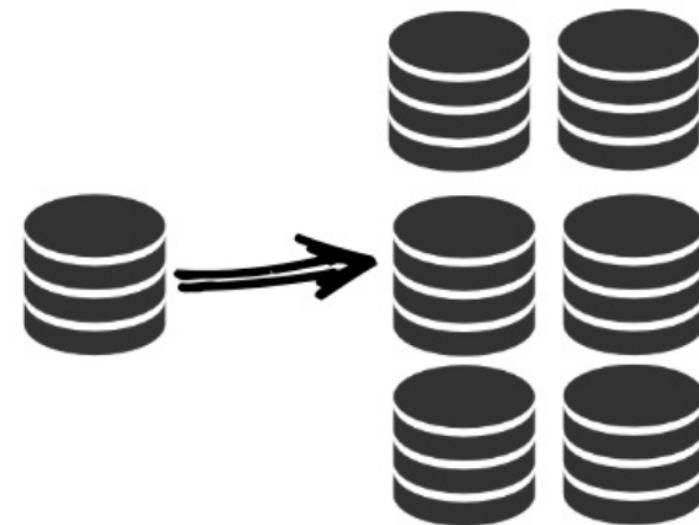
	OldSQL	NewSQL	NoSQL
ACID/Transaction	✓	✓	✗
Language/Algebra	✓	✓	✗
Schema/Constraints	✓	✓	✗
Scalability	✗	✓	✓✓✓
In-memory Tx	✗	✓	✓
Easy replication	✗	✓	✓

Vertical vs. Horizontal Scaling

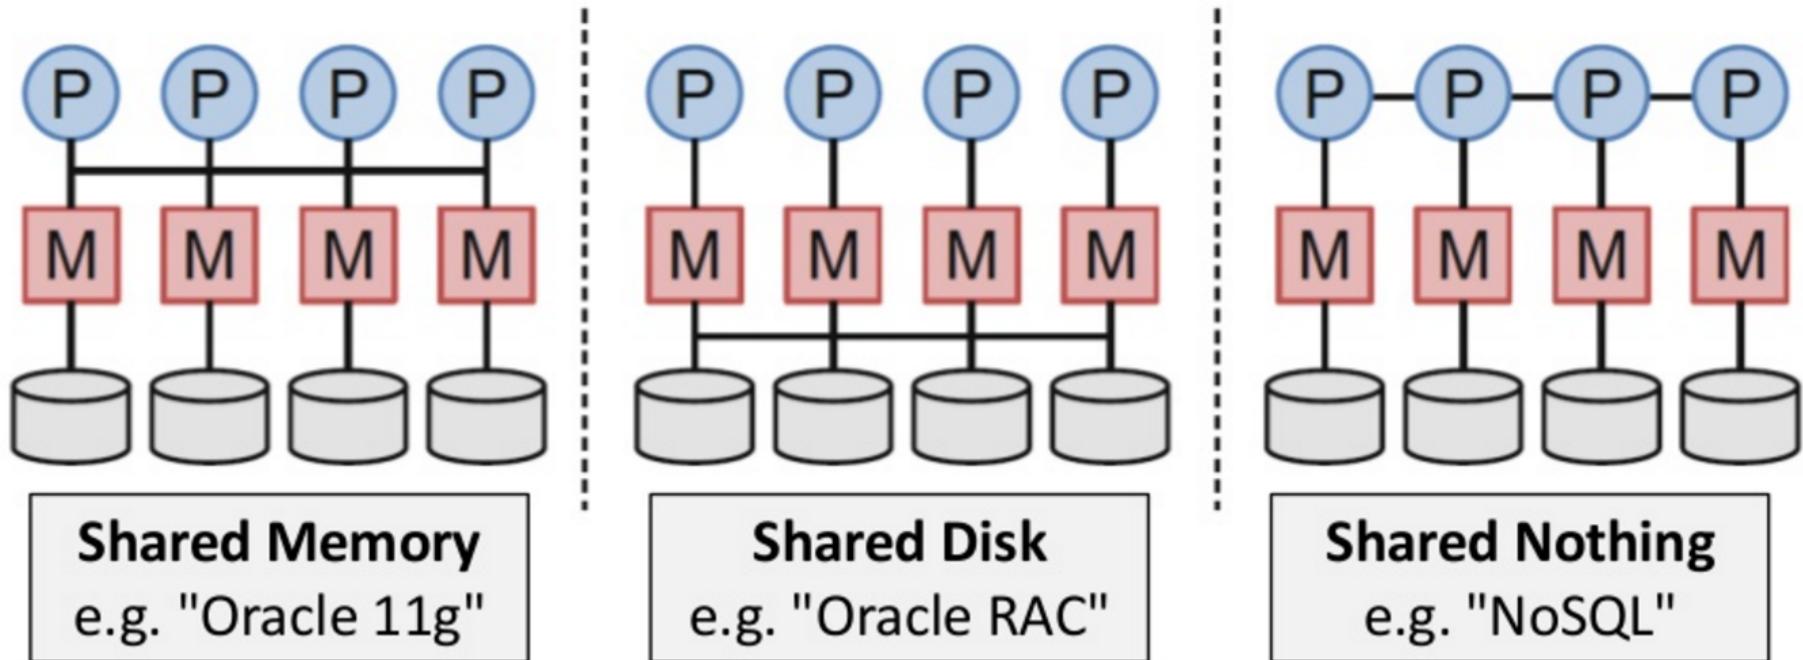
Scale-Up (*vertical* scaling):



Scale-Out (*horizontal* scaling):



Shift toward “Shared Nothing”



When each node is independent and self-sufficient, there is no single point of contention across the system

the potential for scaling is huge !

What if it is schema-dependent ?

Post Table



Reply Table



Do we really need to add a table for it?

Before

After

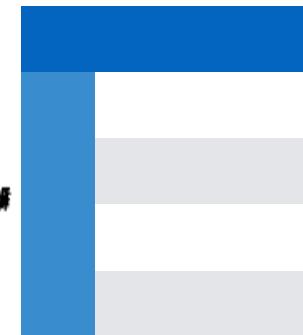
Post Table



Reply Table

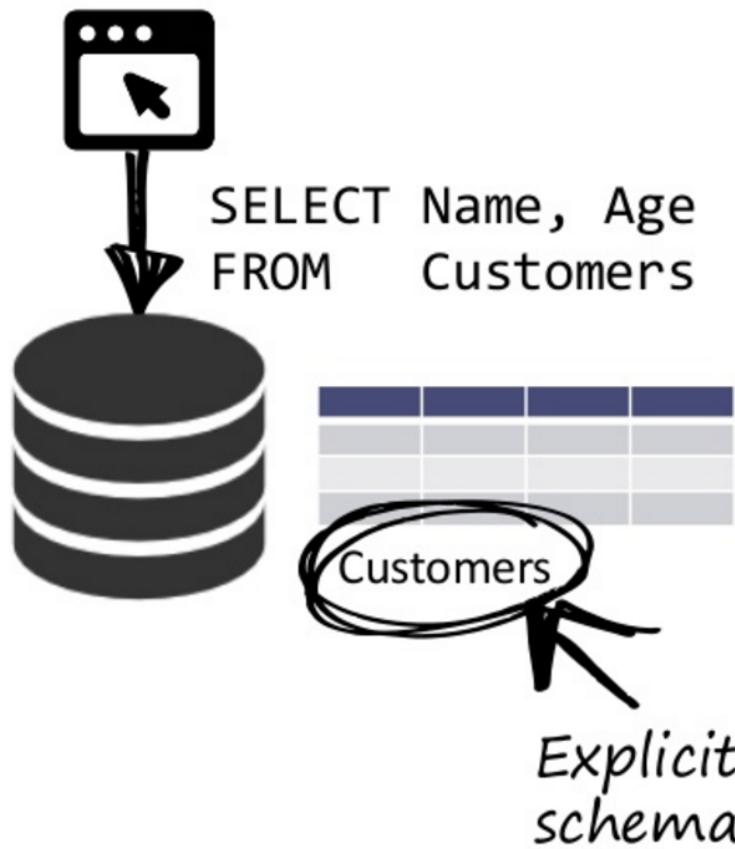


Photo Table

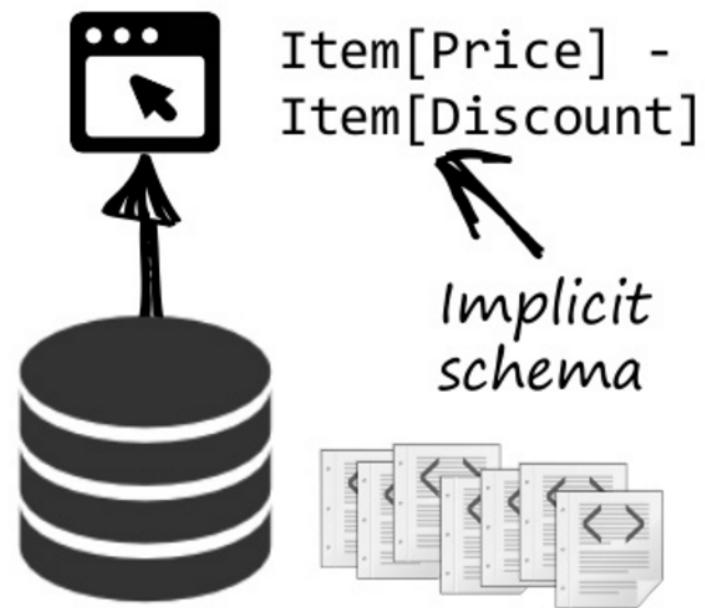


What if it is schema-free ?

RDBMS:



NoSQL DB:



Benefits of schema-free models

Flexibilit
y



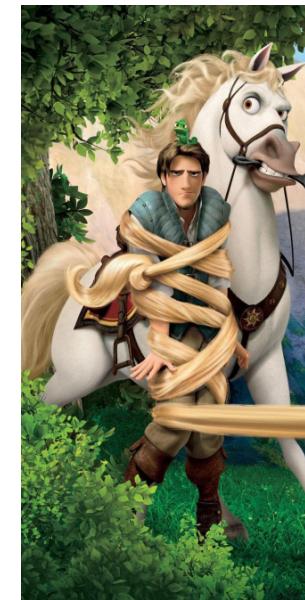
Easy to replicate



Easy to partition

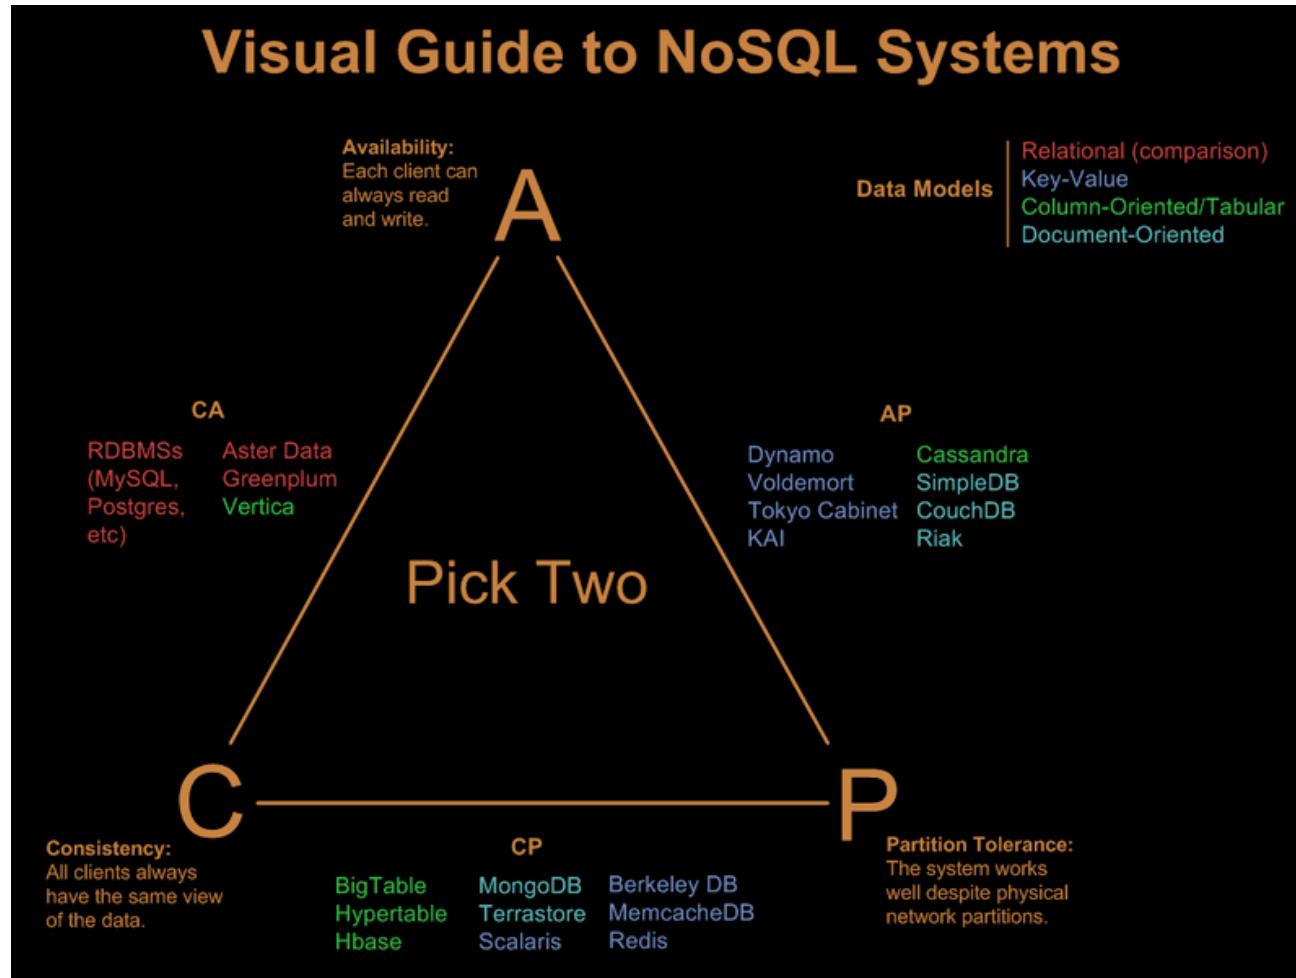


Join-free



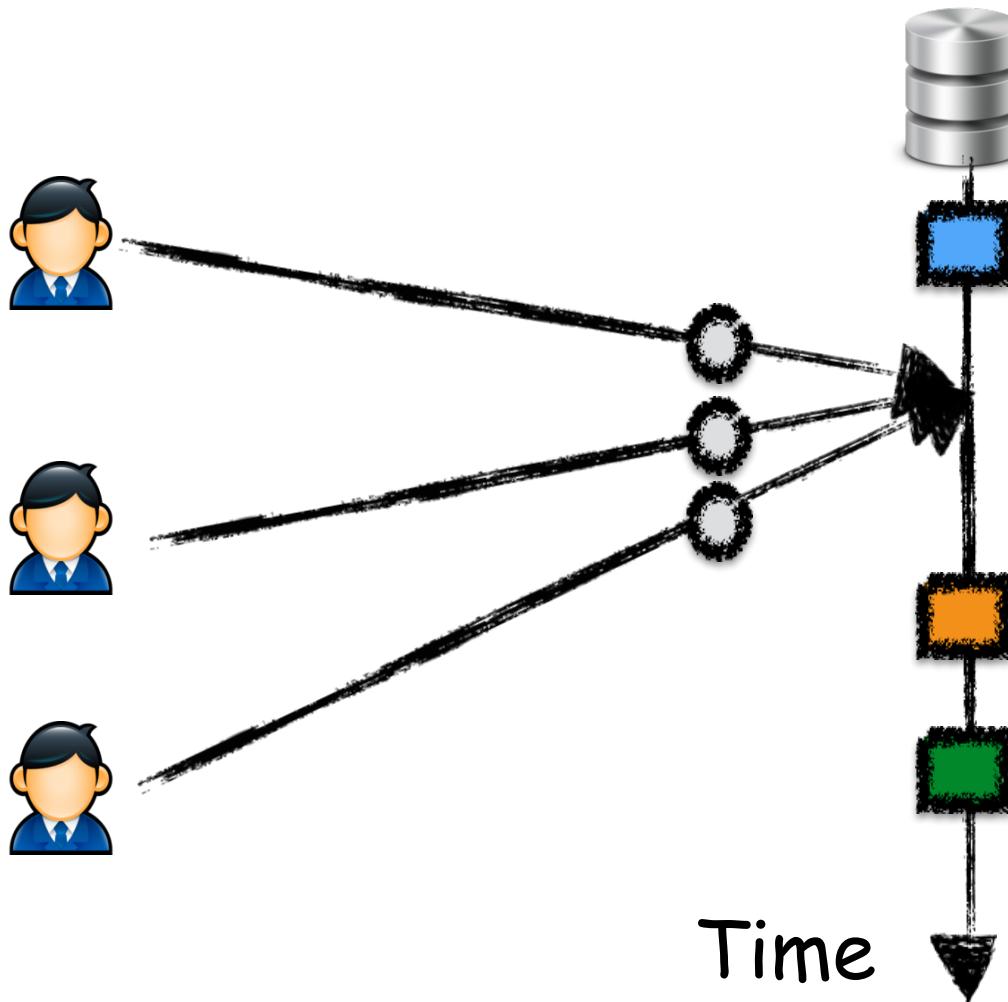
CAP Theorem

CAP Theorem



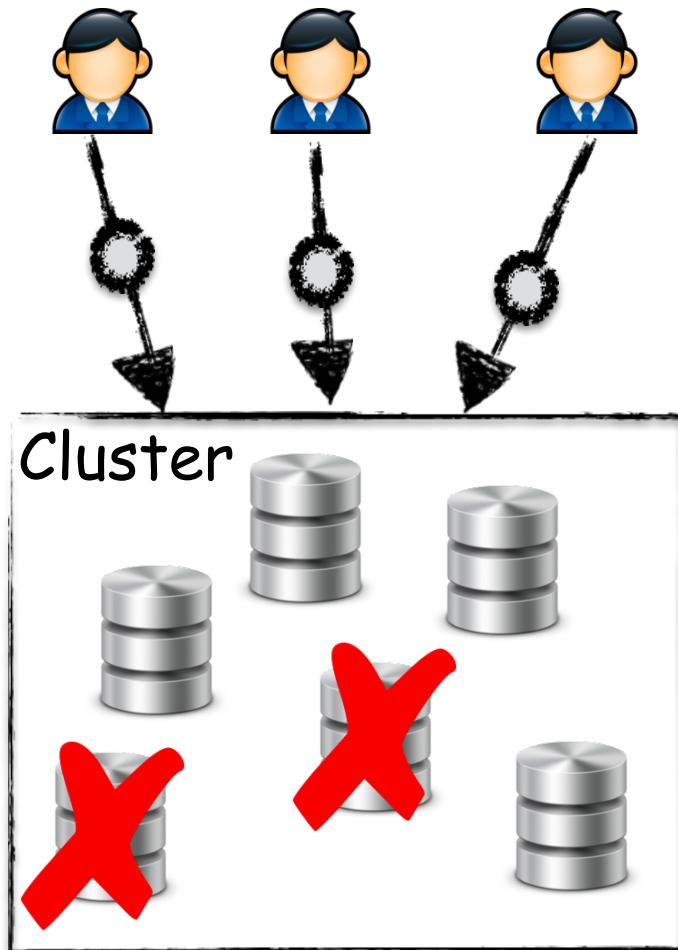
Source: <http://blog.nahurst.com/visual-guide-to-nosql-systems>

Consistency



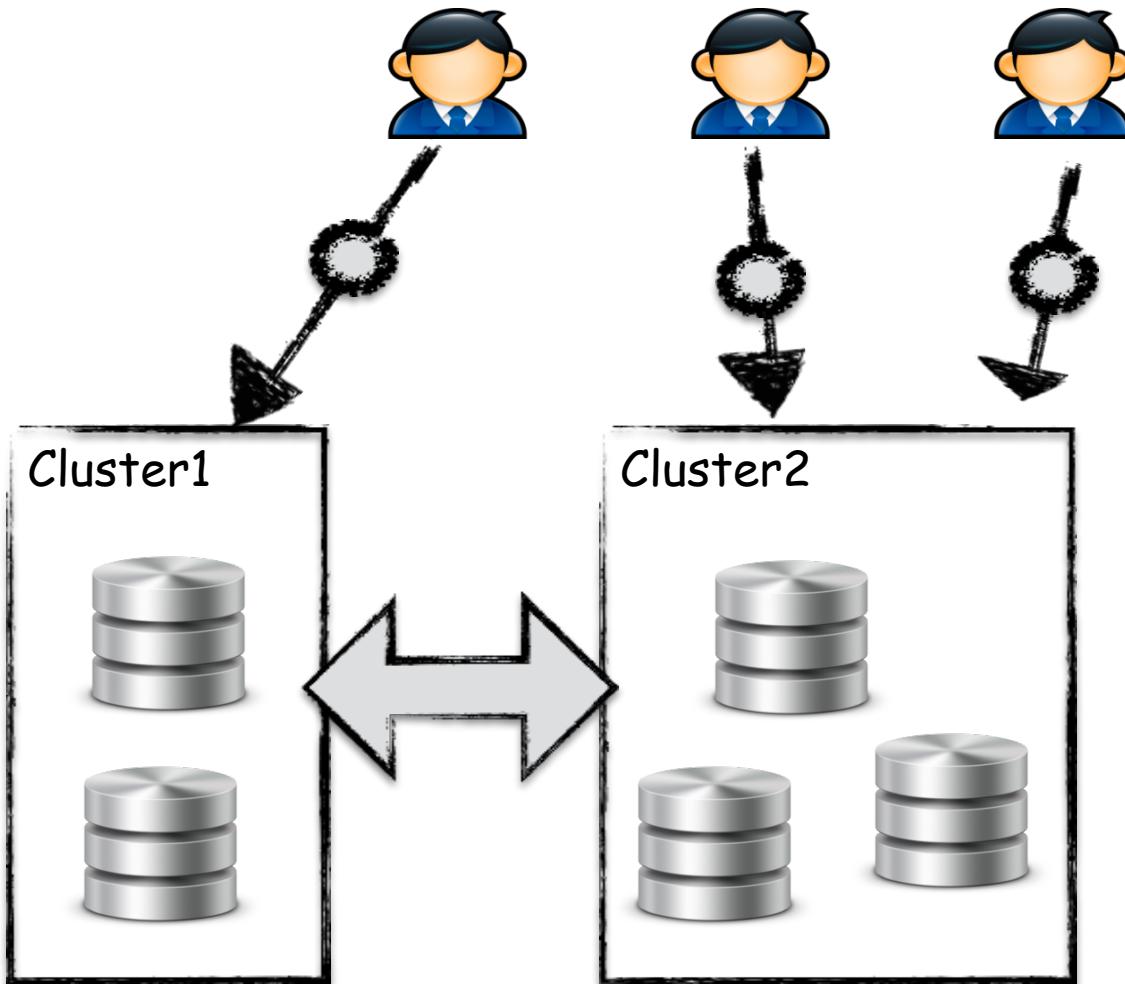
All clients have the same view on the data

Availability



Every request to a non-failed node must result in a correct response

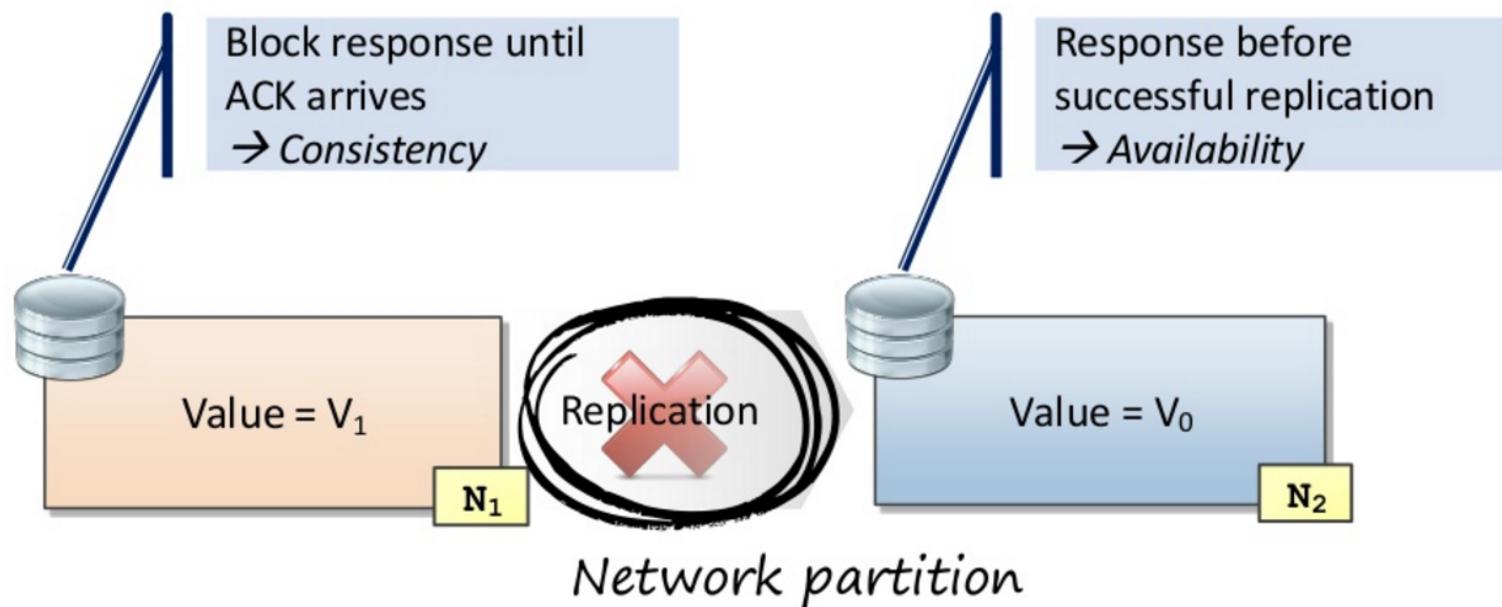
Partition tolerance



The system has to continue working, even under arbitrary network partitions

Simplified proof

**When a network partition occurs,
either consistency or availability has to give up**



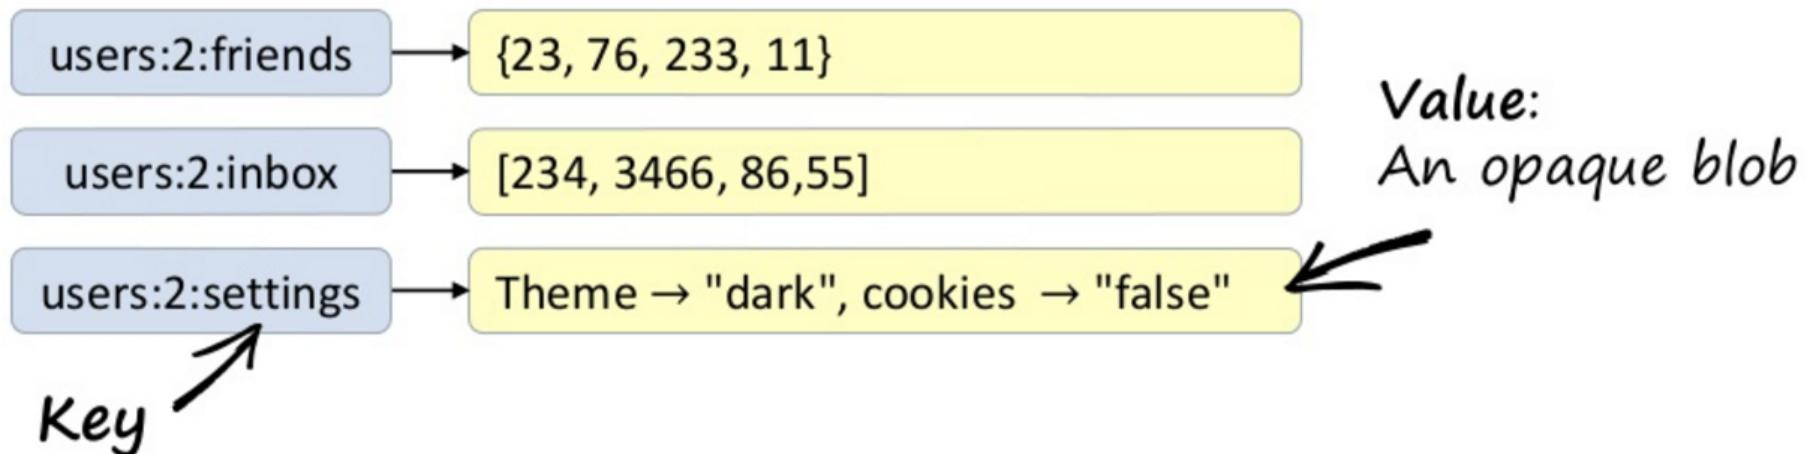
Database Landscape

NoSQL Landscape



Data Model: Key-Value

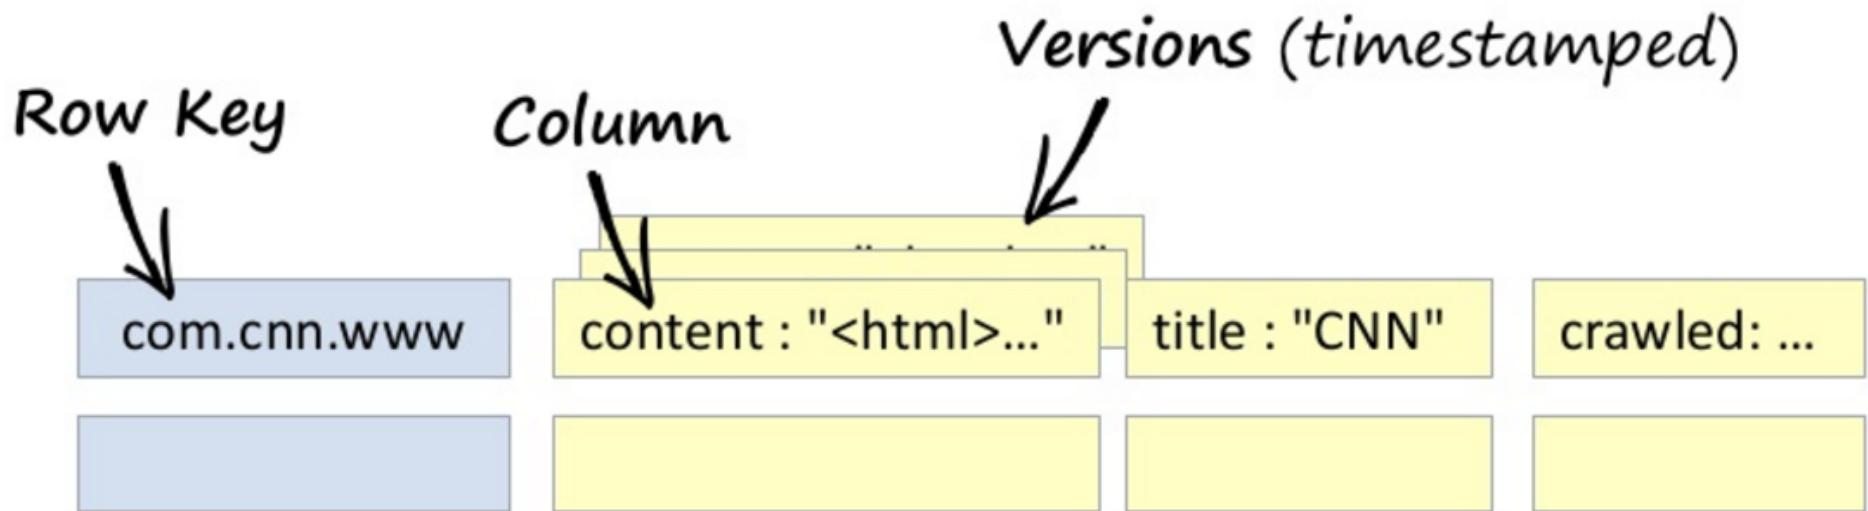
single (key) -> value, no foreign key/value indices
appropriate for instant access to data (shop basket)



Example: Dynamo (AP), Riak (AP), Redis (CP)

Data Model: Wide-Column

data indexed by a 3D map: (row, col, time) -> val
appropriate for aggregation and analytics (reports)

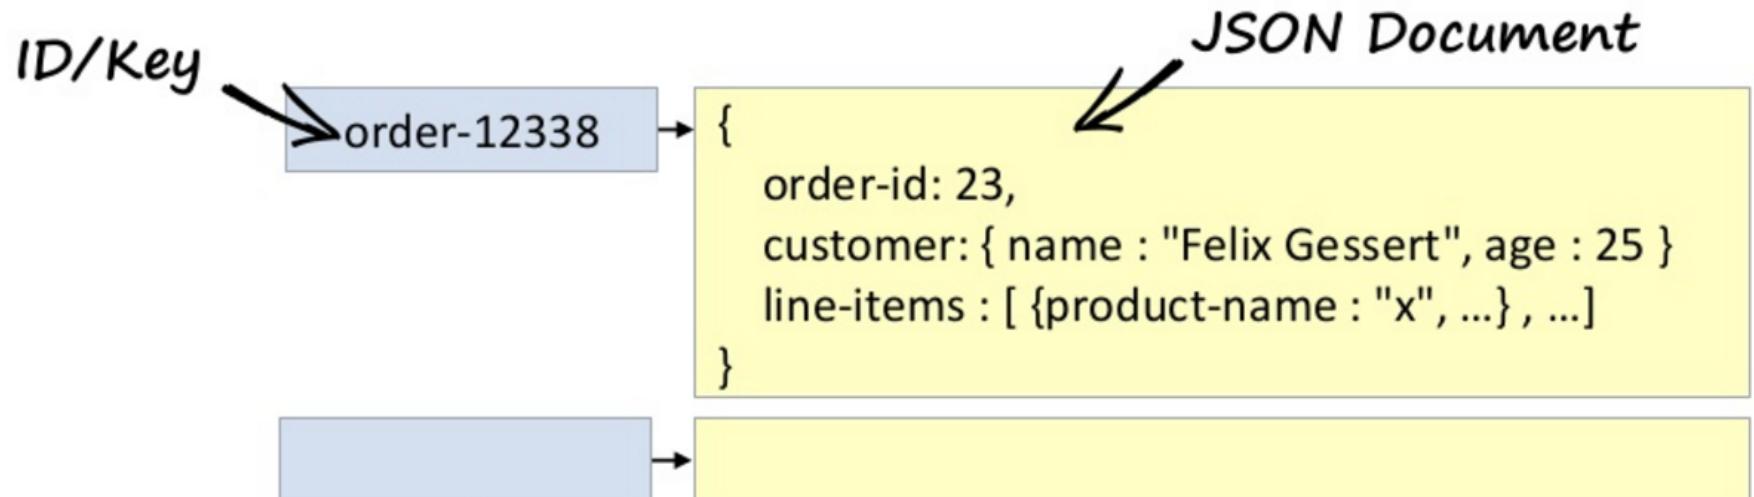


Example: Cassandra (AP), BigTable (CP), HBase (CP)

Data Model: Document

similar to JSON/XML: (coll, key) -> doc

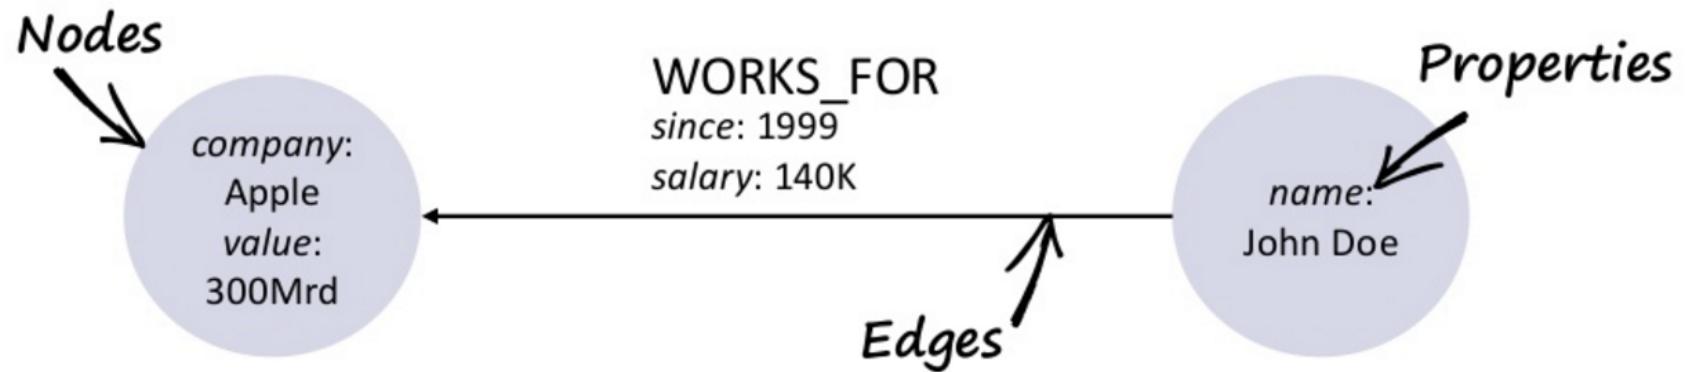
appropriate for data not frequently changing (logs)



Example: CouchDB (AP), SimpleDB (AP), MongoDB (CP)

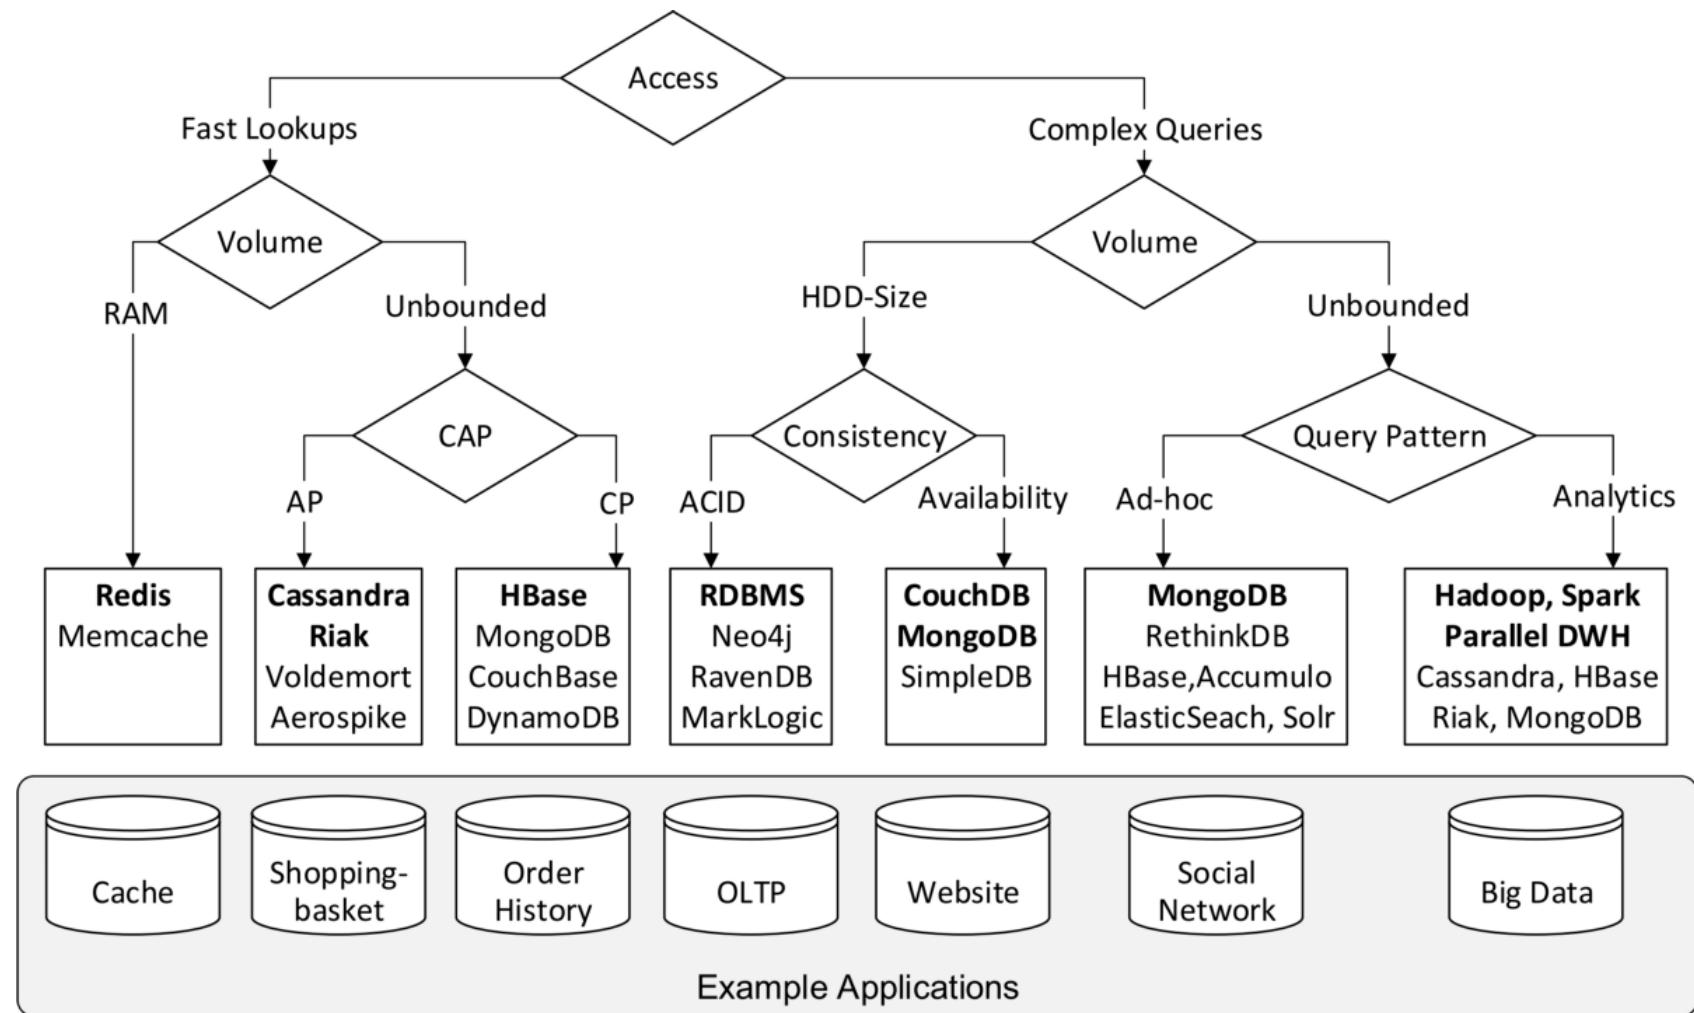
Data Model: Graph

data represented by vertices/edges: $G = (V, E)$
appropriate for complex structures (social network)



Example: Neo4j (CA), InfiniteGraph (CA), OrientDB (CA)

Decision tree for mapping requirements



Source: <https://medium.baqend.com/nosql-databases-a-survey-and-decision-guidance-ea7823a822d>

Database Ranking

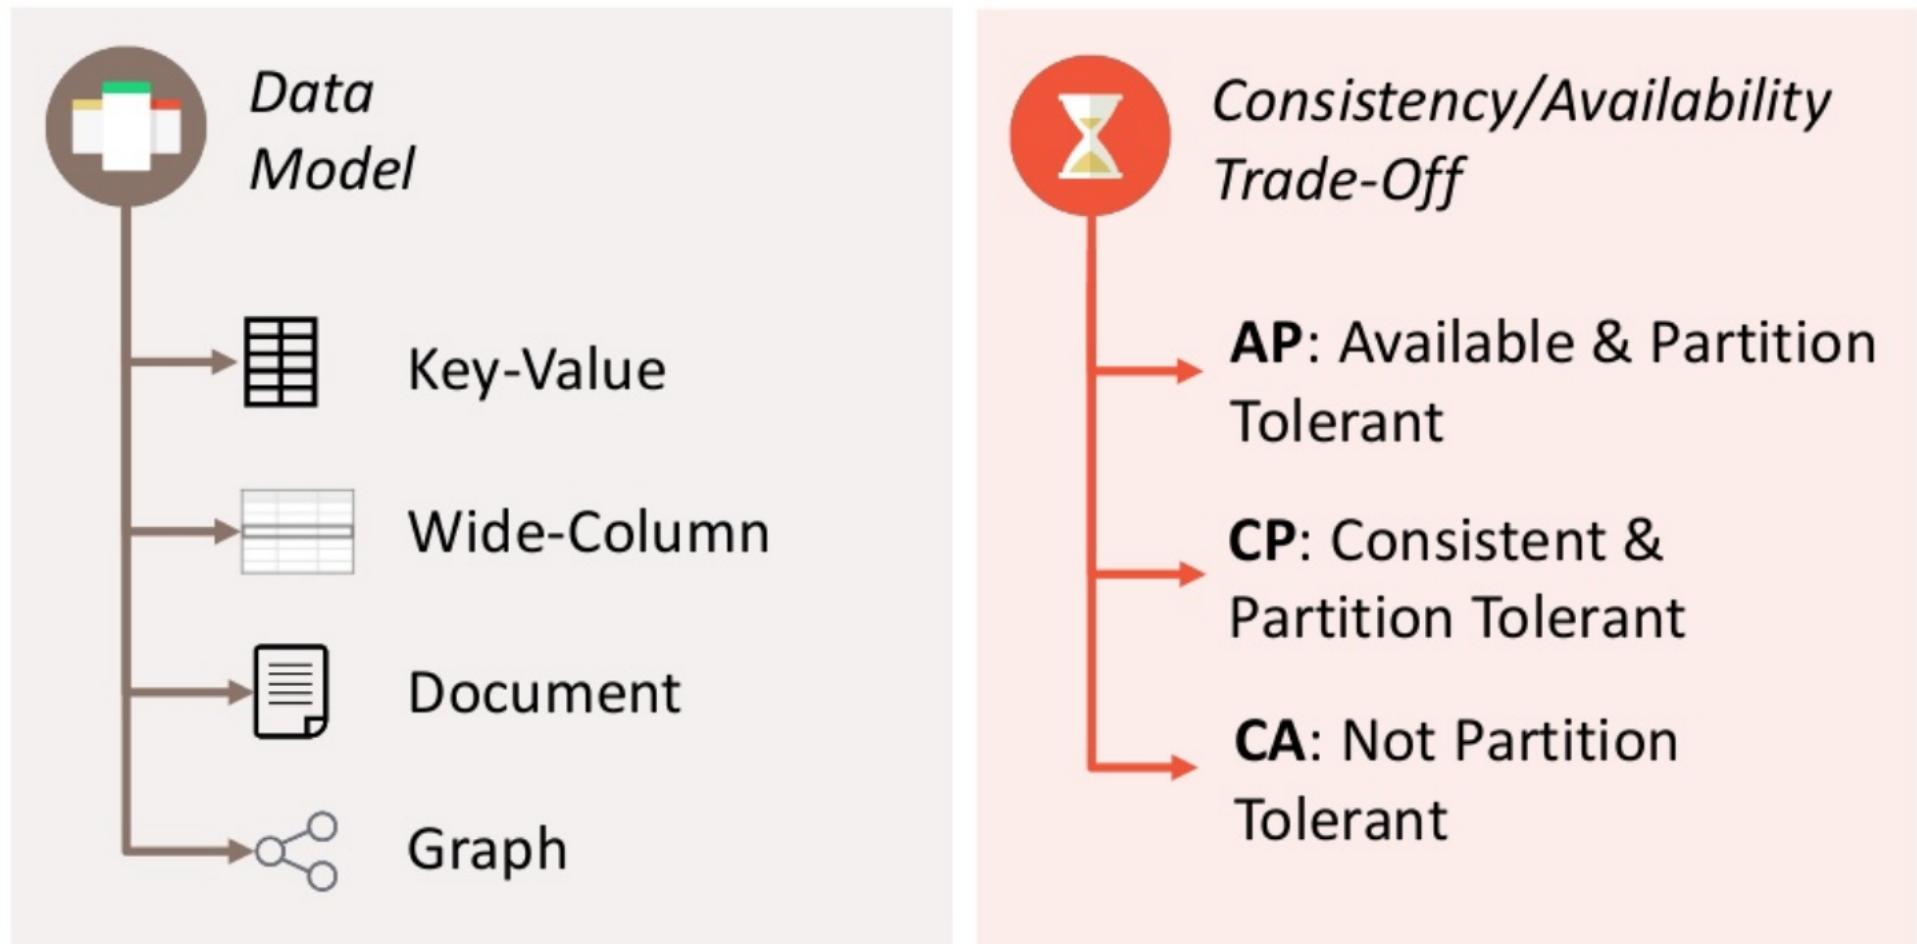
<http://db-engines.com/en/ranking>

316 systems in ranking, October 2016

Rank			DBMS	Database Model	Score		
Oct 2016	Sep 2016	Oct 2015			Oct 2016	Sep 2016	Oct 2015
1.	1.	1.	Oracle 	Relational DBMS	1417.10	-8.46	-49.85
2.	2.	2.	MySQL 	Relational DBMS	1362.65	+8.62	+83.69
3.	3.	3.	Microsoft SQL Server	Relational DBMS	1214.18	+2.62	+90.95
4.	 5.	4.	MongoDB 	Document store	318.80	+2.81	+25.54
5.	 4.	5.	PostgreSQL	Relational DBMS	318.69	+2.34	+36.56
6.	6.	6.	DB2	Relational DBMS	180.56	-0.62	-26.25
7.	7.	 8.	Cassandra 	Wide column store	135.06	+4.57	+6.05
8.	8.	 7.	Microsoft Access	Relational DBMS	124.68	+1.36	-17.16
9.	 10.	 10.	Redis	Key-value store	109.54	+1.75	+10.75
10.	 9.	 9.	SQLite	Relational DBMS	108.57	-0.05	+5.90

Summary

The era of one-size-fits-all database systems is over !



Example: MongoDB

MongoDB in a nutshell

- Created by MongoDB Inc. in Oct. 2007
 - MongoDB <= huMONGous DB
- Document-oriented data model
 - BSON (similar to JSON)
- Supports:
 - indexing
 - replication
 - load balancing
 - server-side



mongoDB

JavaScript

execution.

Insert/Update/Delete

- **Insert:** db.<collection>.insert(<document>)
 - db.users.insert({name:"john",city:"Luxembourg"})
 - no need to create a table upfront !
- **Update:** db.<collection>.update(<query>, <values>)
 - db.users.update({name:"john"}, {\$set: {city:"Seoul"}})
- **Delete:** db.<collection>.remove(<query>)
 - db.users.remove({name:"john"})

Select/Queries

- **Simple Select:** db.<collection>.find(<query>)

- db.users.find({name:"john"})

- **Nested Select:** use dot-notation

- db.users.find({"score.math":30})

- **Show specific columns only**

- db.users.find({}, {"score.math":1, "score.lang":1})

- **Query if to test if exists**

- db.users.find({"score.math":{\$exists:true}}), {name:1})

Other useful operations

- **count()**: counting the number of documents
 - db.users.find().count()
 - db.users.count({name:"john"})
- **sort()**: sorting documents
 - db.users.find().sort({name:-1})
- **List** collections (tables) and databases
 - show collections/tables ; show databases
- **limit()**: Show a specified number of documents
 - db.users.find().limit(5)
(T.F. Bissyandé & M. Hurier)