

MapReduce Model

Responsible:

Dr. Tegawendé F. BISSYANDE

tegawende.bissyande@uni.lu

Course Author:

Edson Ramiro Lucas Filho

Eduardo Cunha de Almeida

Teacher Assistant:

Médéric Hurier

mederic.hurier@uni.lu

Course Features

- Sep. 20th - **Introduction to Big Data**

Part 1. Databases and Query Models for Big Data

- Sep. 27th - **Relational Databases: Reminders**
- Oct. 4th - **Relational Databases: Internals**
- Oct. 11th - **NoSQL & NewSQL Databases**
- **Oct. 18th - MapReduce Model**
- Oct. 25th - **Hadoop and Spark**
- Nov. 8th - **Datalog Model**

Part 2. Data Analysis and Machine Learning

- Nov. 15th - **Statistics and Probabilities**
- Nov. 22th - **Communication and Visualization**
- Nov. 29th - **Features Engineering and Supervised Learning**
- Dec. 5st - **Unsupervised and Reinforcement Learning**
- Dec. 12th - **Homework Time**

What you should know after this class ?

- What is MapReduce and how it works.
- How we can scale with MapReduce
- How we can build a DBMS on top

Section Features

- Introduction
- MapReduce
- How does MapReduce Scale ?
- How do we build DBMSs on top of MapReduce ?

Introduction

What is scalability?

A system whose performance improves after adding hardware, proportionally to the capacity added, is said to be a scalable system.

Poor scalability can result in poor system performance, necessitating the reengineering or duplication of systems.

Database Systems for Advanced Applications: 16th International Conference, DASFAA 2011, Hong Kong, China, April 22-25, 2011, Proceedings. Yu, J.X. and Kim, M.H. and Unland, R.

<http://books.google.lu/books?id=zKHUaaT0Oq0C>

Characteristics of scalability and their impact on performance.

André B. Bondi. AT&T Labs

<http://dl.acm.org/citation.cfm?doid=350391.350432>

MapReduce

What is MapReduce ?

**It's a programming model
to simplify distributed data processing**

What is MapReduce ?

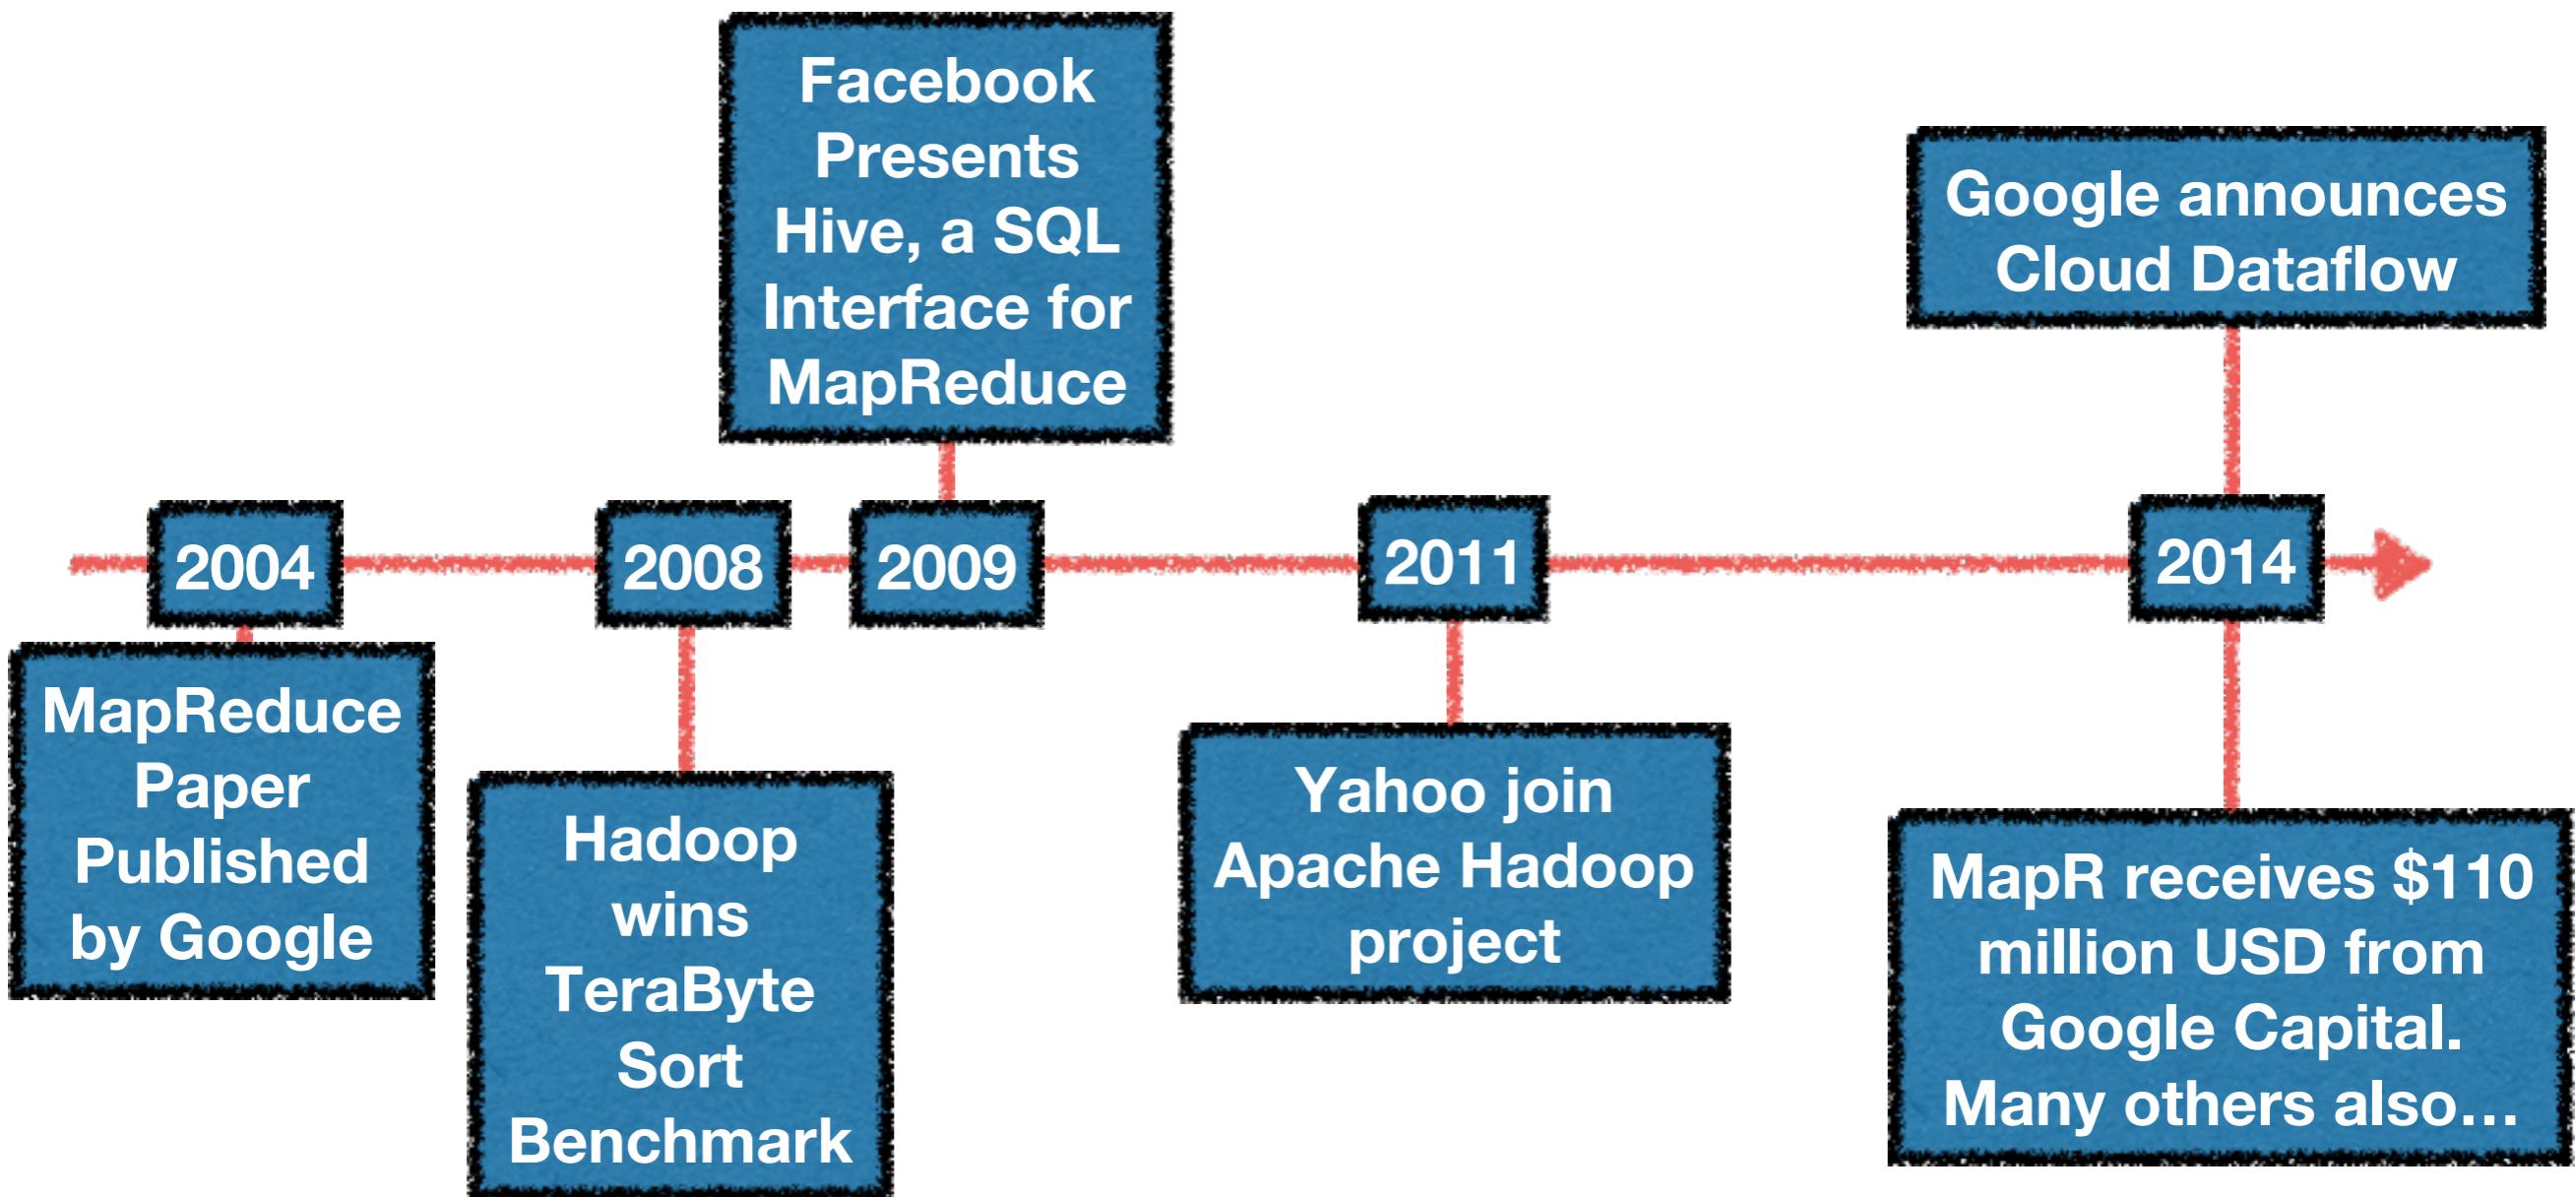
**Based on functional
programming paradigm**

**It's a programming model
to simplify distributed data processing**

**Hide from the user all
problems related to
distributed computation**

**Allows distributed
processing of large
data sets across clusters
of thousands of machines**

History



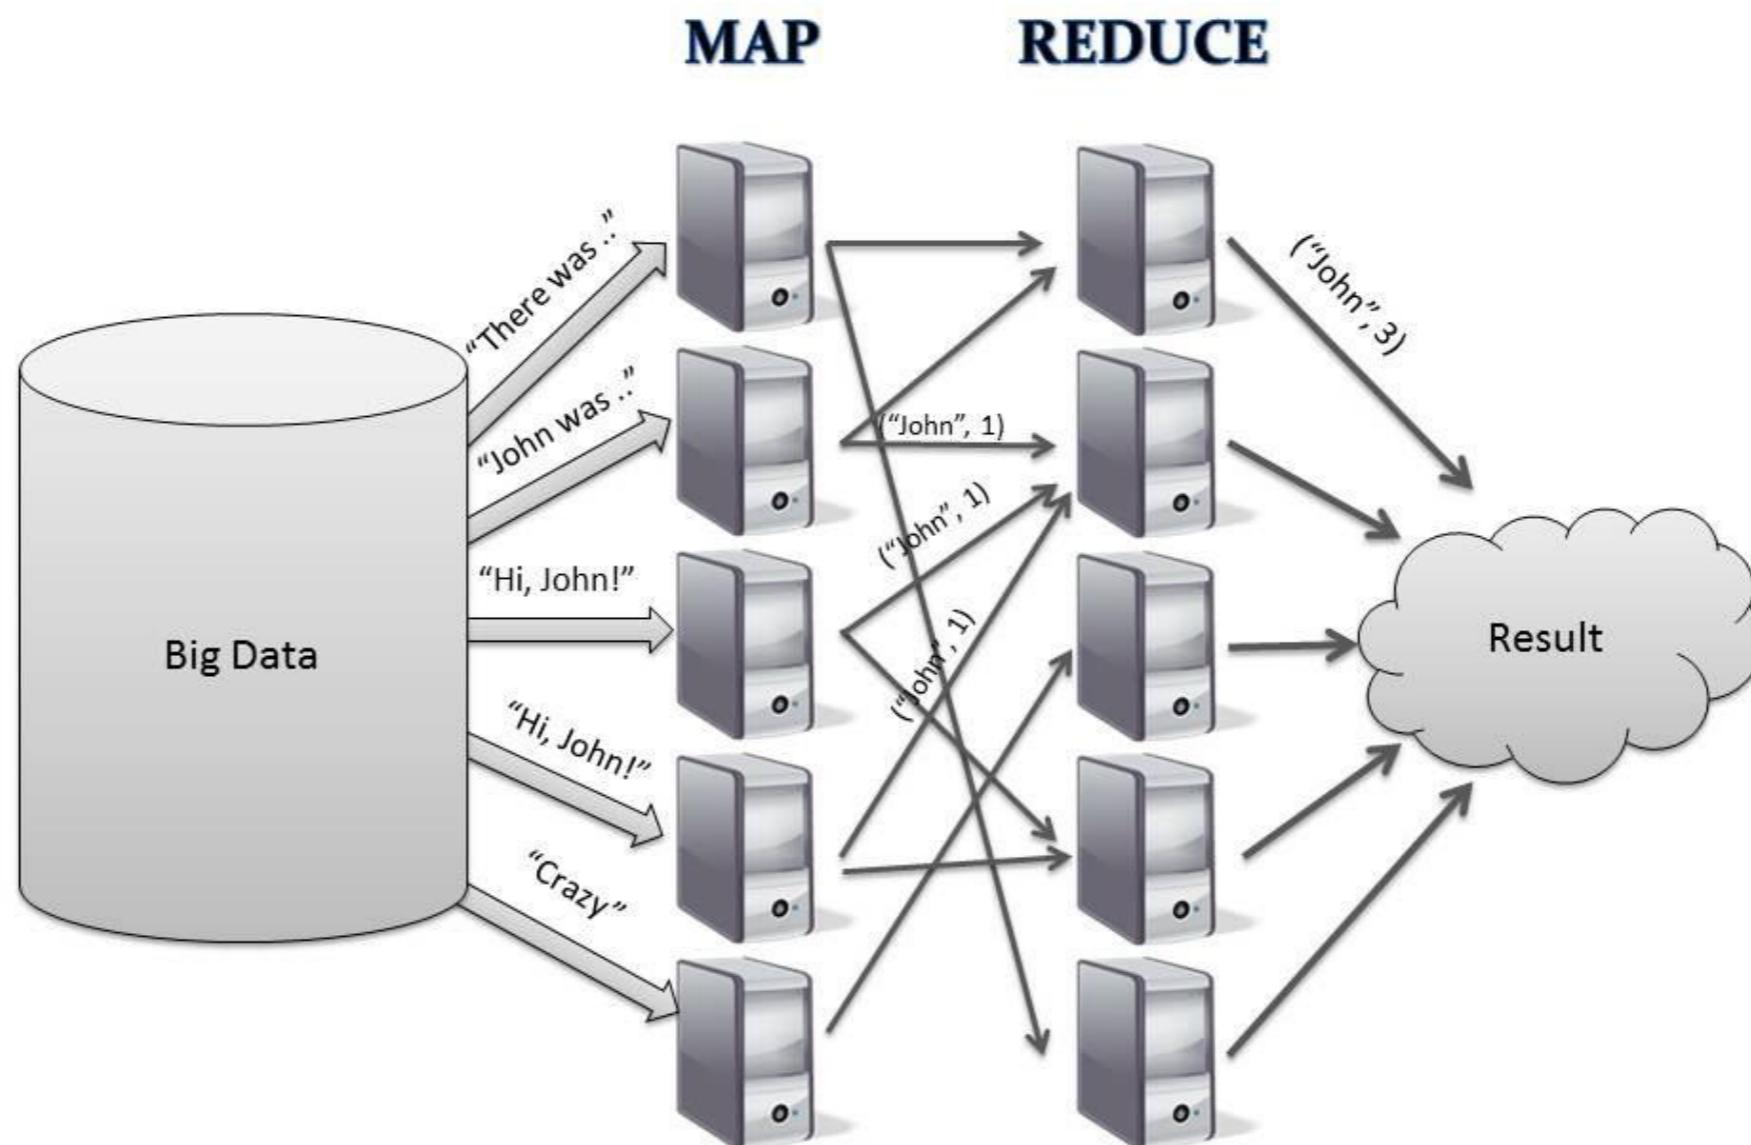
MapReduce Implementations

Owner	Name	Source Code
Microsoft	Dryad	https://github.com/MicrosoftResearch/Dryad
Apache	Hadoop	https://github.com/apache/hadoop
AMPLab (UC Berkley)	Spark	https://github.com/apache/spark

and many others...

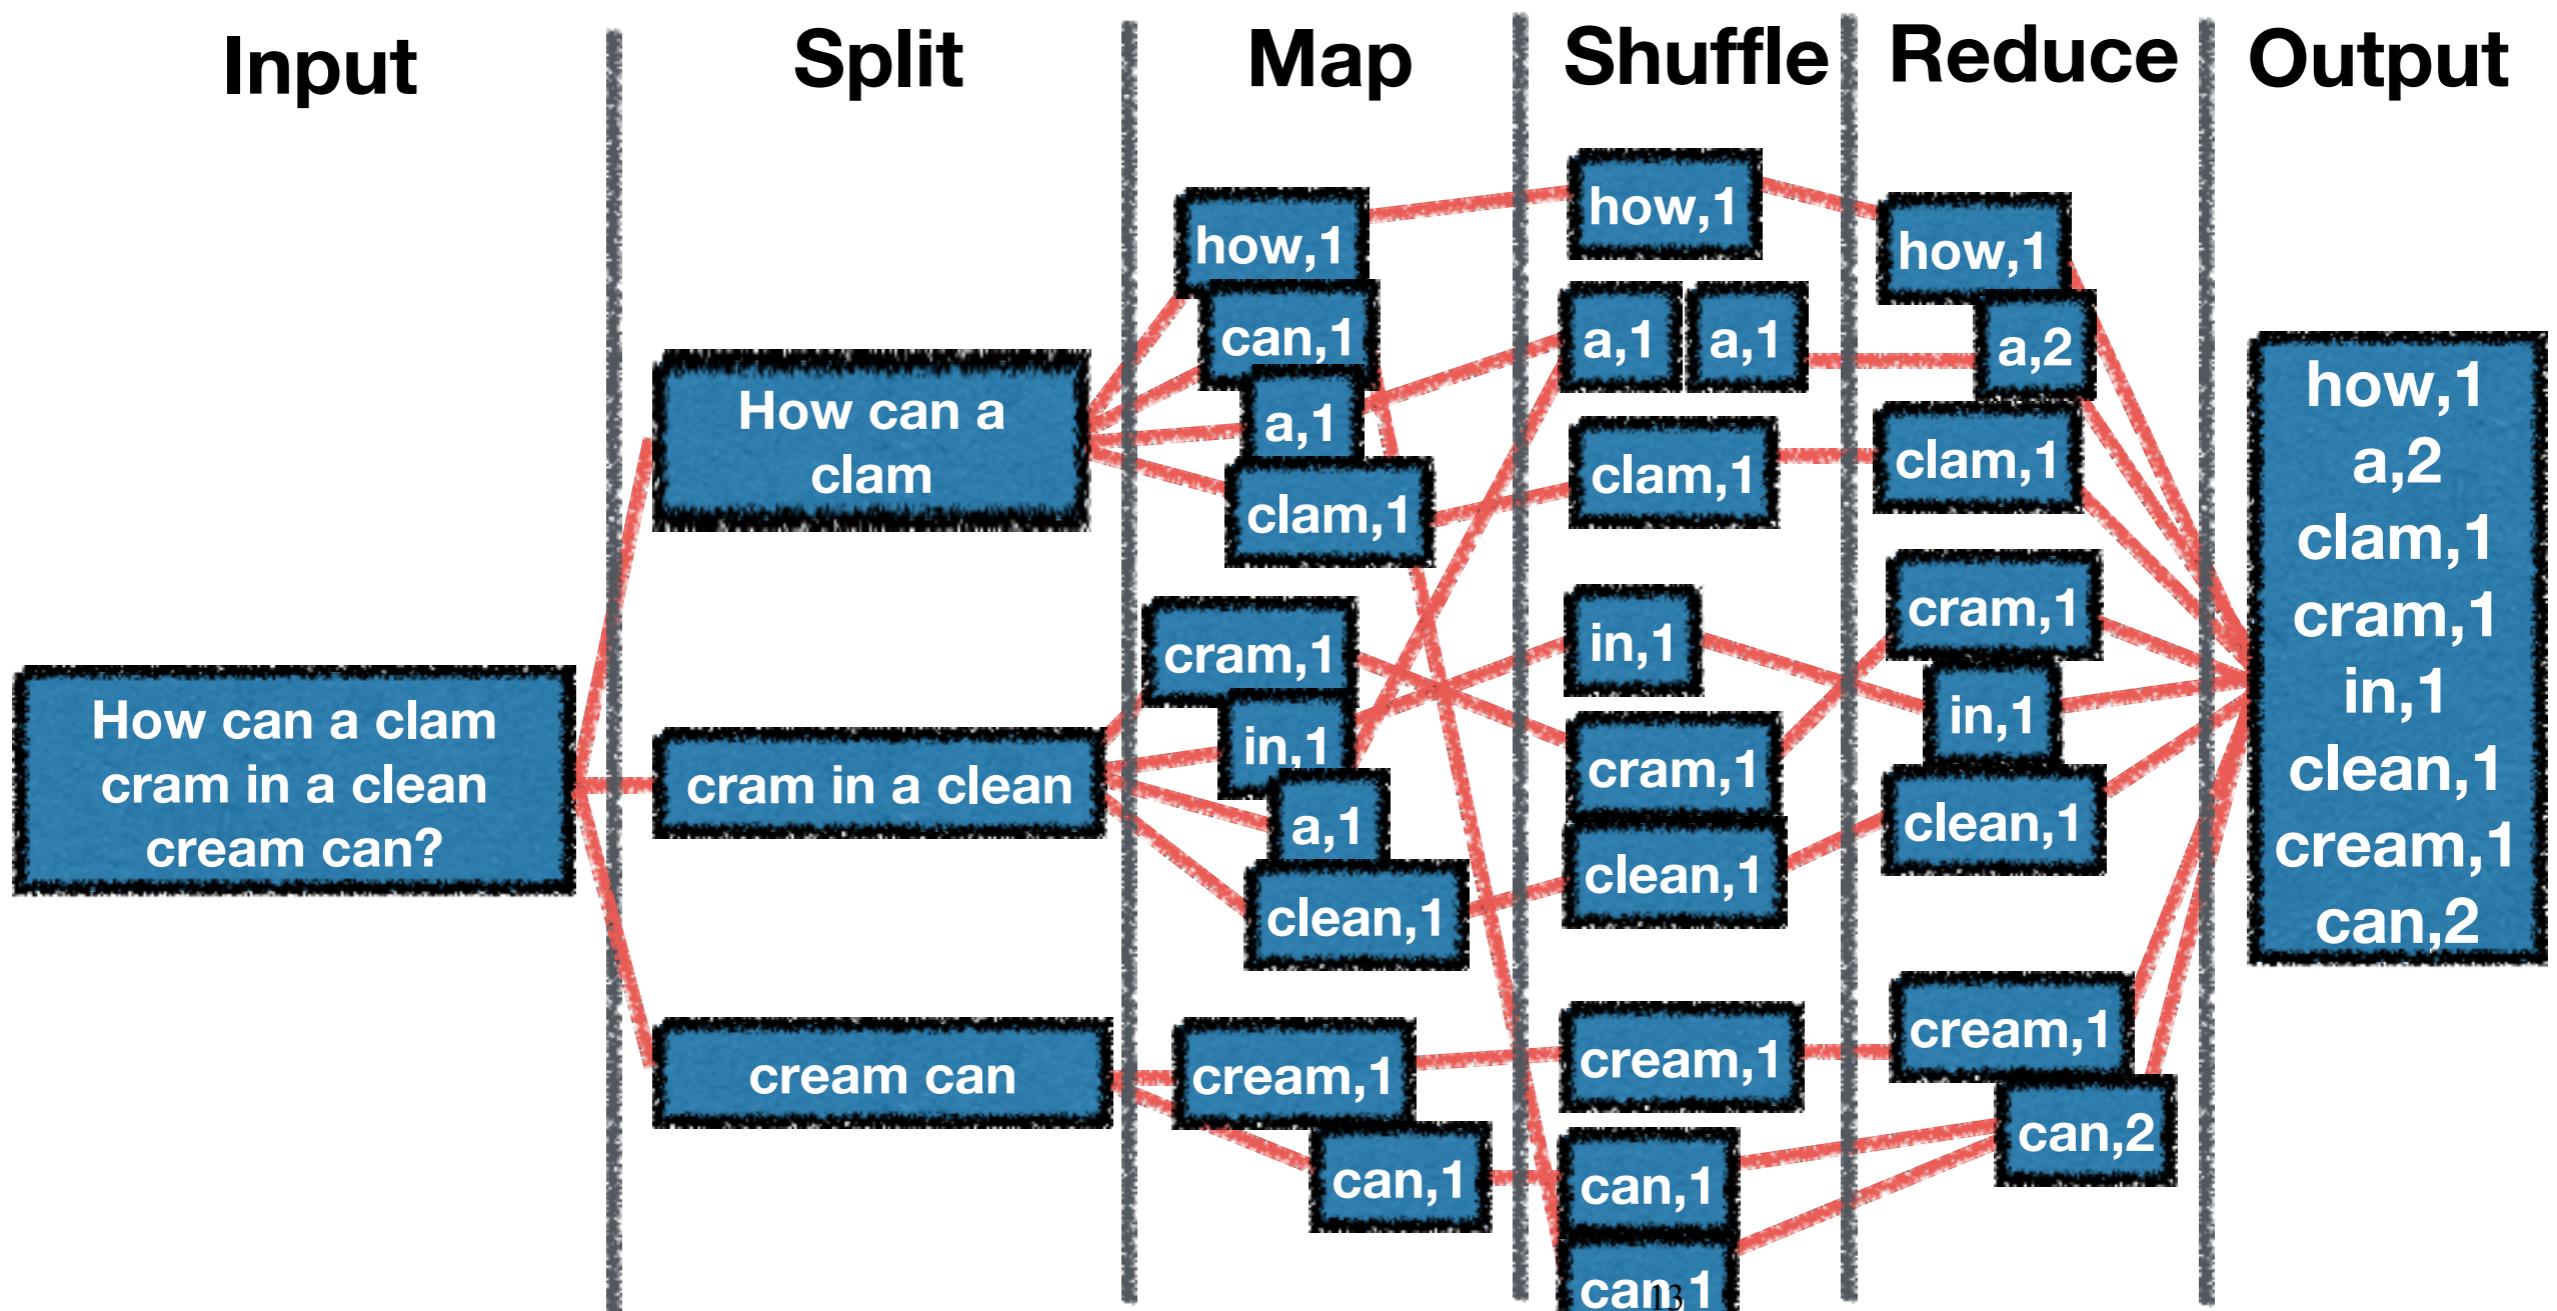
Logical Model

The WordCount Example



Logical Model

The WordCount Example



MapReduce Functions

Based on functional
programming paradigm

It's a programming model

Map(), Reduce(), Shuffle()...

MapReduce Pseudo-Code

```
function map(String name, String document):
    // name: document name
    // document: document contents
    for each word w in document:
        emit (w, 1)

function reduce(String word, Iterator partialCounts):
    // word: a word
    // partialCounts: a list of aggregated partial counts
    sum = 0
    for each pc in partialCounts:
        sum += ParseInt(pc)
    emit (word, sum)
```

MapReduce Code - Mapper

```
class Map extends Mapper<LongWritable, Text, Text, IntWritable> {  
    IntWritable      = new IntWritable(1);  
    Text word = new Text();  
  
    public void map(LongWritable key, Text value, Context context) throws {  
        String line = value.toString();  
        StringTokenizer tokenizer = new StringTokenizer(line);  
        while (tokenizer.hasMoreTokens()) {  
            word.set(tokenizer.nextToken());  
            context.write(word,      );  
        }  
    }  
}
```

MapReduce Code - Reducer

```
class Reduce extends Reducer<Text, IntWritable, Text, IntWritable> {  
  
    public void reduce(Text key, Iterator<IntWritable> values, Context context){  
  
        int sum = 0;  
        while (values.hasNext()) {  
            += values.next().get();  
        }  
        context.write(key, new IntWritable( ));  
    }  
}
```

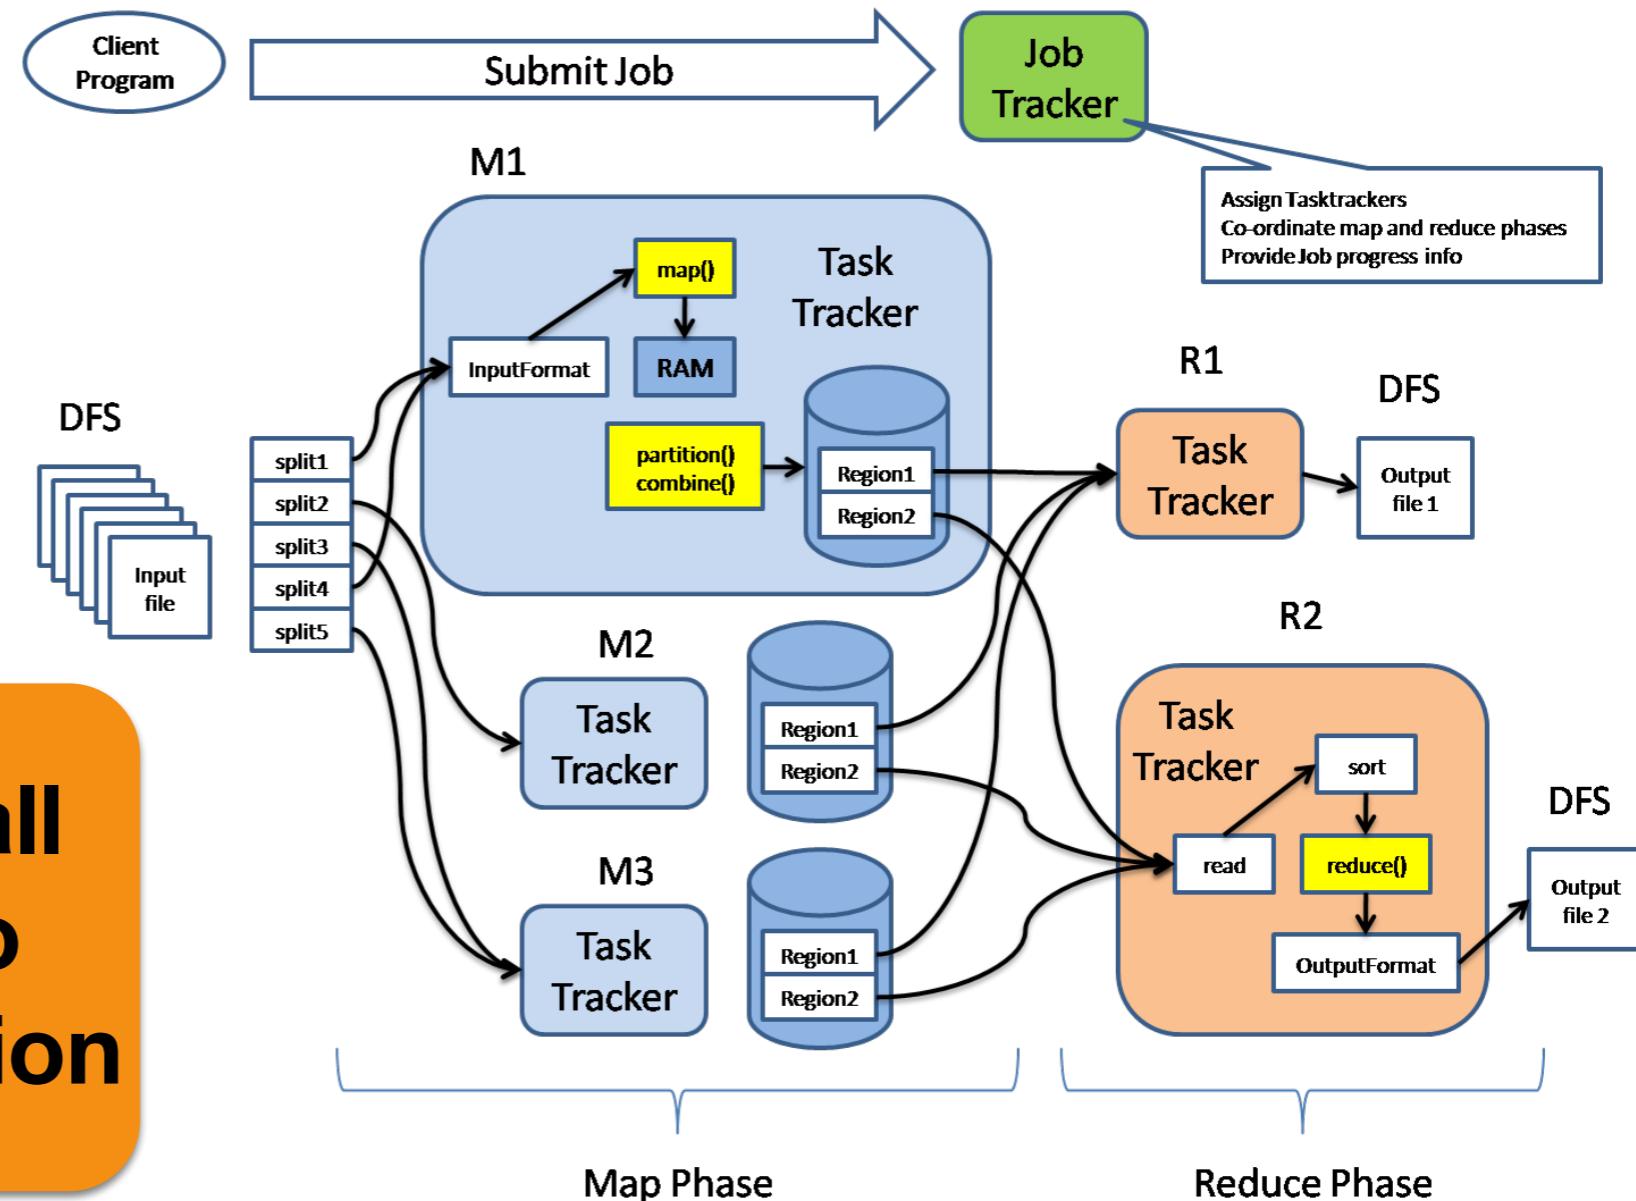
MapReduce Code - WordCount

```
public class WordCount {  
    public static class Map extends Mapper<LongWritable, Text, Text, IntWritable> { ... }  
    public static class Reduce extends Reducer<Text, IntWritable, Text, IntWritable> { ... }  
    public static void main(String[] args) throws Exception {  
        Configuration conf = new Configuration();  
        Job job = new Job(conf, "wordcount");  
        job.setOutputKeyClass(Text.class);  
        job.setOutputValueClass(IntWritable.class);  
        job.setMapperClass(Map.class);  
        job.setReducerClass(Reduce.class);  
        job.setInputFormatClass(TextInputFormat.class);  
        job.setOutputFormatClass(TextOutputFormat.class);  
        FileInputFormat.addInputPath(job, new Path(args[0]));  
        FileOutputFormat.setOutputPath(job, new Path(args[1]));  
        job.waitForCompletion(true);  
    }  
}
```

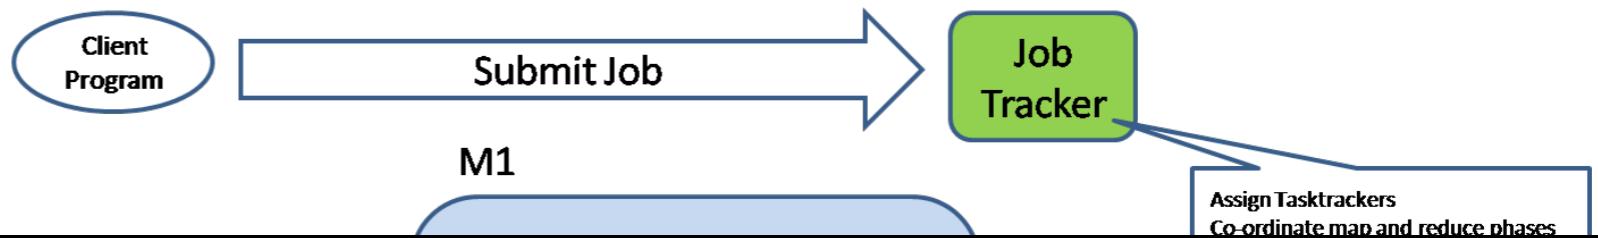
Simplified Distributed Execution

to simplify

Hide from the user all problems related to distributed computation

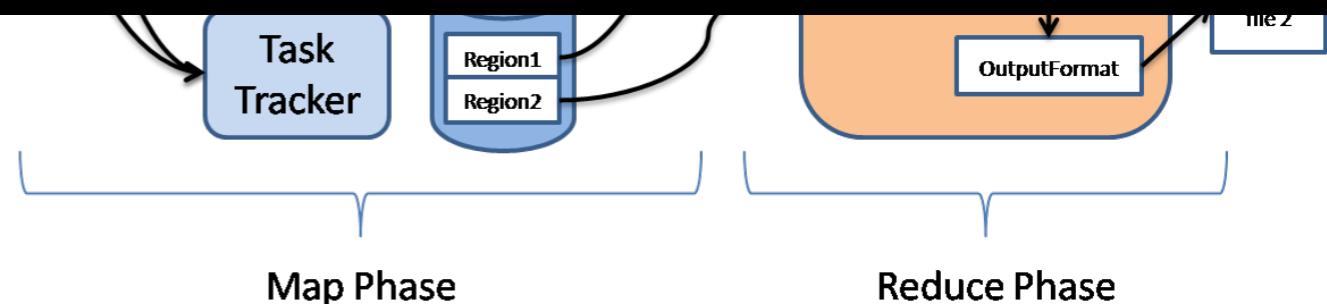


Simplified Distributed Execution



How does it work under the hood?

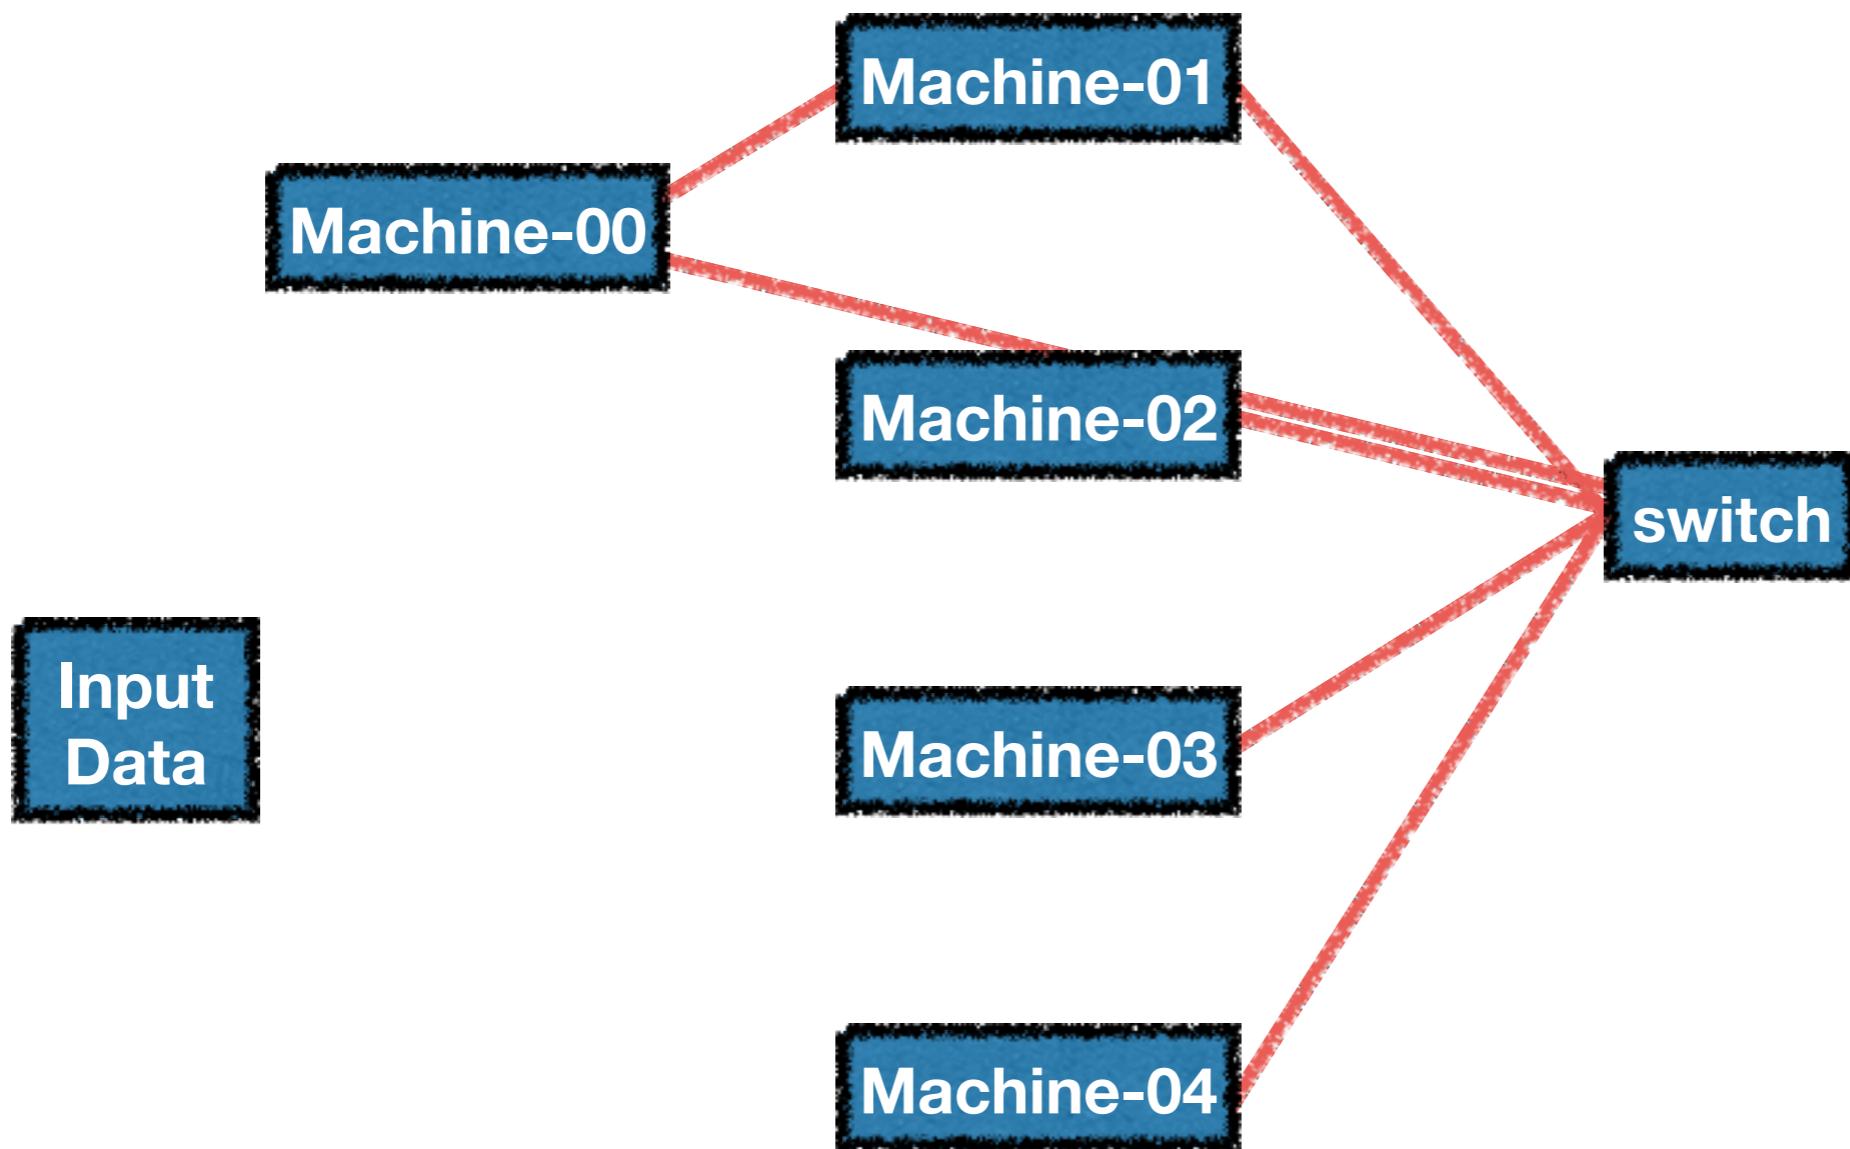
problems related to
distributed computation



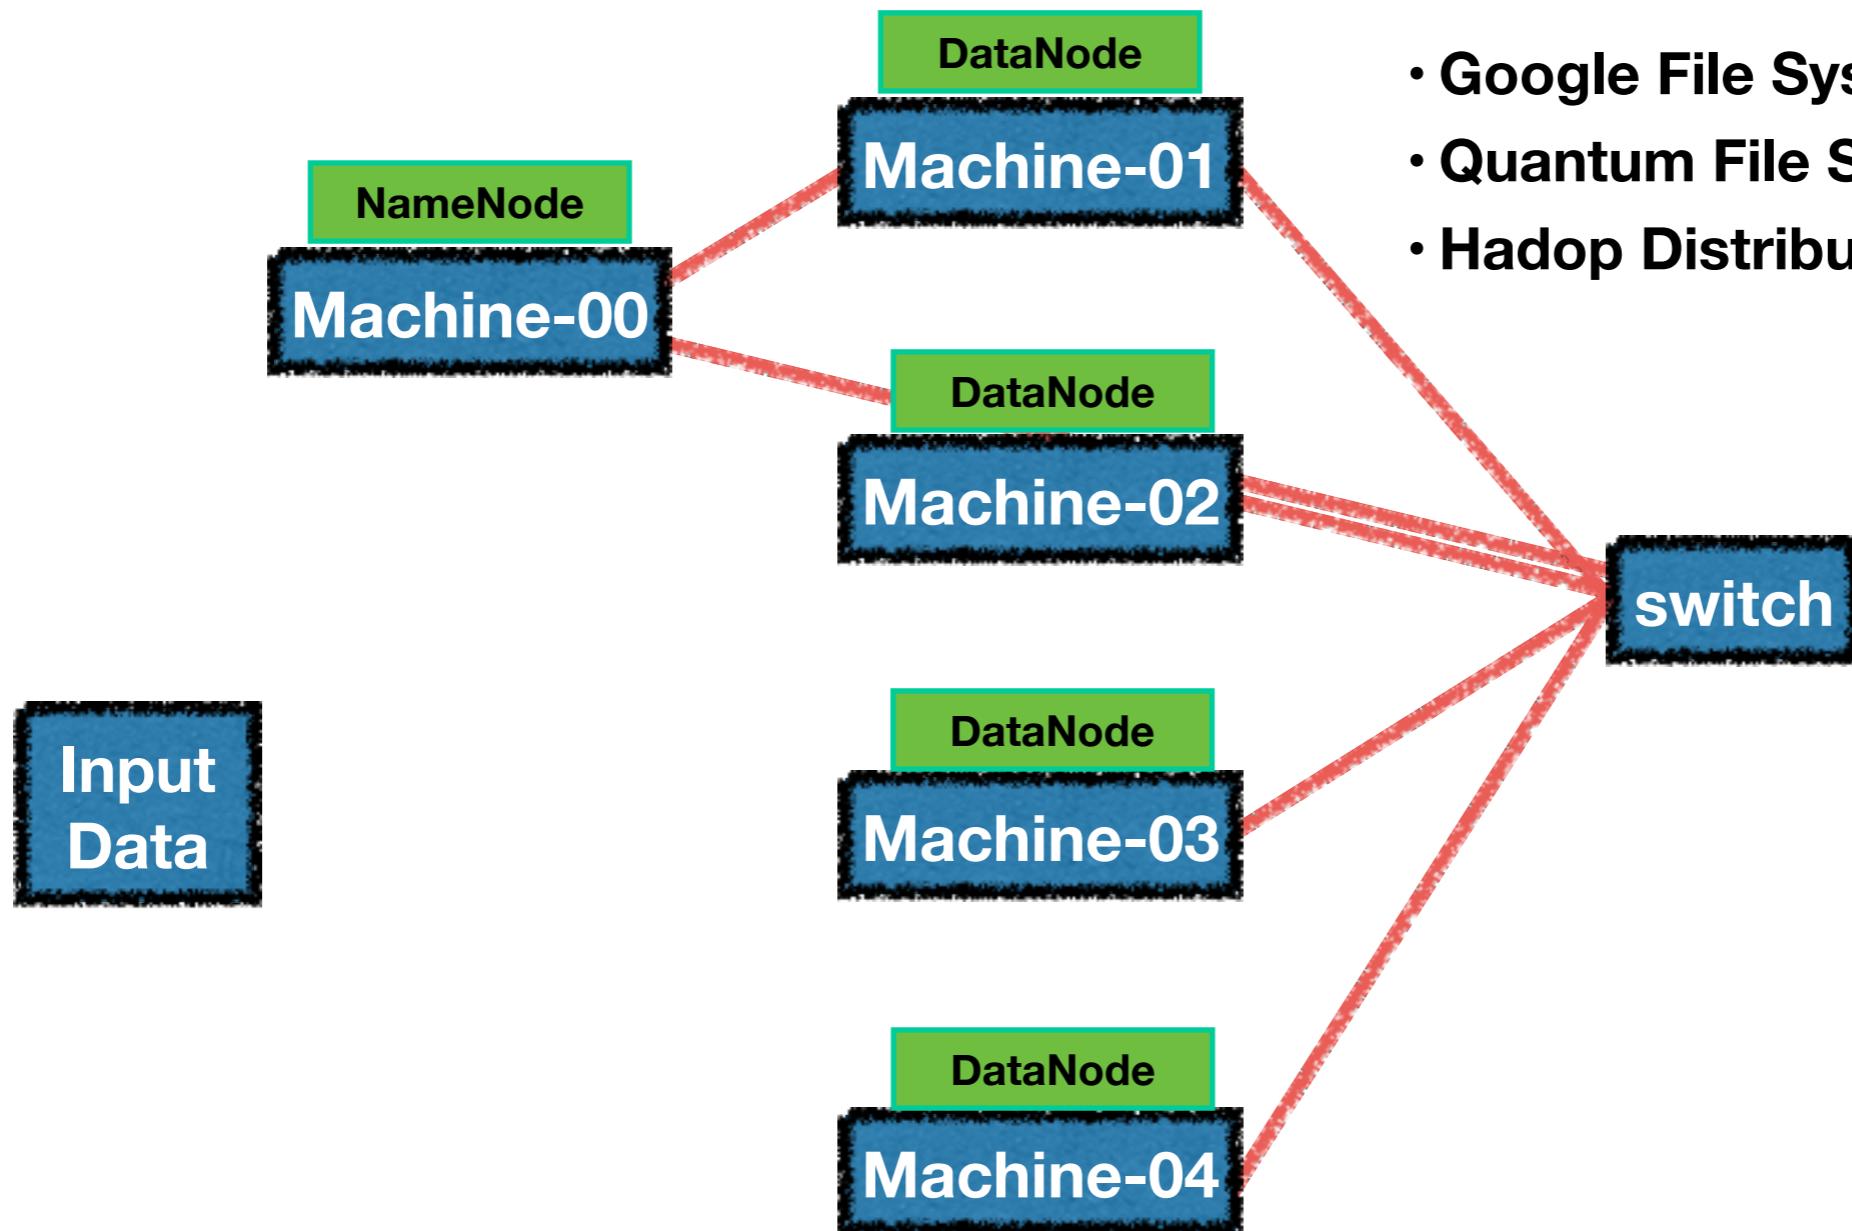
We have data...

Input
Data

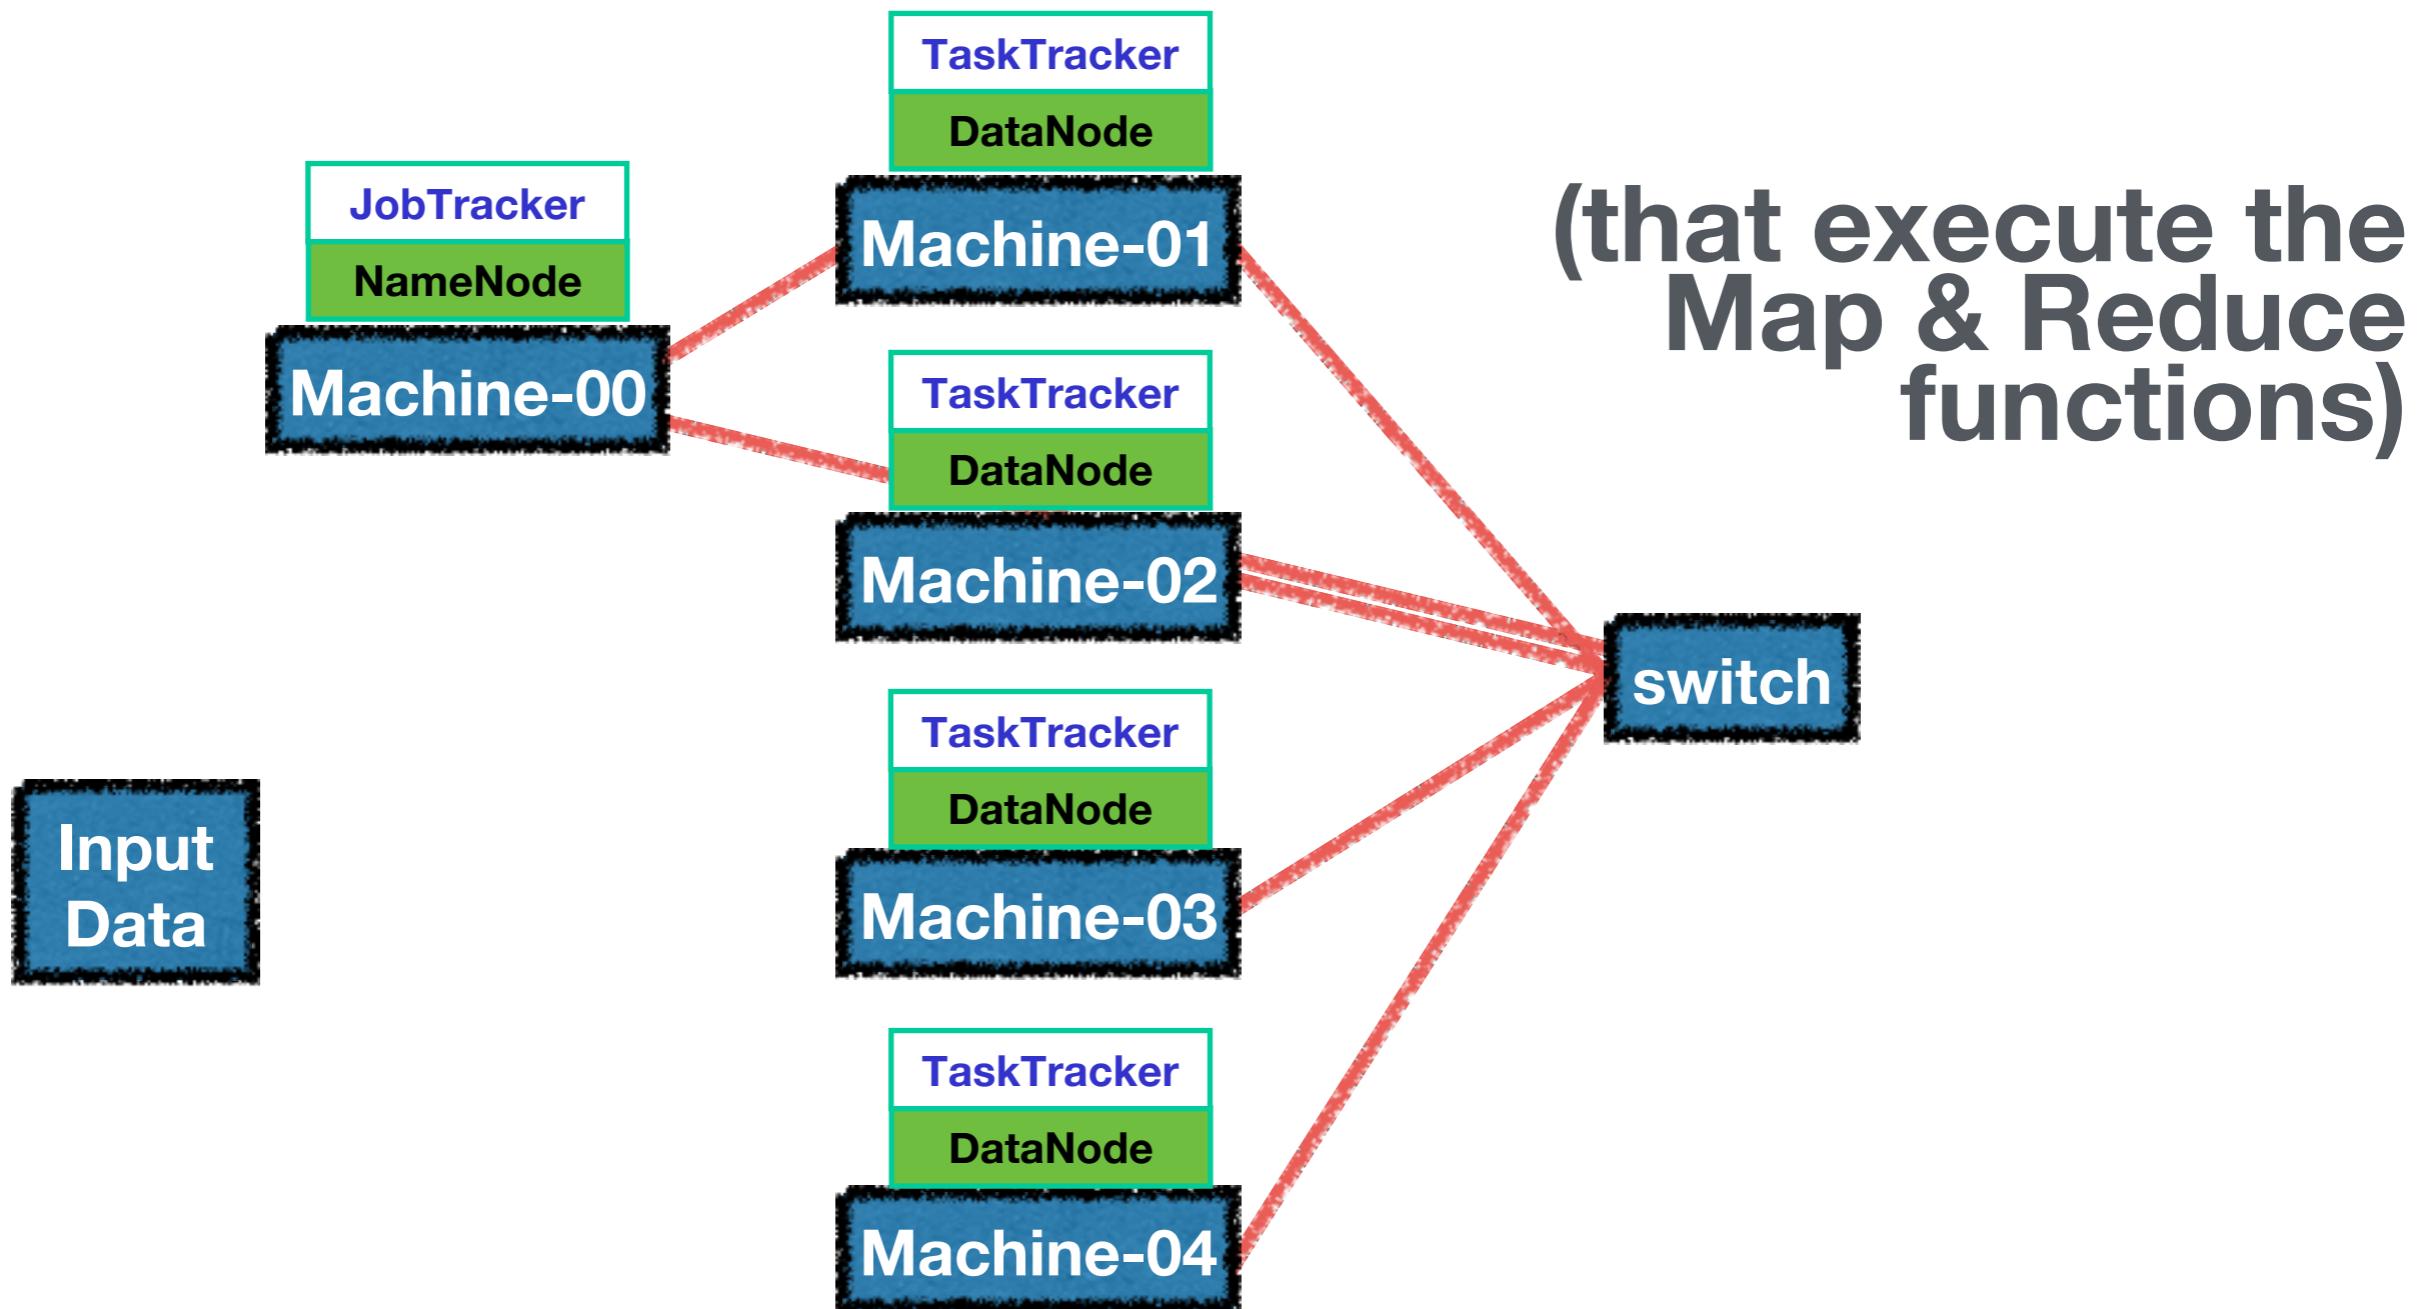
...we have machines,

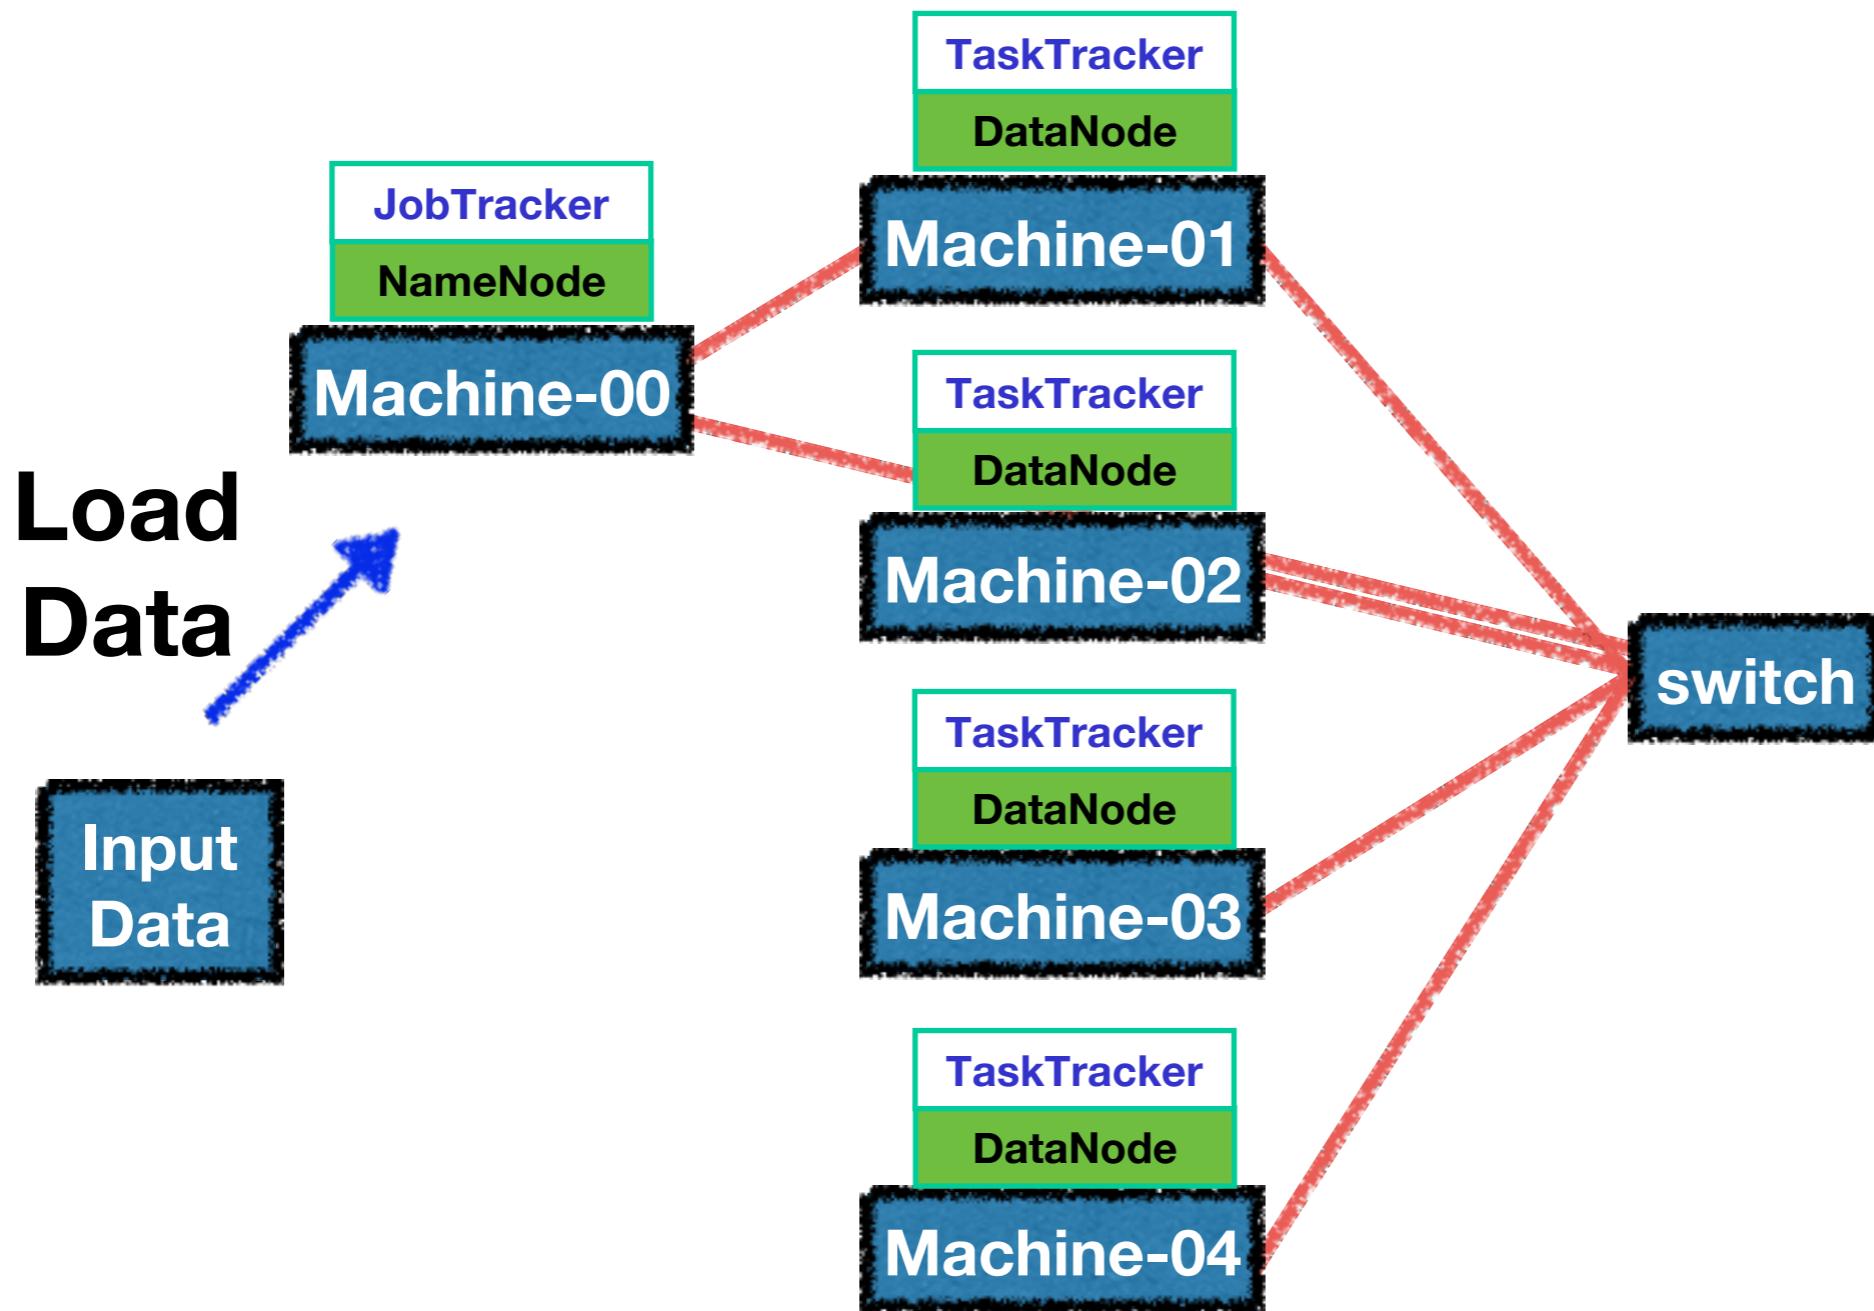


...we have processes for storing data,

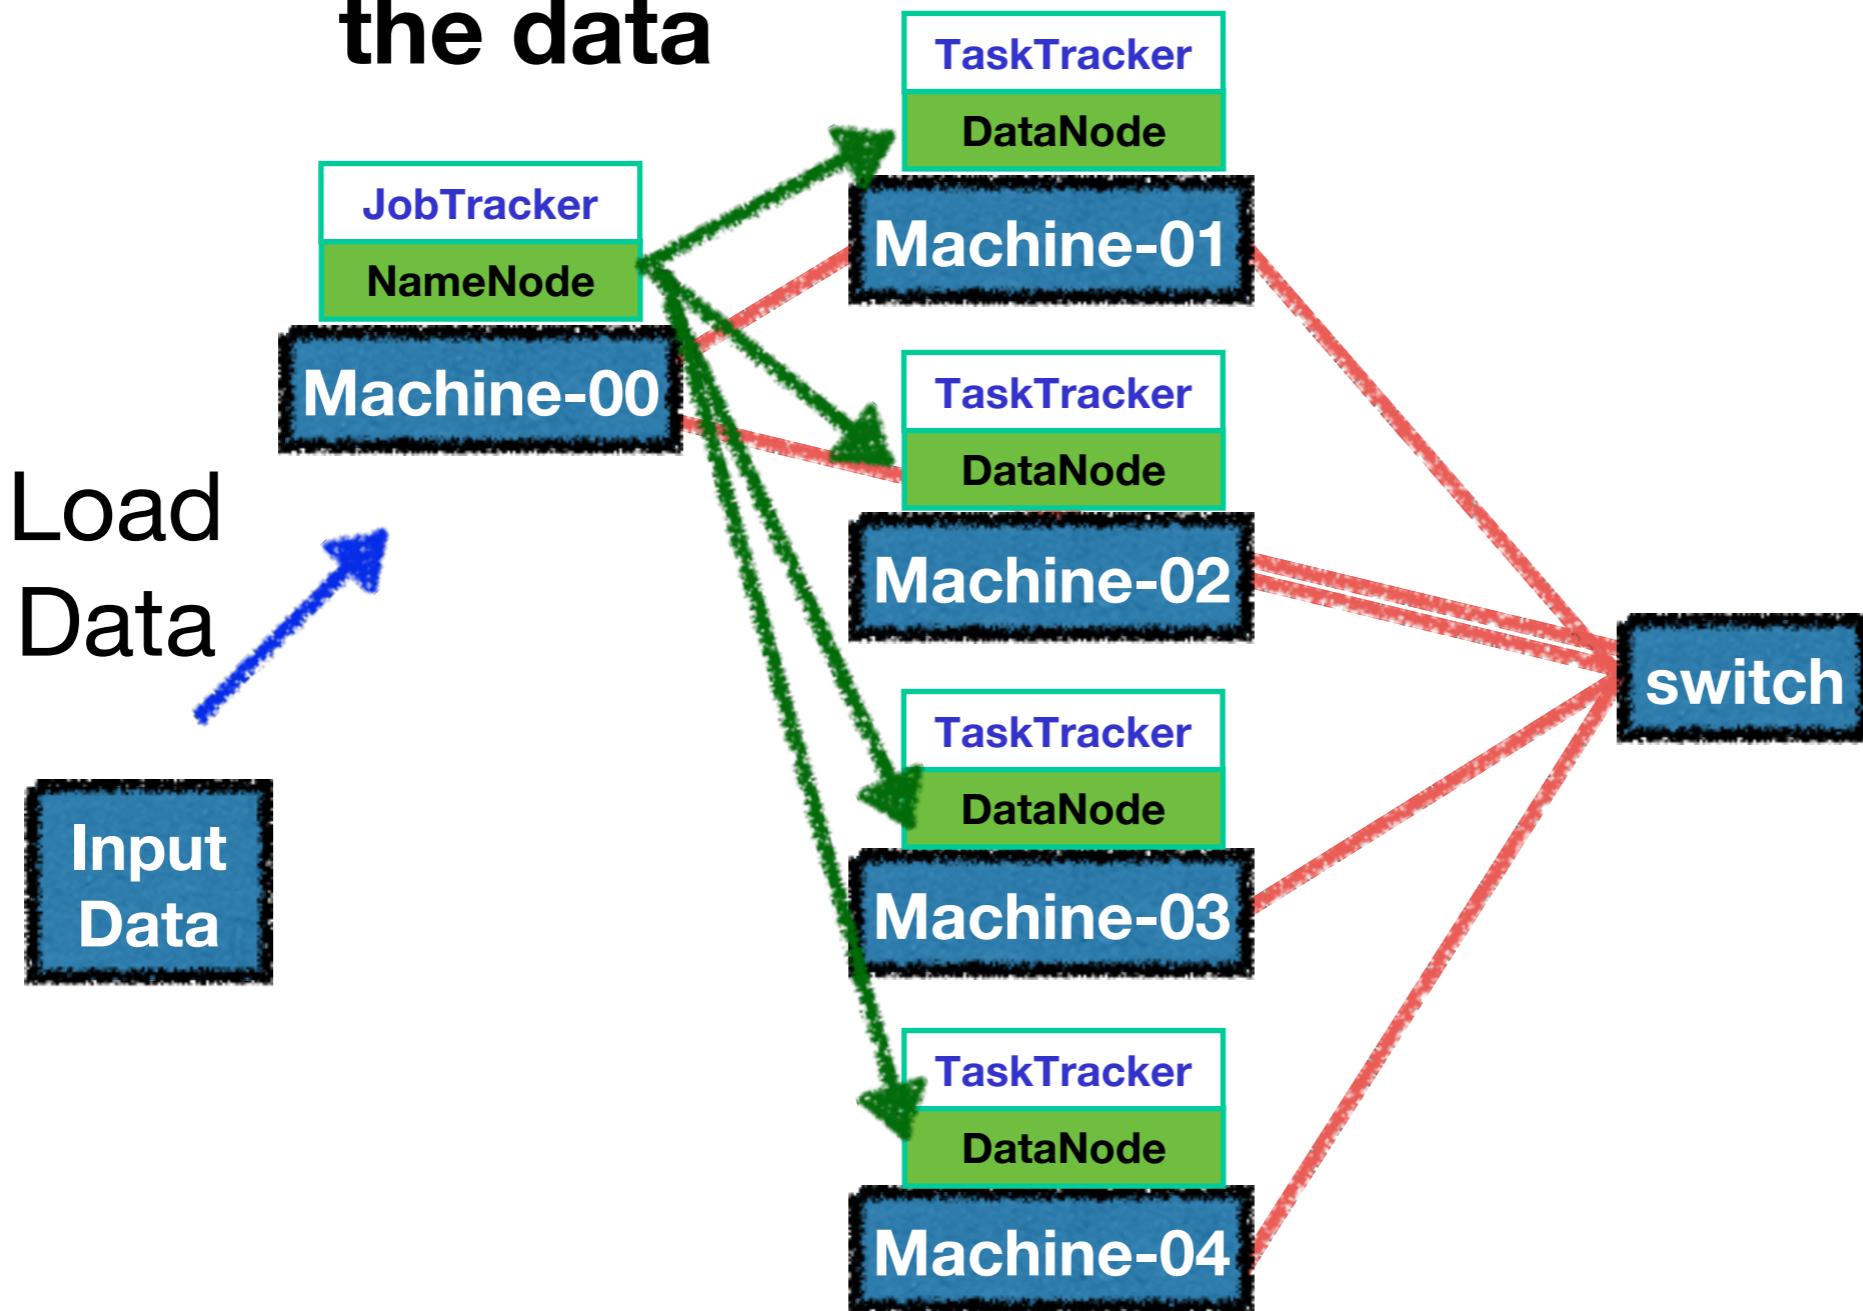


...and processes for processing data,

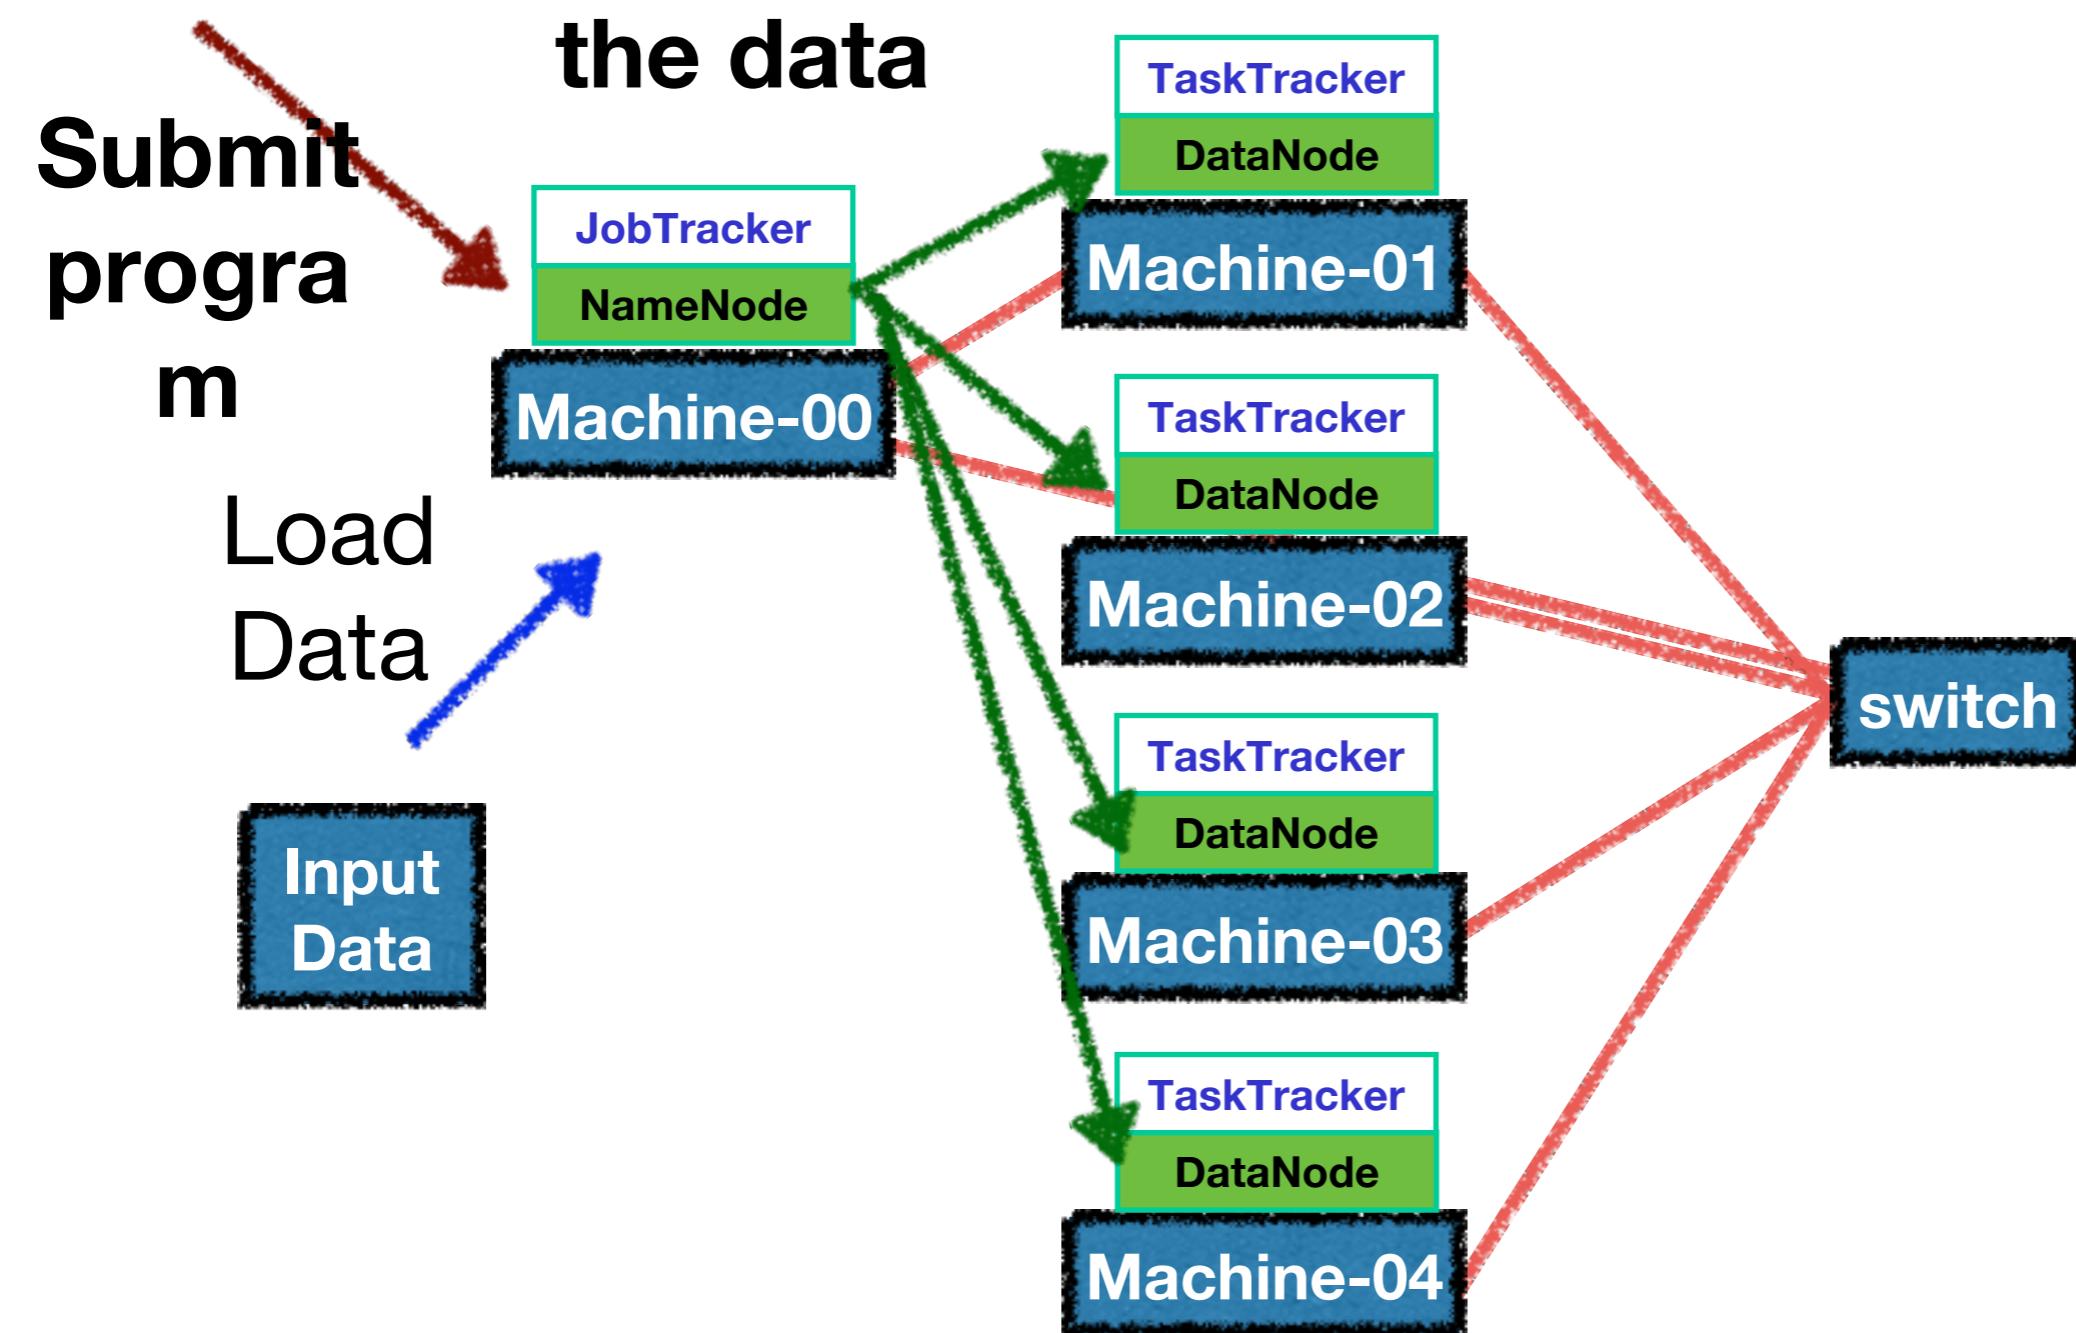




Distributes the data



Distributes the data

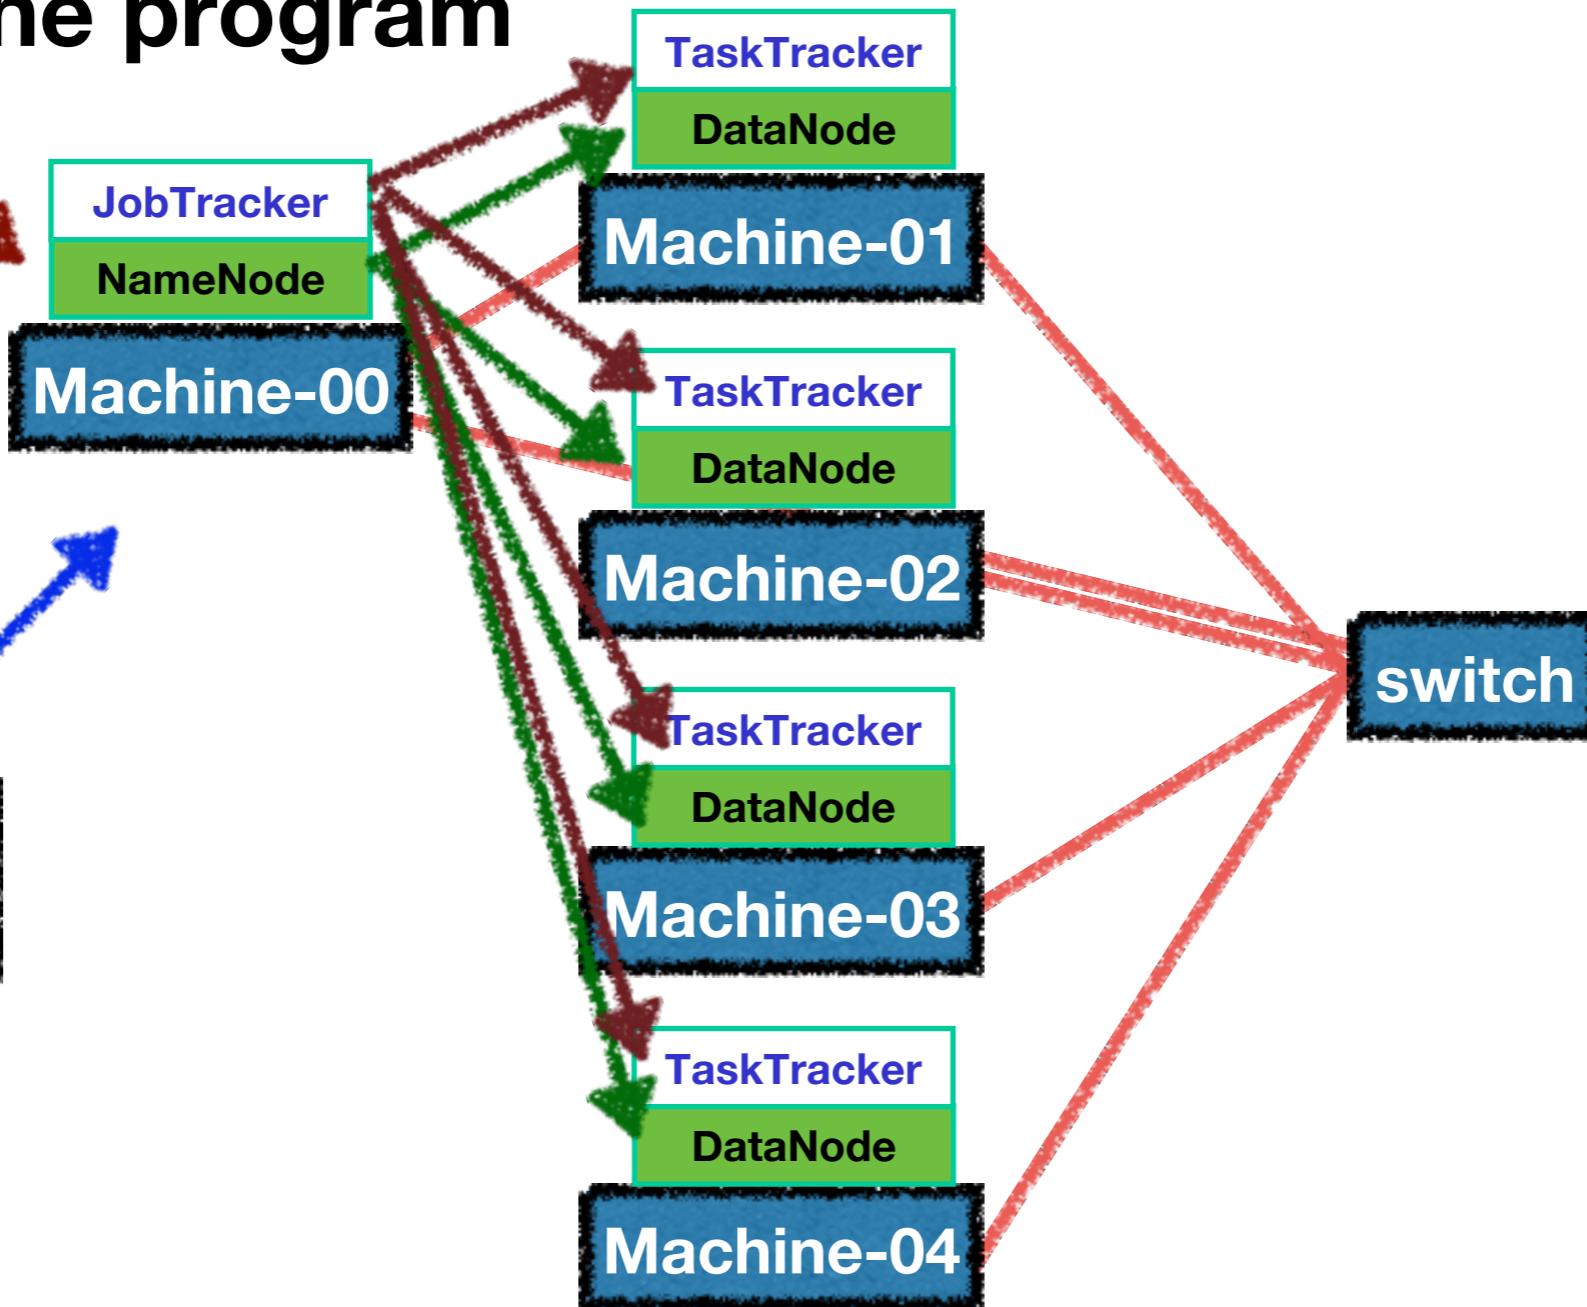


Executes the program

Submit
program

Load
Data

Input
Data

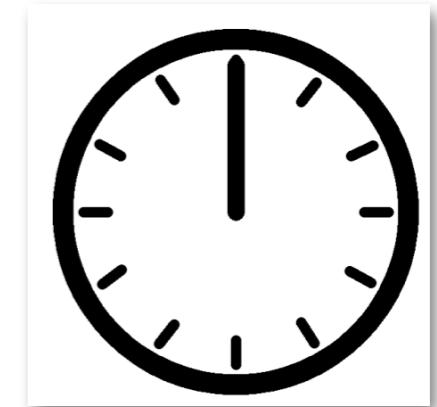
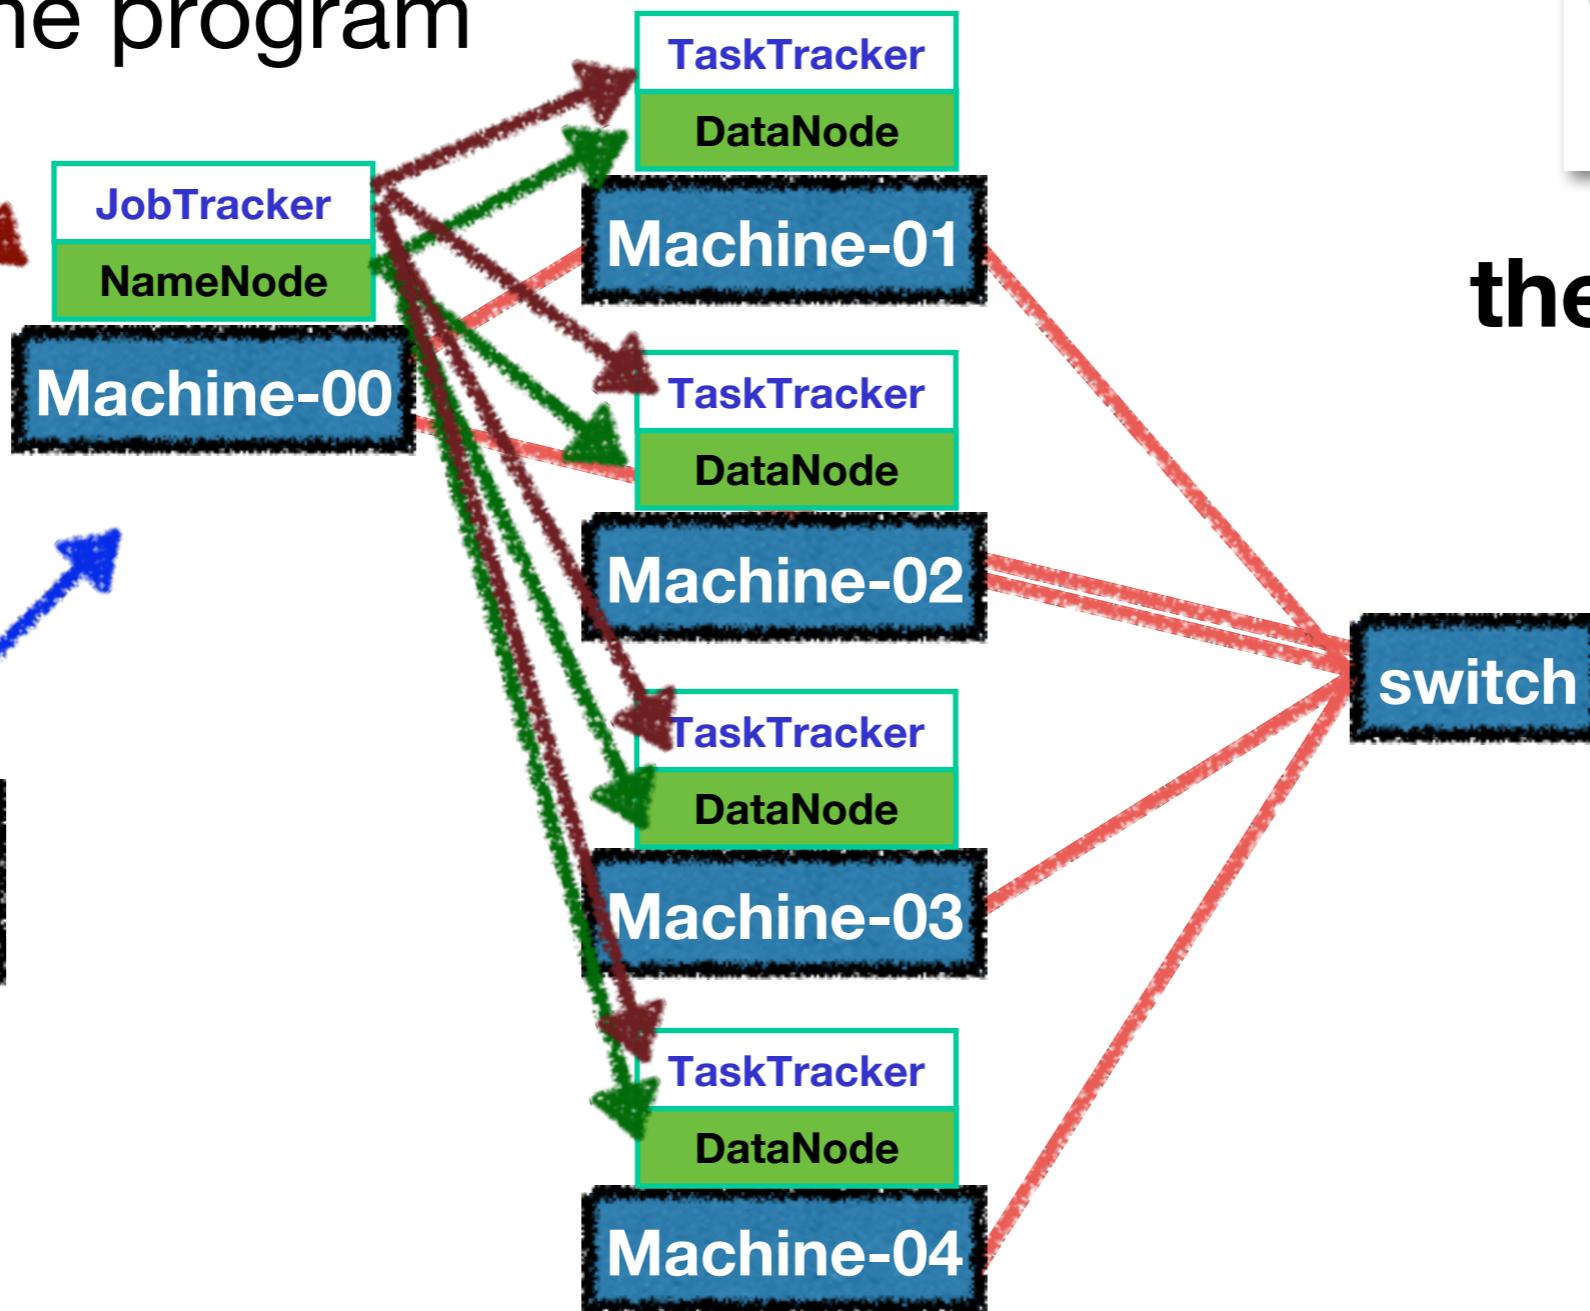


Submit program

Load Data

Input Data

Executes
the program

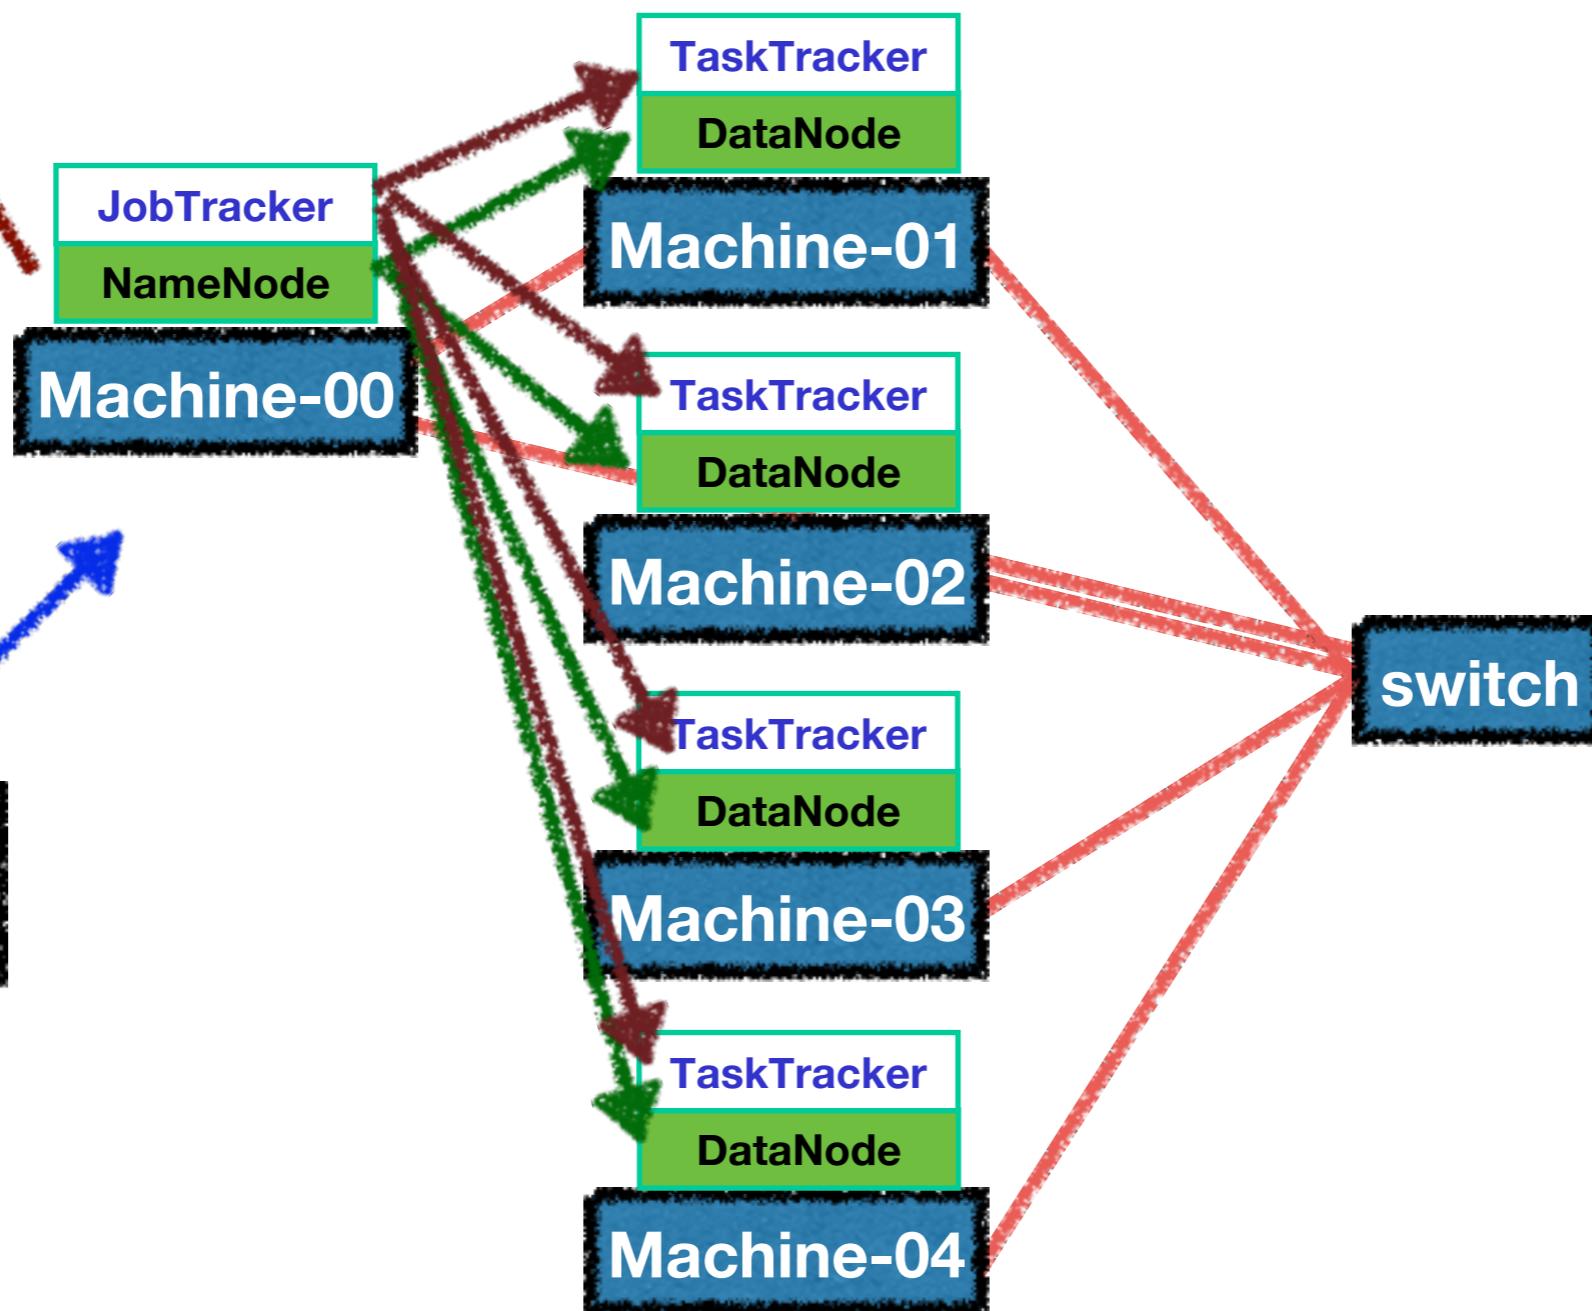


then we wait! :)

...and get
results

Load
Data

Input
Data

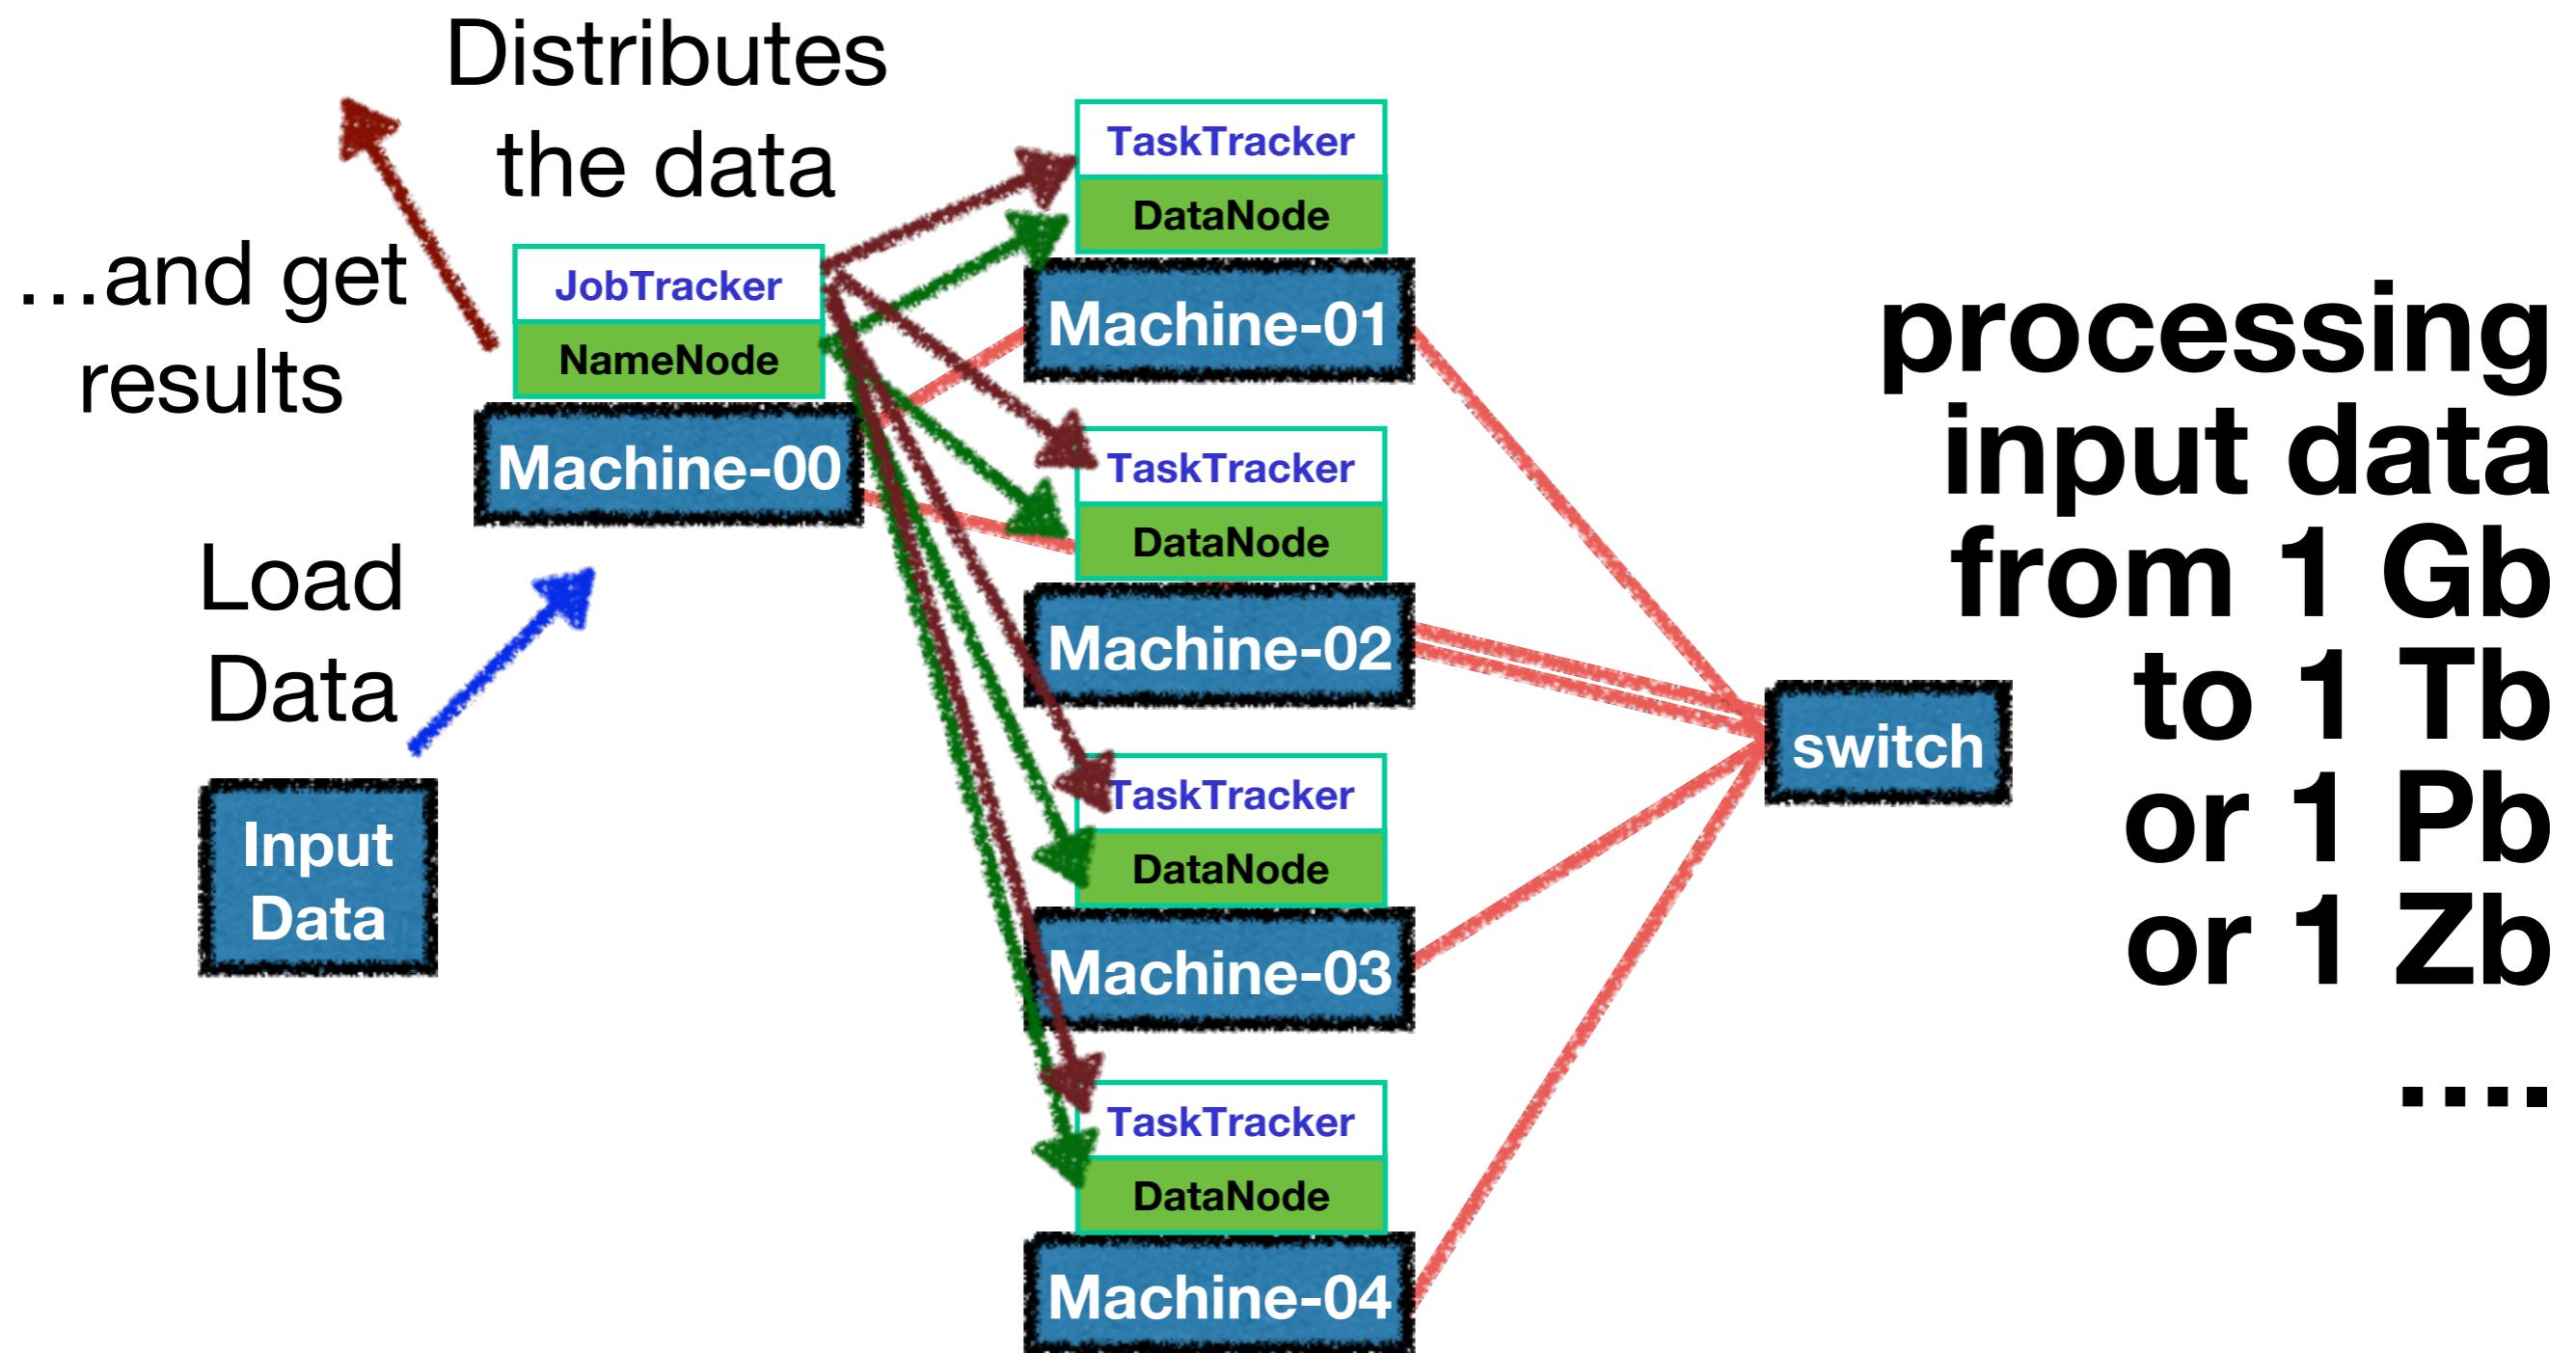


**MapReduce was designed
for batch processing**

**Not interactive or
streaming processing !**

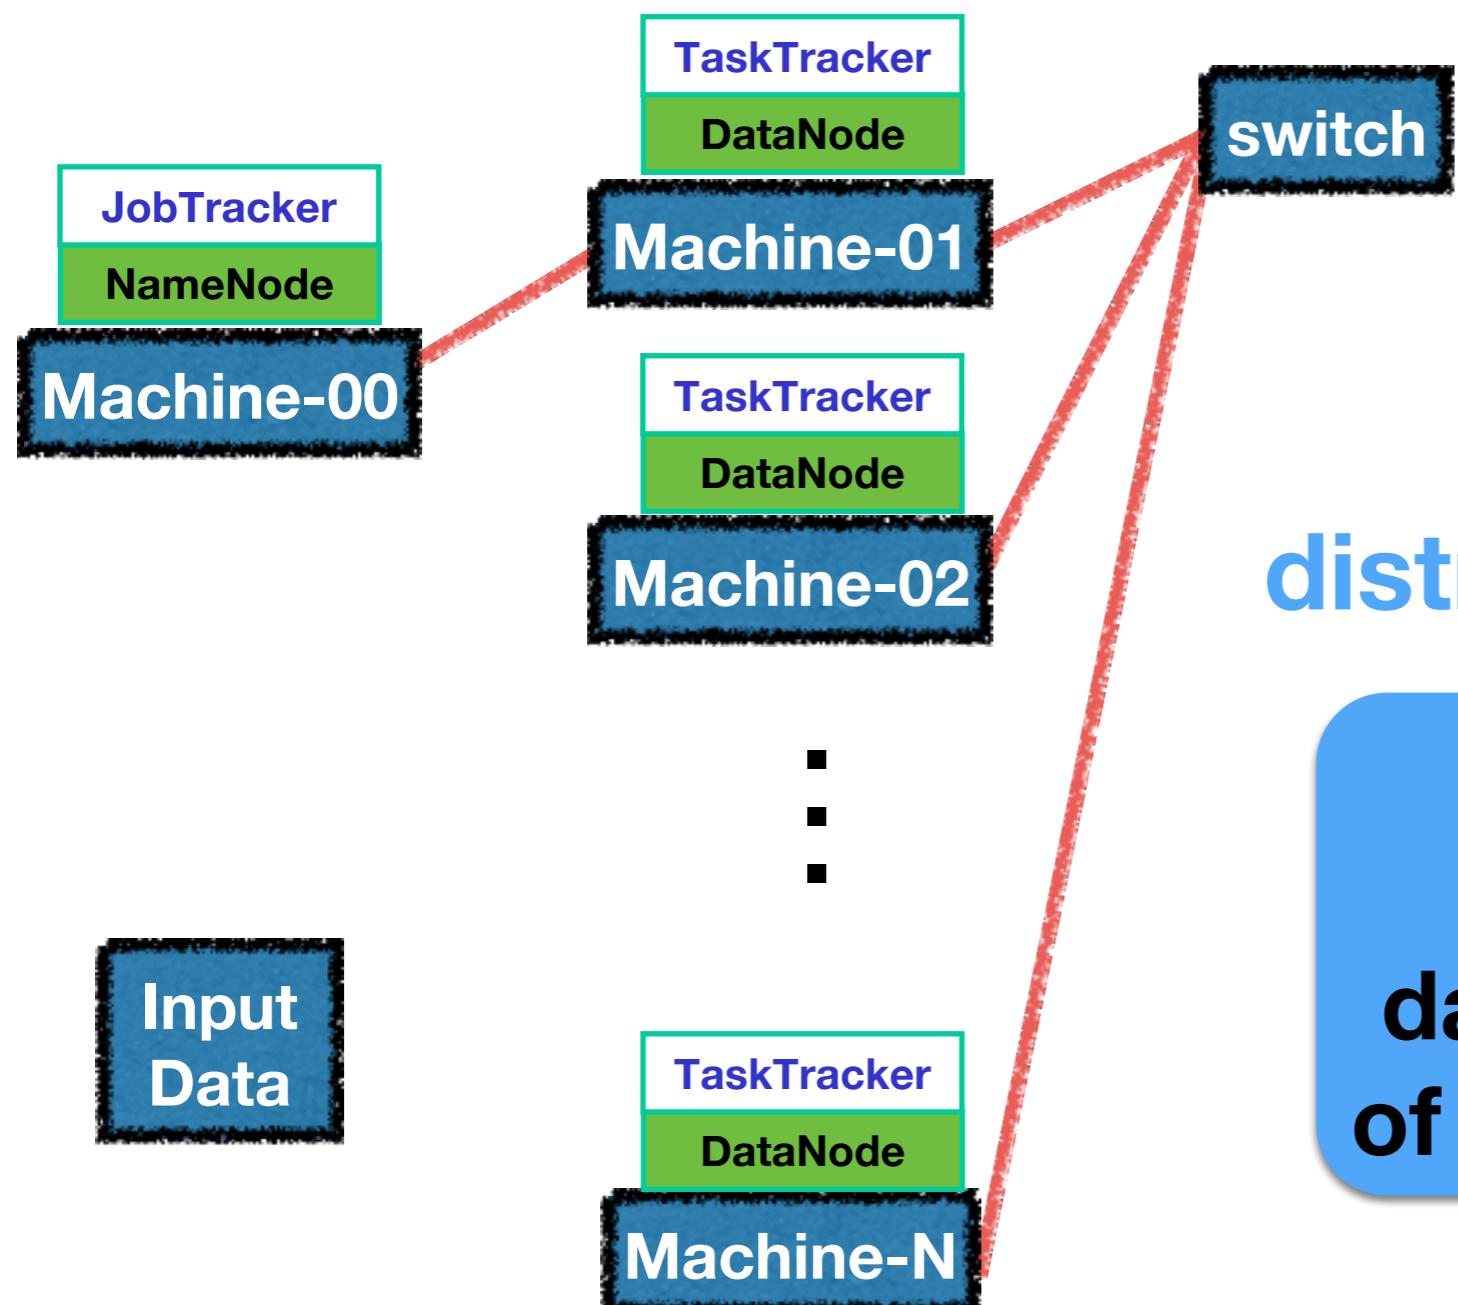
How does MapReduce Scale ?

How does it Scale?



How does it scale?

Just add more (cheap) machines

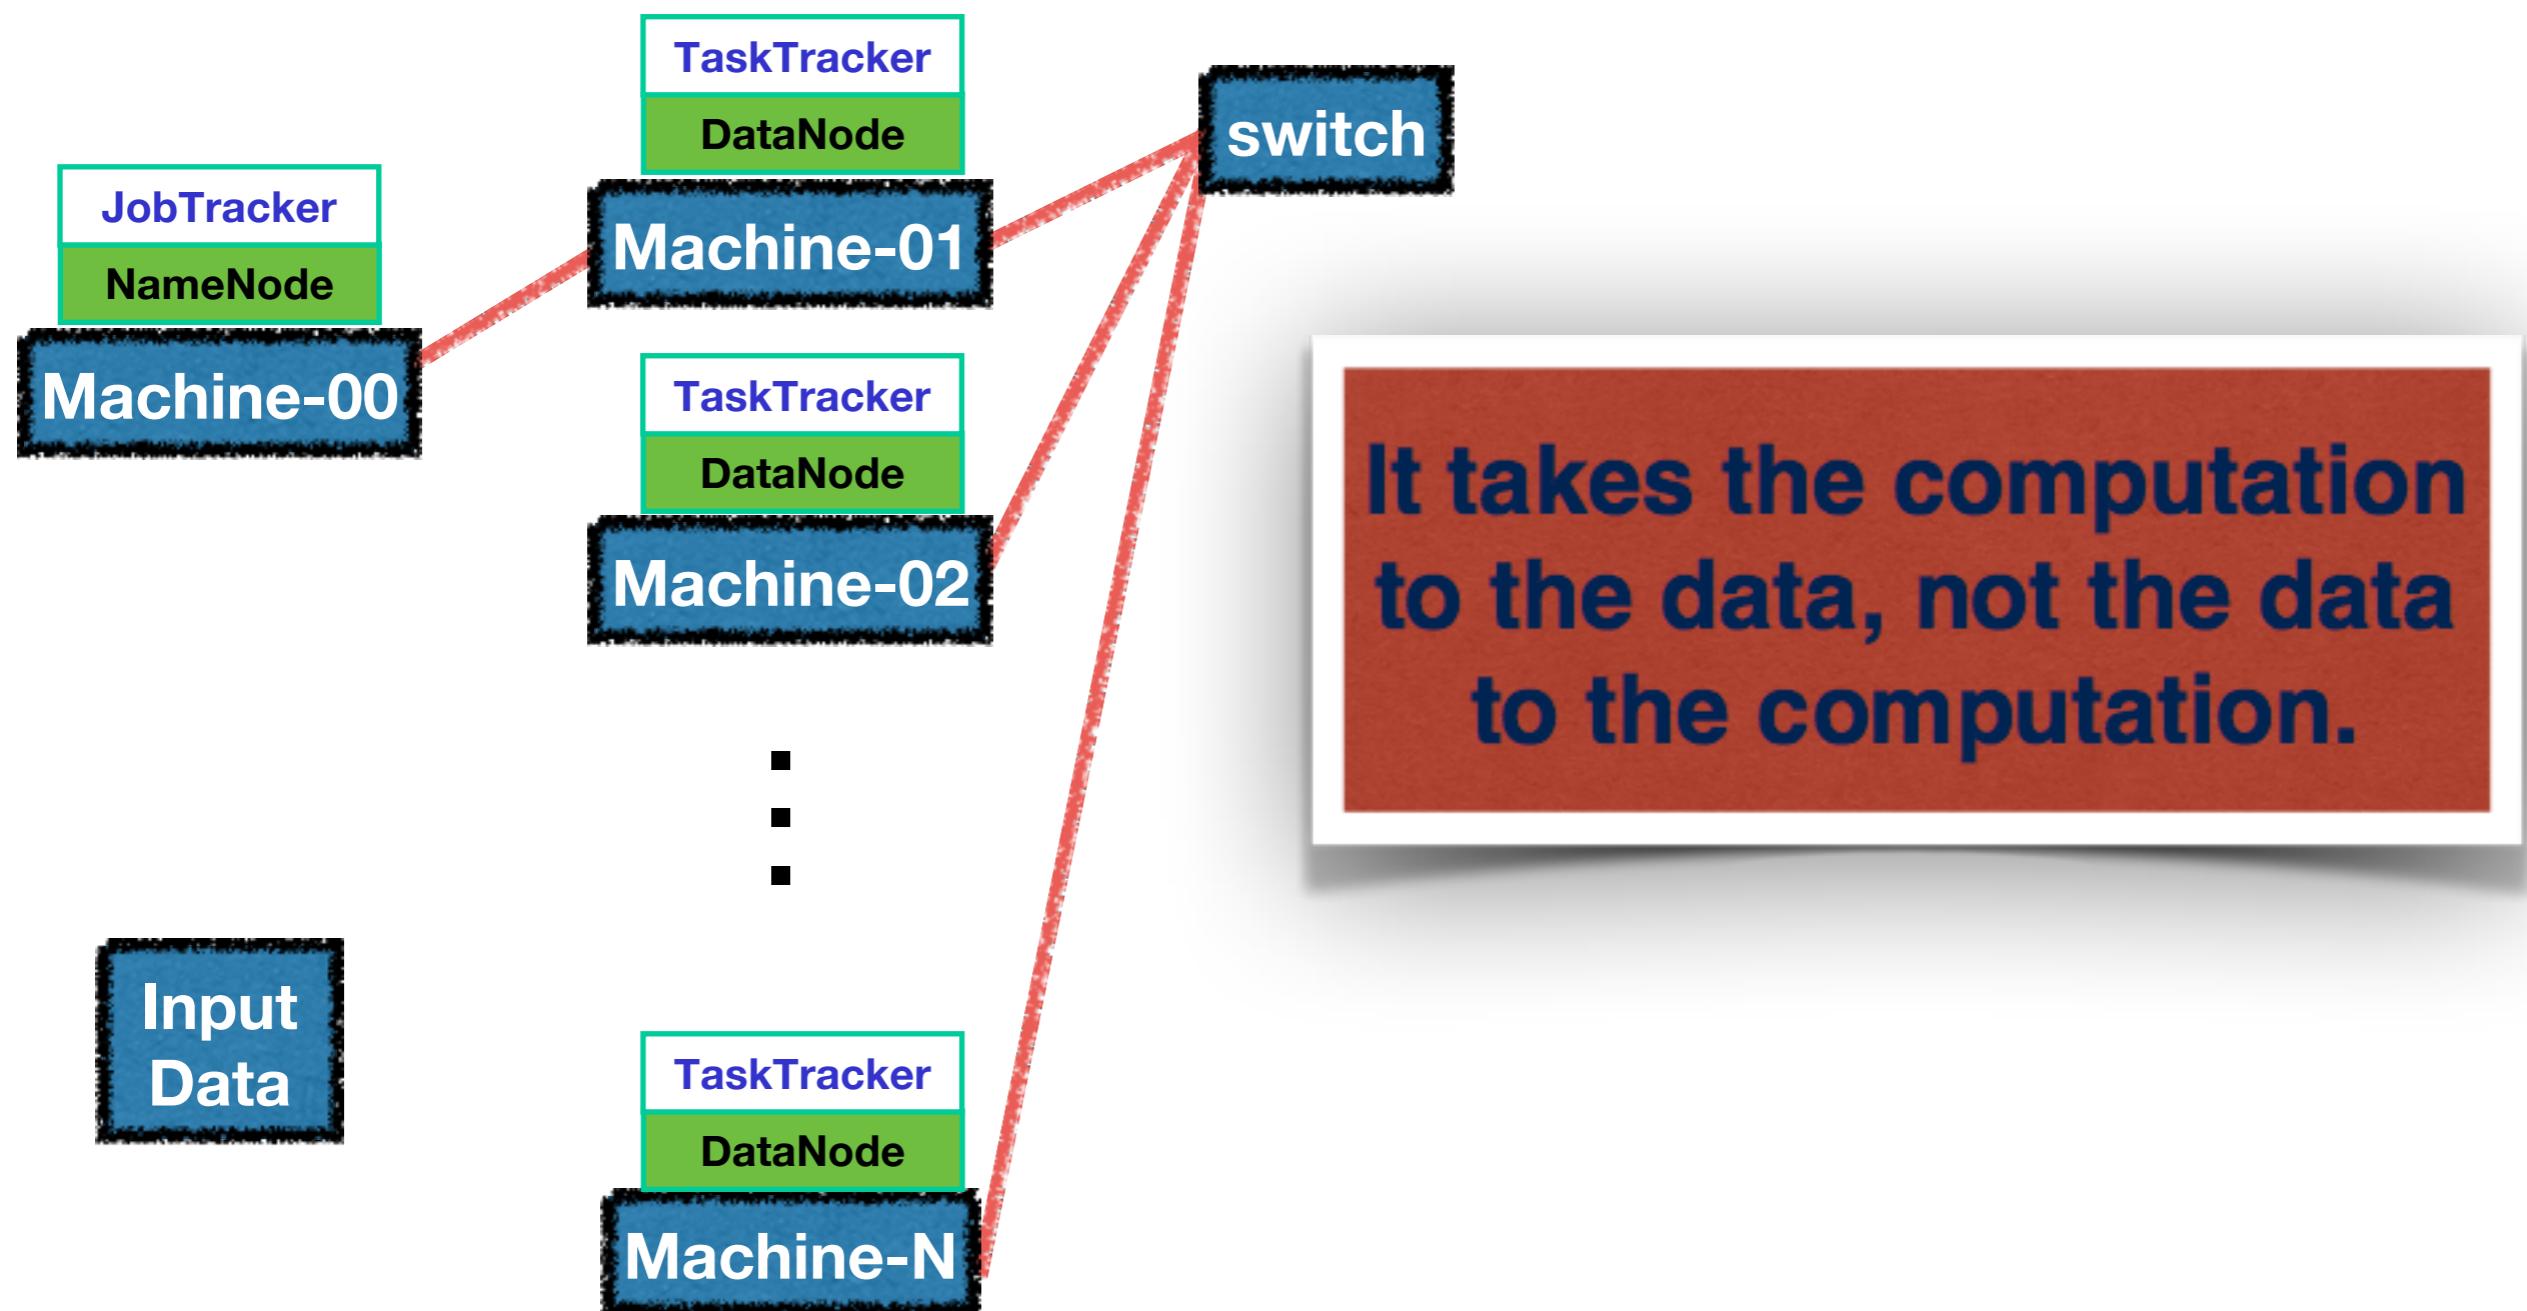


distributed data processing

**Allows distributed
processing of large
data sets across clusters
of thousands of machines**

How does it scale?

Just add more (cheap) machines



How do we build DBMS on top of MapReduce ?

How do we build databases on top of MapReduce?

We translate
SQL statements
into MapReduce
programs.

How do we build databases on top of MapReduce?

```
SELECT name FROM  
people WHERE age = '25';
```

**José,22
Mathew,25
Mirella,24
Letícia,24
Letícia,25**

How do we build databases on top of MapReduce?

```
SELECT name FROM  
people WHERE age = '25';
```

**José,22
Mathew,25
Mirella,24
Letícia,24
Letícia,25**

**May be any
kind of
unstructured
data**

How do we build databases on top of MapReduce?

**SELECT name FROM
people WHERE age = '25';**

```
function map(String name,  
           String document):  
    // name: document name  
    // document: document contents  
    array = document.split(',')  
    if (array[1] == '25')  
        emit (array[0], true)  
    else  
        emit (array[0], false)
```

José,22
Mathew,25
Mirella,24
Letícia,24
Letícia,25

How do we build databases on top of MapReduce?

**SELECT name FROM
people WHERE age = '25';**

```
function map(String name,  
           String document):  
    // name: document name  
    // document: document contents  
    array = document.split(',')  
    if (array[1] == '25')  
        emit (array[0], true)  
    else  
        emit (array[0], false)
```

José, false
Mathew, true
Mirella, false
Letícia, false
Letícia, true

How do we build databases on top of MapReduce?

```
SELECT name FROM  
people WHERE age = '25';
```

```
function reduce(String word,  
               Iterator idXOld):  
    // word: a word  
    // idXOld: list of ages that are 25  
    for each age in idXOld:  
        if (age == true)  
            emit (word)
```

José, false
Mathew, true
Mirella, false
Letícia, false
Letícia, true

How do we build databases on top of MapReduce?

```
SELECT name FROM  
people WHERE age = '25';
```

```
José,false  
Mathew,true  
Mirella,false  
Letícia,false  
Letícia,true
```

How do we build databases on top of MapReduce?

```
SELECT name FROM  
people WHERE age = '25';
```

- We can use just Map function
- We can perform many kinds of computation
(e.g., Graph, Machine Learning, SQL)

Readings

- [paper] Dean, Jeffrey and Ghemawat, Sanjay. MapReduce: Simplified Data Processing on Large Clusters (OSDI'04). Material at:
<http://dx.doi.org/10.1145/1327452.1327492>
- [implementation] The Apache™ Hadoop® project develops open-source software for reliable, scalable, distributed computing. <http://hadoop.apache.org/>

Questions?

Eduardo Cunha de Almeida <eduardo@inf.ufpr.br>
Edson Ramiro Lucas Filho <edson.lucas@uni.lu>